

Intro

- lessons related to my graduation project guided by xww
- bilibili links: <https://www.bilibili.com/video/BV1yt411w7h8>
- cuda c is used
- **thinking in a parallel way is important**

1 GPU Programming Model

- How to be faster?
 - faster clock
 - more work per clock cycle
 - parallel computing
- 2 kinds of PROCESSORS
 - CPU - complex control hardware
 - flexibility & performance
 - expensive in terms of power
 - more for control
 - GPU - simpler control hardware
 - more hardware for computation
 - potentially more power efficient
 - more restrictive programming model
- what can the programming model do
 - send data from CPU memory to GPU memory
 - send data from GPU memory to CPU memory
 - allocate GPU memory space
 - CPU launches a kernel on GPU
 - dynamic parallelism (may not be widely supported): GPU launches a kernel on GPU
 - a typical GPU program
 - `cudaMalloc`: CPU allocates storage on GPU
 - `cudaMemcpy`: CPU copies input data from CPU to GPU
 - `launchKernel`: CPU launches kernels on GPU to process the data
 - `cudaMemcpy`: CPU copies results back to CPU from GPU
 - ...
 - **BIG IDEA**
 - kernels look like serial programs
 - write the program as if it will run on one single thread, but the GPU will run it on many threads according to the number specified
- Square example: `square.cu`
 - compile: `nvcc -o square square.cu`
 - convention:

- variables for host start with `h_`
- variables for device start with `d_`
- launch a kernel: `square<<<blocks, threads per block[, shared memory]>>>`
 - GPU can run many blocks at once
 - maximum number of threads per block
 - older: 512
 - newer: 1024
 - optional: blocks and threads per block can be dimensional: `dim3(w, 1, 1) == dim3(w) == w`
 - optional: stream
 - [how to set arguments](#)

2 GPU Hardware and Parallel Communication Patterns

- communication takes place through memory

Parallel Communication Patterns

- map: tasks read from and write to specific data elements - **one input for one output**
 - square
- gather: each calculation gathers input data from different places together - **many input for one output**
 - pooling
- scatter: tasks compute where to write output - **one input for many output**
 - [map & scatter](#)
- stencil: tasks read input from a fixed neighborhood in an array and much data reuse - **several input for one output**
 - a special kind of gather



- transpose: reorder the data in memory - **one input for one output**
 - a special kind of scatter

GPU Hardware

- [GPU](#)

- multiple SMs: streaming multiprocessor and one block is allocated to one SM, but one SM may deal with several blocks, which depends on the available resources
- a GPU may have multiple SMs
 - a SM has multiple streaming processors, which is responsible for a thread
 - memory
 - local memory: accessible to a specific thread
 - shared memory: accessible to a thread block
 - global memory: accessible to all kernels
 - host/CPU memory
- different thread blocks should not cooperate
- synchronization
 - barrier: point in the program where threads stop and wait. When all threads have reached the barrier, they can proceed
- CUDA makes few guarantees about when and where thread blocks will run
 - pros
 - hardware can run things efficiently
 - no waiting on slowpokes
 - scalability
 - cons
 - no assumptions for block execution sequence
 - no communication between blocks - dead lock
- CUDA guarantees
 - all threads in a block run on the same SM at the same time
 - all blocks in a kernel finish before any blocks from the next kernel run
- How to write efficient program from high level
 - maximize arithmetic intensity $\frac{\text{math}}{\text{memory}}$
 - reduce the time waiting for operands
 - maximize compute ops per thread
 - minimize time spent on memory per thread
 - speed: local memory > shared memory >> global memory >> CPU memory
 - coalesce global memory accessed
 - speed: coalesced > strided > random
 - avoid thread divergence - `if`, `switch`
- atomic operation
 - Problem: multiple threads read/write the same memory location
 - cons
 - only certain operations and data types are supported
 - can work around using atomic comparison and swap
 - still no ordering constraints
 - floating-point arithmetic is non-associative
 - serializes access to memory
 - slow


3 Fundamental GPU Algorithms

Analyze


- step complexity
 - the depth of the workflow/tree
- work complexity
 - total number of nodes of the workflow/tree

Reduce

- inputs
 - set of elements
 - binary & associative reduction operator
- serial implementation
 - each iteration is dependent on the previous iteration



- step complexity: $n - 1 / O(n)$
 - work complexity: $n - 1 / O(n)$
- parallel reduce



- step complexity: $O(\log_2 n)$
 - work complexity: $n - 1 / O(n)$
 - [Brent Theorem](#)
- see `reduce.cu`


Scan

- address set of problems otherwise difficult to parallelize
- not useful in serial world but very useful in parallel world
- inputs
 - set of elements
 - binary & associative scan operator
 - identity element I (for a certain operator) with property $[I \text{ op } a = a]$
- 2 kinds:
 - exclusive scan: all elements before without the current element

- result starts with identity element
 - the last element is ignored
- inclusive scan: all elements before with the current element
 - result starts with the first element
 - the last element is used
- serial implementation
 - step complexity: $n/O(n)$
 - work complexity: $n/O(n)$
- Hillis & Steele's scan
 - inclusive scan
 - more step-efficient
 - step complexity: $\log n/O(\log_2 n)$
 - work complexity: $O(n \log_2 n)$



- Blelloch scan
 - exclusive scan
 - more work-efficient
 - step complexity: $2 \log n/O(\log_2 n)$
 - work complexity: $2n/O(n)$



- which to choose
 - more work than processors: processors limited the performance, so prefer the work-efficient one
 - more processors than work: steps limited the performance, so prefer the step-efficient one

Histogram

- divide the array into several bins for statistics
- cumulative distribution function
 - input: histogram

- output: cumulative distribution function
- operation: exclusive scan
- problems: synchronization
 - atomic operation(addition)
 - slow
 - produce per-thread local histograms and then reduce
 - no need for atomic operations
 - sort the data and then reduce by key

Compact / Filter

COMPACT (MORE FORMALLY)

INPUT	s_0	s_1	s_2	s_3	s_4	...
PREDICATE	T	F	T	F	T	... "is my index even?"
OUTPUT	s_0	—	s_2	—	s_4	← SPARSE

- dense computation (usually with `map()`) is better than sparse computation (usually with `if`), since less threads are required and less idle threads are generated
 - clipping in graphics pipeline
 - matrix multiplication
- compact algorithm
 - predicate
 - scan: produce 1 for TRUE and produce 0 for FALSE
 - exclusive-sum-scan: produce addresses
 - scatter input into output using addresses
- a good strategy for allocation (in step scatter)
 - input: allocation request per input element
 - output: location in array to write the thread's output
 - example of clipping

A GOOD STRATEGY FOR ALLOCATE

- INPUT: ALLOCATION REQUEST PER INPUT ELEMENT
- OUTPUT: LOCATION IN ARRAY TO WRITE YOUR THREAD'S OUTPUT

IN: 1 0 1 2 1 0 3 0
OUT: 0 1 1 2 4 5 5 8

into addresses 5, 6, and 7. And so on. And look. This is exclusive scan as well.

- segmented scan
 - group the large scan into many small scans
 - each is launched independently

新版
反馈

回到
旧版

- combined as segments
- application: multiplication of matrix and vector
 - [CSR](#) (Compressed Sparse Row) representation of sparse matrix
 - value vector
 - column vector
 - rowptr vector (index of column vector)

not all the elements in a row are 0?

- steps:
 - create segmented representation of value vector from rowptr vector
 - gather vector values using column
 - pairwise multiply the above vector
 - (backward) inclusive segmented sum scan

Sort

brick sort & merge sort

- odd-even sort (brick sort)

Unsorted array: 2, 1, 4, 9, 5, 3, 6, 10

Step 1(odd): 2 1 4 9 5 3 6 10

Step 2(even): 1 2 4 9 3 5 6 10

Step 3(odd): 1 2 4 3 9 5 6 10

Step 4(even): 1 2 3 4 5 9 6 10

Step 5(odd): 1 2 3 4 5 6 9 10

Step 6(even): 1 2 3 4 5 6 9 10

Step 7(odd): 1 2 3 4 5 6 9 10


Step 8(even): 1 2 3 4 5 6 9 10

Sorted array: 1, 2, 3, 4, 5, 6, 9, 10

- step complexity: $O(n)$
- work complexity: $O(n^2)$
- [parallel merge sort](#)
 - parallel merge
 - each thread deals with one element
 - use the index in the list and binary search result in the other list to obtain the final index
 - break up the last big merge into smaller ones
 - otherwise many SMs are idle!
 - choose splitters and merge the small list between
 - step complexity: $O(n \log n)$
 - work complexity: $O(n^2)$

sorting networks

- a kind of data-independent sort
- bitonic sequence:
 - only one extreme point & left side is non-strictly monotonically decreasing & right side is non-strictly monotonically increasing
 - sequence that can be rotated to satisfy the above requirement
- [bitonic sort](#)
 - basic operation: `min` + `max` / `swap`
 - one of the fastest and relatively simple
 - steps:
 - split the bitonic sequence into 2 halves
 - cross the two halves and do pairwise comparisons, so that the biggest elements form a new bitonic sequence and the smallest elements form a new bitonic sequence
 - iterate



- particularly efficient when the data can fit into the share memory, so sorting network is often used to sort the data in a block together with other sorting algorithms sorting the blocks
- step complexity: $O(\log^2 n)$
- work complexity: $O(n \log^2 n)$

radix sort

- GPU performance leader
- only for integer
- implemented by scans
- steps
 - start with LSB
 - split the input into 2 sets based on bit. Otherwise preserve order (compact & scatter)
 - move to the next LSB, repeat
 - to reduce the passes: testing multiple bits in a pass
- work complexity: $O(kn)$
 - k : bits in representation

others

- quick sort
 - one of the fastest in serial world
 - steps
 - choose pivot element
 - compare all elements vs. pivot
 - split into 3 arrays: <P, =P and >P
 - recurse on each array
 - implement
 - comparison & split: compact
 - recurse: segment
 - there is no recursion in GPU
 - hard to implement
- key value sort
 - sort the pointer rather than the complex data structure
 - move the key and the value together
 - most GPUs handle 64-bit data, it's good to use a 32-bit key and a 32-bit value

4 Optimizing GPU Algorithms

Levels of optimization:

1. picking good algorithms
 - 3-10x speedup
2. basic principles for efficiency
 - 3-10x speedup
3. architecture-specific detailed optimization
 - 30%-80% speedup
4. μ -optimization of instruction level

APOD - systematic optimization

- analyze
- parallelize
- optimize
- deploy

analyze

- understanding hotspots
 - do not rely on intuition
 - run a profiler to know which to parallelize
 - *Amdahl's law* to know the speedup limit: $S = \frac{1}{(1-P) + \frac{P}{n}}$
- understanding strong scaling vs. weak scaling

- weak scaling: run a bigger problem or more small problems in parallel
- strong scaling: run a problem faster
- time limits
 - memory access
 - use `CUDA Device Query` to check the information of the device
 - 40-60 usage: okay
 - 60-75 usage: pretty well
 - >75 usage: excellent
 - `GB/s` is usually used
 - compute performance

parallelize

see `transpose.cu`

- maximum parallelism does not always bring the best performance -> granularity coarsening
- v1: thread per matrix
- v2: thread per row
- v3: thread per element
- v4: thread per element with tiling
 - consider memory access limit and set `K=32`
- v5: thread per element with tiling and `K=16`
 - consider occupancy
- v6: consider shared memory bank conflicts & padding

optimize

see `transpose.cu`

For code:

- most GPU codes are memory limited
 - tools to visualize performance: nSight / NVVP
 - parallel per element version only achieves 31.1% peak memory bandwidth, so it needs to make use of coalescing
 - writing `out` in `out[j + i*N] = in[i + j*N];` is bad
 - method: each tile is transposed in shared memory to achieve coalescing writes
 - *Little's Law*: the number of bytes delivered = the average latency of each memory transaction × the bandwidth
 - increase the number of bytes delivered or decrease the latency to increase the bandwidth
 - the synchronization limits the latency, so reduce `K` to 16 helps to increase efficiency
 - occupancy: usage of the resources on a SM
 - increasing occupancy usually helps to a point
 - methods:
 - control amount of shared memory can affect occupancy

- change the blocks and threads in `<<< , >>>`
 - compilation options to control register usage
 - further optimization: [shared memory bank conflicts](#) & padding
- optimize compute performance
 - minimize time waiting at barriers
 - minimize thread divergence
 - [warp](#): set of threads (32 threads) that execute the same instruction at the same time

warp is the basic unit of scheduling and operation and not all the threads are executed at the same time

- CPU - SIMD: SSE/AVX
- GPU - [SIMT](#)
 - be aware of branch divergence
 - avoid branchy code
 - beware large imbalance in thread workloads
 - switch

CUDA assigns thread IDs to warps:

x varies fastest
y varies slower
z varies slowest

x →

0	1	2	3	4	...	30	31	32	33	...	61	62	63
0	1	2	3	4	...	30	31	32	33	...	61	62	63
:	:	:	:	:	:	:	:	:	:	:	:	:	:
0	1	2	3	4	...	30	31	32	33	...	61	62	63

16

Quiz: what will be the slowdown for each of the following expressions in switch statement?

32 switch (threadIdx.x % 32) { case 0..31 } ...
kernel <<< 1024, 1 >>> ();

32 switch (threadIdx.x % 64) { case 0..63 } ...
kernel <<< 1024, 1 >>> ();

1 switch (threadIdx.y) { case 0..31 } ...
kernel <<< 64*16, 1 >>> ();

2 switch (threadIdx.y) { case 0..31 } ...
kernel <<< 16*16, 1 >>> ();

the threads in a warp should belong to the same block?

- loops
- assorted math optimizations
 - use double precision only when you mean it
 - use intrinsics

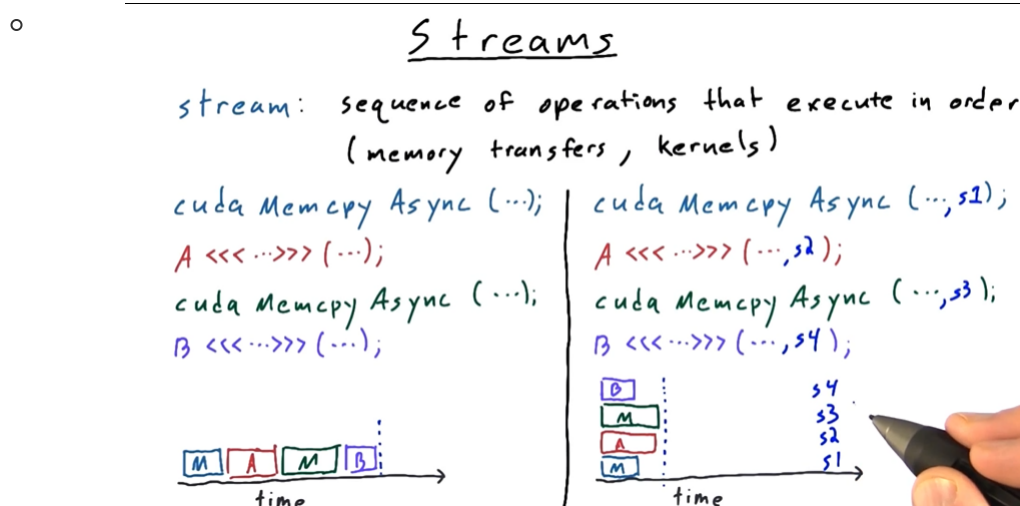
For system:

- host and device communicate through PCI/PCIe

- use `cudaHostMalloc()` or `cudaHostRegister()` to avoid copy to staging area of PCI (PCIe can only transfer data in staging area)
- enables asynchronization `cudaMemcpyAsync()`
- streams: sequence of operations that execute in order
 - each operation/instruction can be assigned a stream

```
cudaStream_t s1;
cudaStreamCreate(&s1);
cudaStreamDestroy(s1);
```

- one fashion to use:
 - break the data that fits the GPU memory into 2 halves
 - transfer one half and begin the next transfer while beginning processing the ready half
- overlap memory transfers & compute
- help to fill GPU with smaller kernels



5 Parallel Computing Patterns

N-body Computation

Problems:

- all-pairs N-body
 - simplest
 - brute-force: $O(N^2)$
 - tree methods (*Barnes-Hut*): $O(N \log N)$
 - fast multipole method: $O(N)$
- target: to obtain the force on each object, and move each object accordingly to simulate

Solutions:

see `nbody.cu`

- NxN matrix

- x-axis: source object
 - y-axis: destination object
- better: tiling
 - PxP tiles
 - calculate each tile and sum up the horizontal tiles
 - one thread for one destination object in a tile
 - still must share source object parameters across all threads
 - no sharing necessity for destination objects
 - no communication necessity to sum up results
 - reduce parallelism actually
 - to obtain best efficiency, the size of the tile should be chosen carefully
 - **why to make tradeoff between more parallelism and more work per thread?**
 - the communication cost with a thread is much smaller than the cost between threads
 - more parallelism may bring more communication and more synchronization

Sparse Matrix-Vector Multiply

Problems:


- [multiply a sparse matrix by a dense vector](#)
 - to use less memory
 - to use less computation
- what is the role of a thread?
 - thread per row
 - calculate one entry of the result
 - more work per thread
 - thread per element
 - calculate one partial product of one entry
 - more threads
 - communication between threads
 - **again to make tradeoff between more parallelism and more work per thread**

Solutions:

see `spmv.cu`

- thread per row
 - the lengths of rows may be different and the run time depends on the longest row
 - better when the lengths of rows are similar
- thread per element
 - run time depends only on the number of non-zero elements
 - better when the lengths are varying
- hybrid approach
 -

HYBRID APPROACH [DELL & GARLAND, SUPERCOMPUTING 2009]



- how to draw the line dividing the matrix?

- keep threads busy
- closer communication is cheaper

BFS

Problems:


- implement BFS on graphs in parallel
- calculating distance from root for each node

Solutions:

see [BFS.cu](#)

- simple version with no frontier
 - steps:
 - in each iteration, visit all the edges
 - if a vertex is found unvisited, set its depth
 - iterate V times at most
 - what if a vertex is set multiple times in one iteration: it's OK since the value to be set is the same
 - when to terminate: when no vertex is set in one iteration
 - work complexity: $O(VE)$
 - consider a linear chain
 - usually in a graph edges are far more than vertices, so the work complexity is at least $O(V^2)$
 - weak scalability
- with a frontier
 -

GRAPH DATA STRUCTURE



- C storing the neighbors, whose length is $O(E)$
- R storing the start offset of neighbors of each vertex, whose length is $O(V)$
- D storing the depth

- steps:

- in parallel, for each vertex on the frontier, find the starting point of its neighbors
- for each vertex on the frontier, calculate the number of its neighbors by $R[v+1] - R[v]$
- allocate space to store the new frontier
- copy each active edge list to the allocated array and the storing location is calculated by an exclusive scan
- call the list, removing/compact items that have been visited, which can be determined by D

List Ranking

In a link list, how to set `chum` (a pointer pointing to the last node) ?

in each iteration, use `chum[k] = chum[chum[k]]` instead of `chum[k] = next[k] + 1`

Problems:

- output: all nodes in order
- input
 - each node knowing its successor
 - starting node
 - organized as a link list

Solutions:

- construct a table recording the node k -hop away from each node
 - work complexity: $O(N \log N)$ (more than serial algorithm)
 - step complexity: $O(\log N)$ (less than serial algorithm)
- steps:
 - construct the table
 - scatter and recording output positions
 - in n -th iteration, check the 2^n -hop-away row to activate threads from currently active threads

INPUT: STRUCT

NODE #	0	1	2	3	4	5	6	7	8	9	ANSWER
+1	5	6	7	8	9	2	3	4	0	1	0 5 2 7 4
+2	2	3	4	0	1	7	8	9	5	6	9 1 6 3 8
+4	4	0	1	2	3	9	5	6	7	8	
+8	3	4	0	1	2	8	9	5	6	7	
OUTPUT	0		2			1		3			

get position number 3. On the next iteration, we're going to use the next line

Hash

Problems:

- chaining is bad for parallelism

- load imbalance: the lengths may be varying, so some threads may be waiting for a long time
- contention in construction
- implement parallel hashing

Solutions:

- [cuckoo hashing](#)
 - multiple hash tables and one hash function per table
 - steps:
 - all items use h_1 and write into t_1 ; kick out the existed items if necessary
 - for those failed ones, use h_2
 - repeat
 - if there are too many iterations, choose new hash functions and start over
 - implementation
 - constant time lookup
 - construction is simple & fast
 - kick out should be atomic
 - alternative algorithm
 - construction: sort
 - lookup: binary search

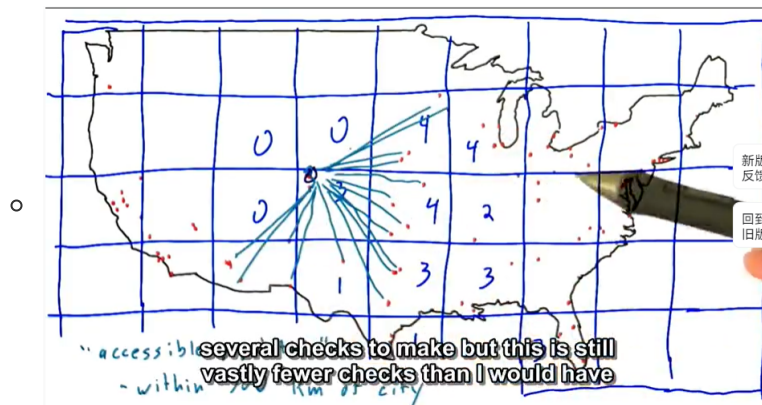
6 Additional Parallel Computing Topics

- [reference](#)

Parallel Optimization Patterns

7 basic techniques from Stratton's taxonomy:

1. data layout transformation
 - burst utilization like AoS -> SoA
 - [partition camping](#) (bank conflict?)
 - arrays of structures of tiled arrays (ASTA)
2. scatter-to-gather transformation
 - gather: many overlapping reads
 - scatter: many potentially conflicting writes
3. tiling
 - buffering data into fast on-chip storage for repeated access/reuses
4. privatization
 - avoid different threads/blocks writing to the same output location
5. binning/spatial data structure
 - binning: build data structure that maps output locations to (a small subset of) the relevant input data
 - group the input into bins, and only consider relevant bins when calculating



6. compaction

- when most input elements require no computation
- consider the structure of warps, the speedup is usually at most 32x

7. regularization

- reorganizing input data to reduce load imbalance

Libraries

- **cuBLAS**
 - cuda Basic Linear Algebra Subroutines
- **cuFFT**
 - 1D, 2D, 3D
 - cuda Fast Fourier Transform
- **cuSPARSE**
 - BLAS-like routines for sparse matrix formats
- **cuRAND**
 - pseudo-random and quasi-random generation routines
- **NPP**
 - low-level image processing primitives
- **Magma**
 - GPU and multiple CPU implementations of LAPACK routines
- **CULA**
 - implementations of Eigensolvers, matrix factorization, and matrix solvers for dense matrices similar to the LAPACK API
- **ArrayFire**
 - framework for data-parallel array manipulation

How to use:

- substitute library calls
- manage data locality
- rebuild and link

CUDA C++ Programming Power Tools

- **Thrust**
 - analogous to STL

- automatically installed when installing CUDA toolkits
- [CSDN](#) and [documents](#)
- **CUB**
 - CUDA Unbound
 - accessing memory can be complicated and this tool helps to manage global memory, shared memory and local variables of threads
- **CUDA DMA**
 - make it easier to use shared memory
 - specifically focus on shared memory and global memory
 - high performance
 - explicit transfer patterns: sequential, strided and indirect

Platforms

Other languages:

- wrap CUDA C++
 - PyCUDA
- target CUDA directly
 - copperhead - data-parallel subset of Python
 - CUDA Fortran
 - Halide - image processing DSL
 - Matlab

see `whitebalance.m`

Cross-platform:

- OpenCL
- OpenGL compute
- OpenACC

Dynamic Parallelism

see `dynamic.cu`

- bulk parallelism
 - every thread is executing in parallel without any dependency
- nested parallelism
 - a kernel launches another kernel
 - watch out
 - multiple launches caused by every thread executing the same launch call
 - each block is independent
 - a block's private/shared data is private and cannot be inherited
 - recursive parallelisms
- task parallelism
 - use multiple tasks to fill the whole GPU

- multiple tasks operate their own data

composability: every kernel is independent and does not care about what other kernels are doing