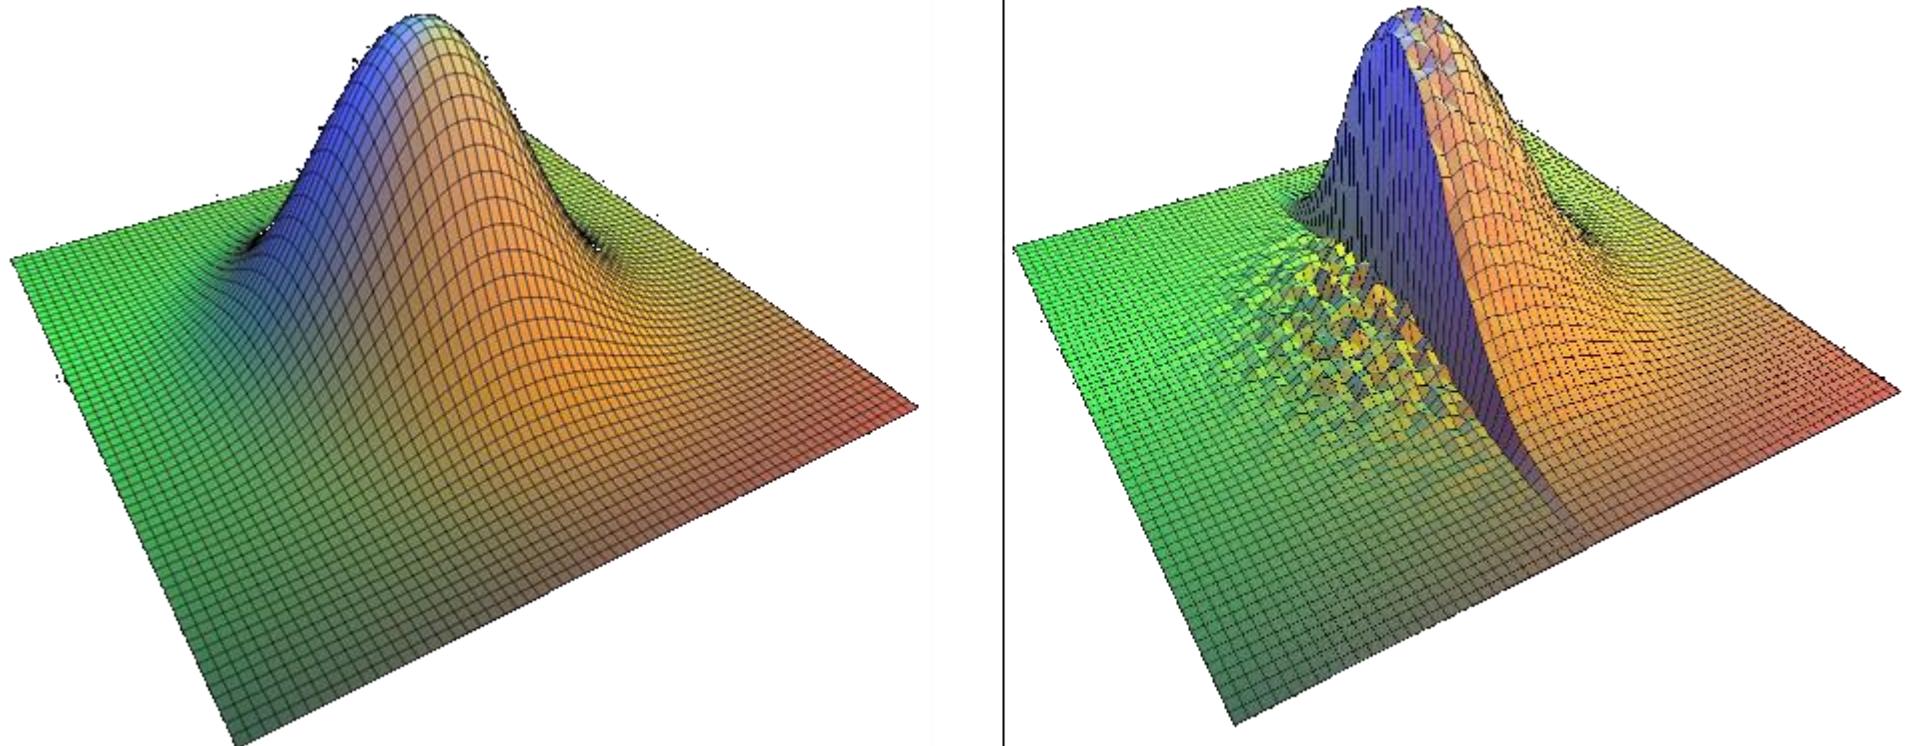


# 5. Filters



# Image Filtering

- An image is a 2D array of pixel values

- Filtering:

- Replace each pixel by a *linear* combination of its neighbors
- The combination is determined by the filter's *kernel*
- Often spatially-invariant, the same kernel is applied to all pixel locations

$$g[\cdot, \cdot] = \frac{1}{16}$$

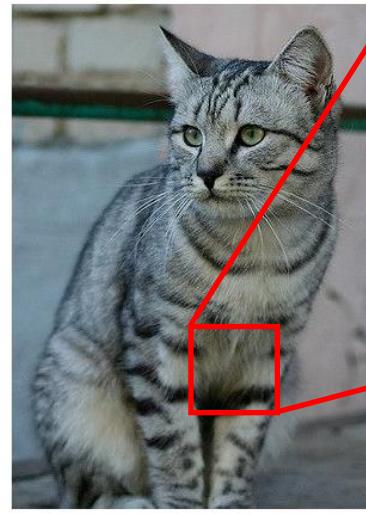
1	2	1
2	4	2
1	2	1

Gaussian Filter

$$g[\cdot, \cdot] = \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Box Filter



105	112	108	111	104	99	106	99	96	103	112	119	104	97	93	87
91	98	102	106	104	79	98	103	99	105	123	136	118	105	94	85
76	85	90	102	98	105	87	96	95	99	123	136	118	105	99	85
89	95	81	93	123	131	127	100	95	104	107	96	96	93	104	94
106	91	64	69	95	88	85	104	107	109	96	75	94	95	96	95
114	108	85	55	55	59	44	64	87	112	120	98	74	84	91	91
133	137	147	103	65	81	88	65	52	84	74	86	103	99	85	82
128	137	144	140	109	95	86	70	62	65	63	63	68	73	85	101
125	133	148	137	119	121	117	94	65	79	88	65	54	64	72	98
137	125	131	147	133	127	126	131	111	98	89	75	61	64	72	84
115	114	109	123	158	148	131	118	113	109	100	92	74	65	72	76
109	93	98	97	108	147	131	118	132	114	113	109	105	95	77	98
163	77	86	81	77	79	102	123	117	115	117	125	125	130	115	87
162	55	82	89	78	71	88	181	124	126	119	181	187	114	131	119
163	65	75	88	83	71	62	81	128	138	135	185	81	98	118	118
187	65	71	87	106	95	69	45	76	130	126	187	92	94	105	112
118	97	82	88	117	123	116	66	42	51	95	93	89	95	102	107
164	146	112	88	82	120	124	184	76	48	45	66	88	181	102	109
157	170	157	120	93	86	114	132	112	97	69	55	70	92	99	94
130	128	134	161	139	109	119	118	121	134	114	87	65	53	69	86
128	112	98	117	158	144	128	115	104	187	102	93	87	81	72	79
123	107	95	86	83	112	153	149	122	189	104	75	88	187	112	99
122	121	102	88	82	86	94	117	145	148	153	102	58	78	92	107
122	164	148	103	71	56	78	83	93	103	119	139	102	61	69	84

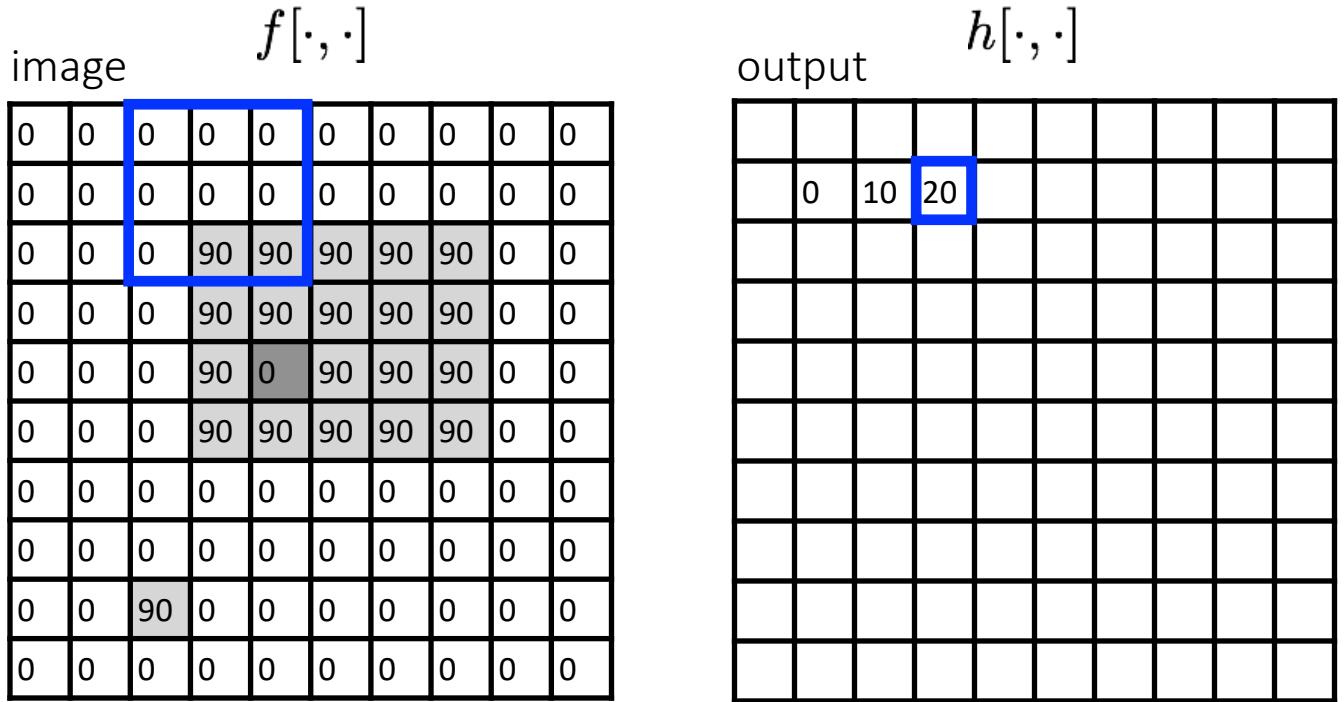


# Box filter example

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

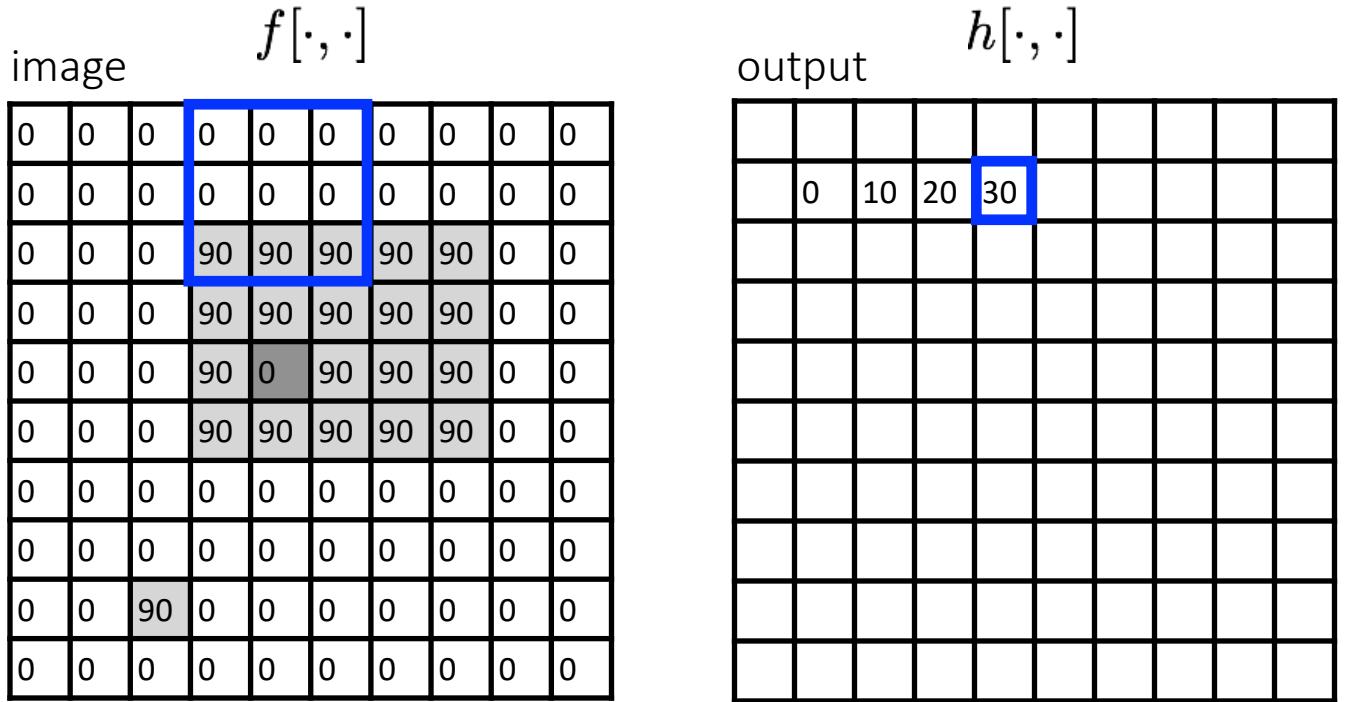
output       $k, l$       filter      image (signal)

# Box filter example

$$g[\cdot, \cdot]$$

kernel

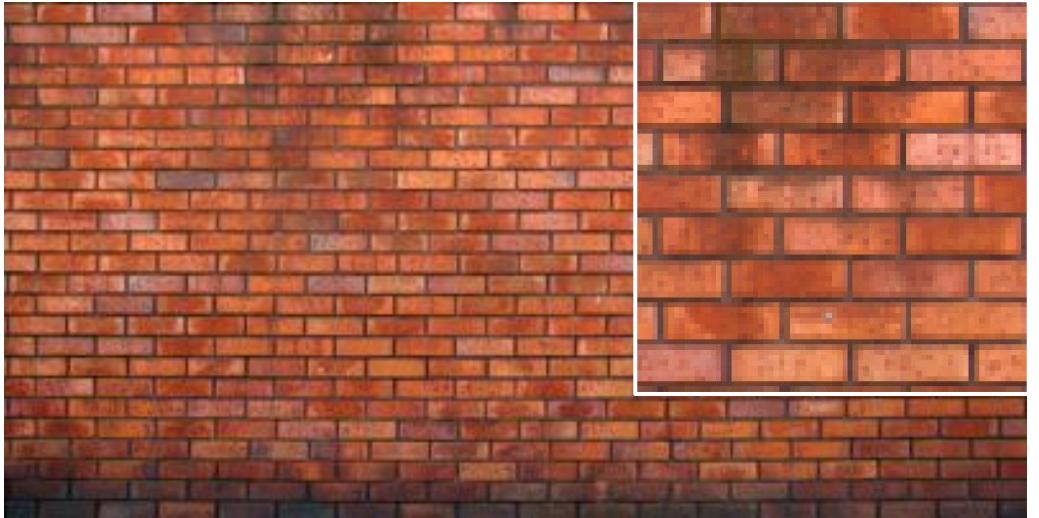
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

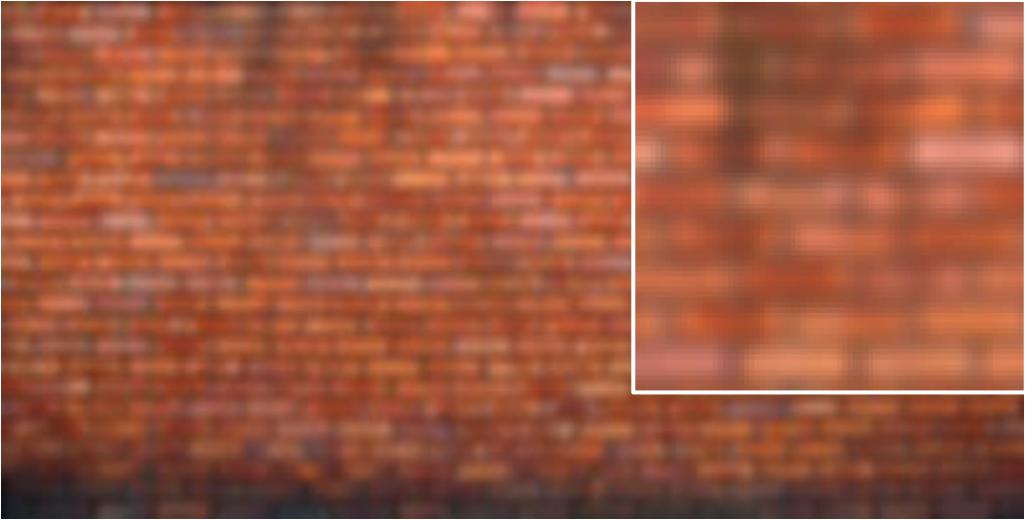
output      filter      image (signal)

# Gaussian vs box filtering

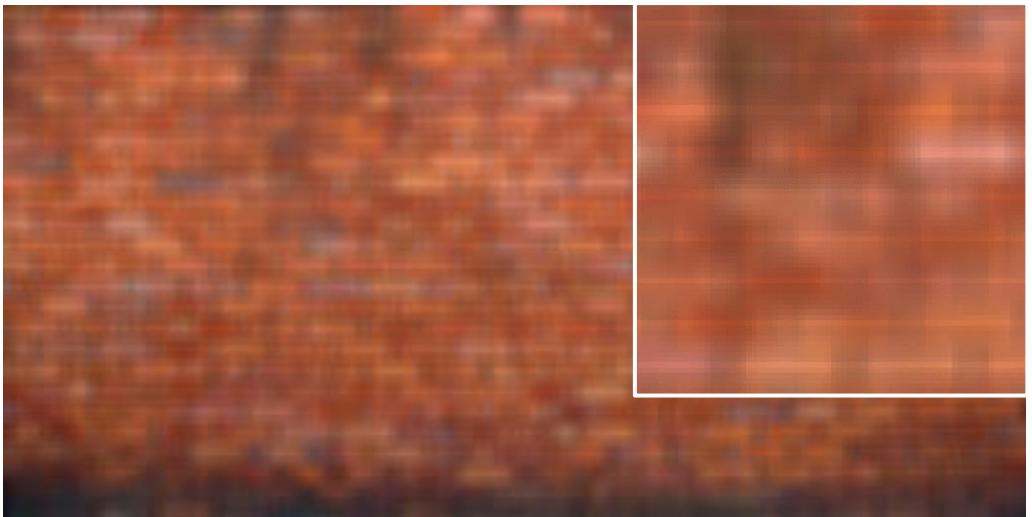


original

Which blur do you like better?



7x7 Gaussian



7x7 box



# How to create a soft shadow effect?

CMU

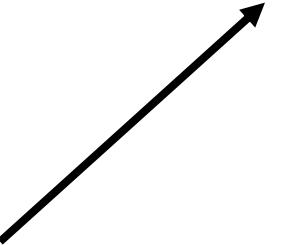
CMU

Gaussian blur



CMU

shift + overlay





# Fun with filters

input



filter

0	0	0
0	1	0
0	0	0

output



unchanged

input



filter

0	0	0
0	0	1
0	0	0

output



shift to left  
by one



# Detecting Edges

- How would you detect edges in an image (i.e., discontinuities in a function)?
  - Take derivatives: derivatives are large at discontinuities

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$

- How do you differentiate a discrete image (or any other discrete signal)?
  - Use finite differences

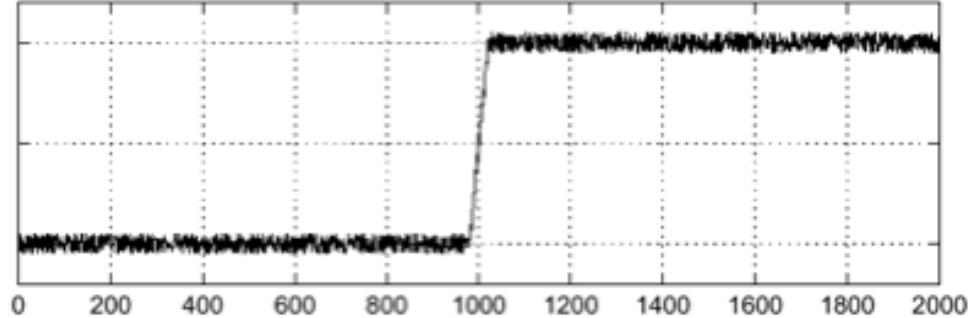
-1	0	1
----	---	---

Remove limit and set  $h = 2$

1	0	-1
---	---	----

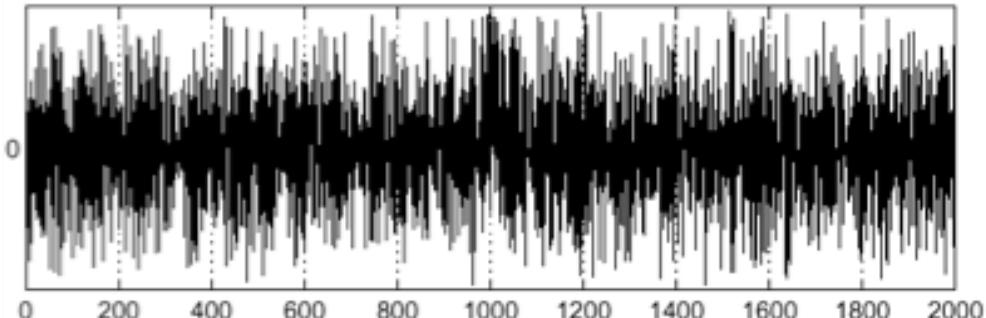
# But Images are Noisy

intensity plot



Using a derivative filter:

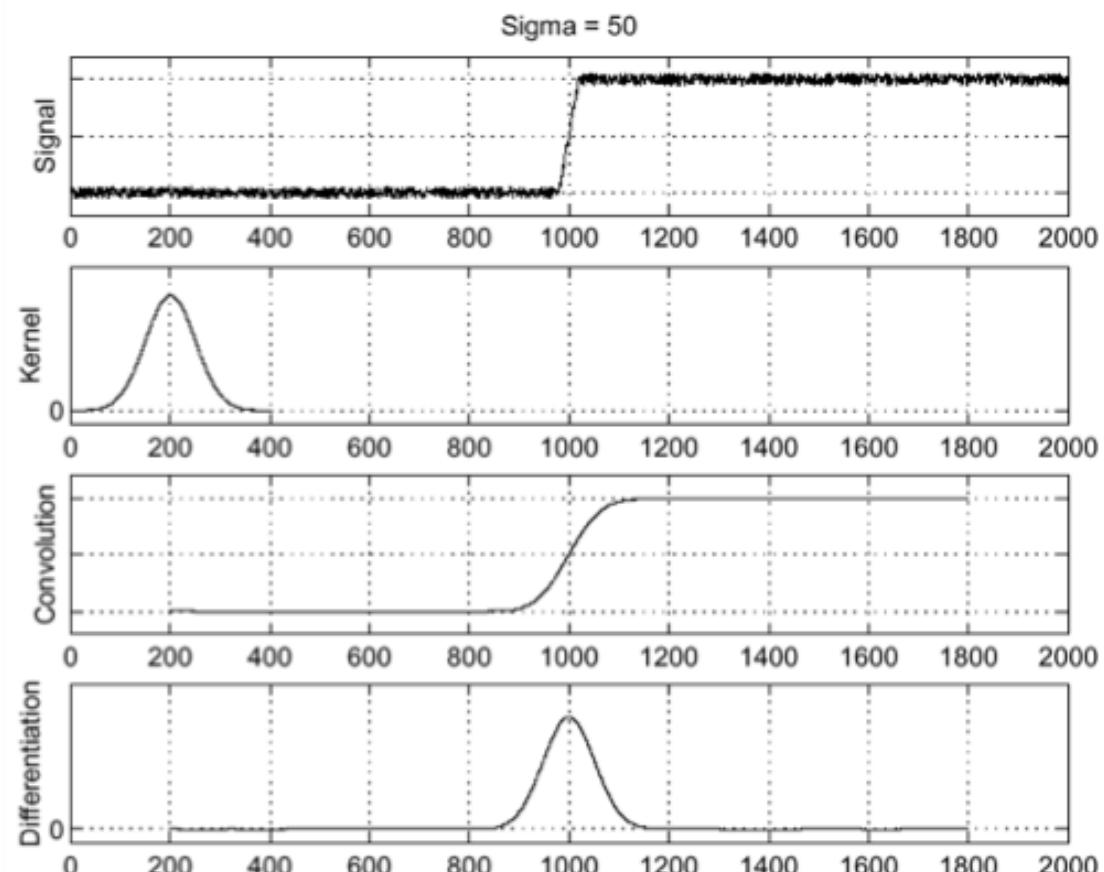
derivative plot



# Blur Before Taking Derivatives

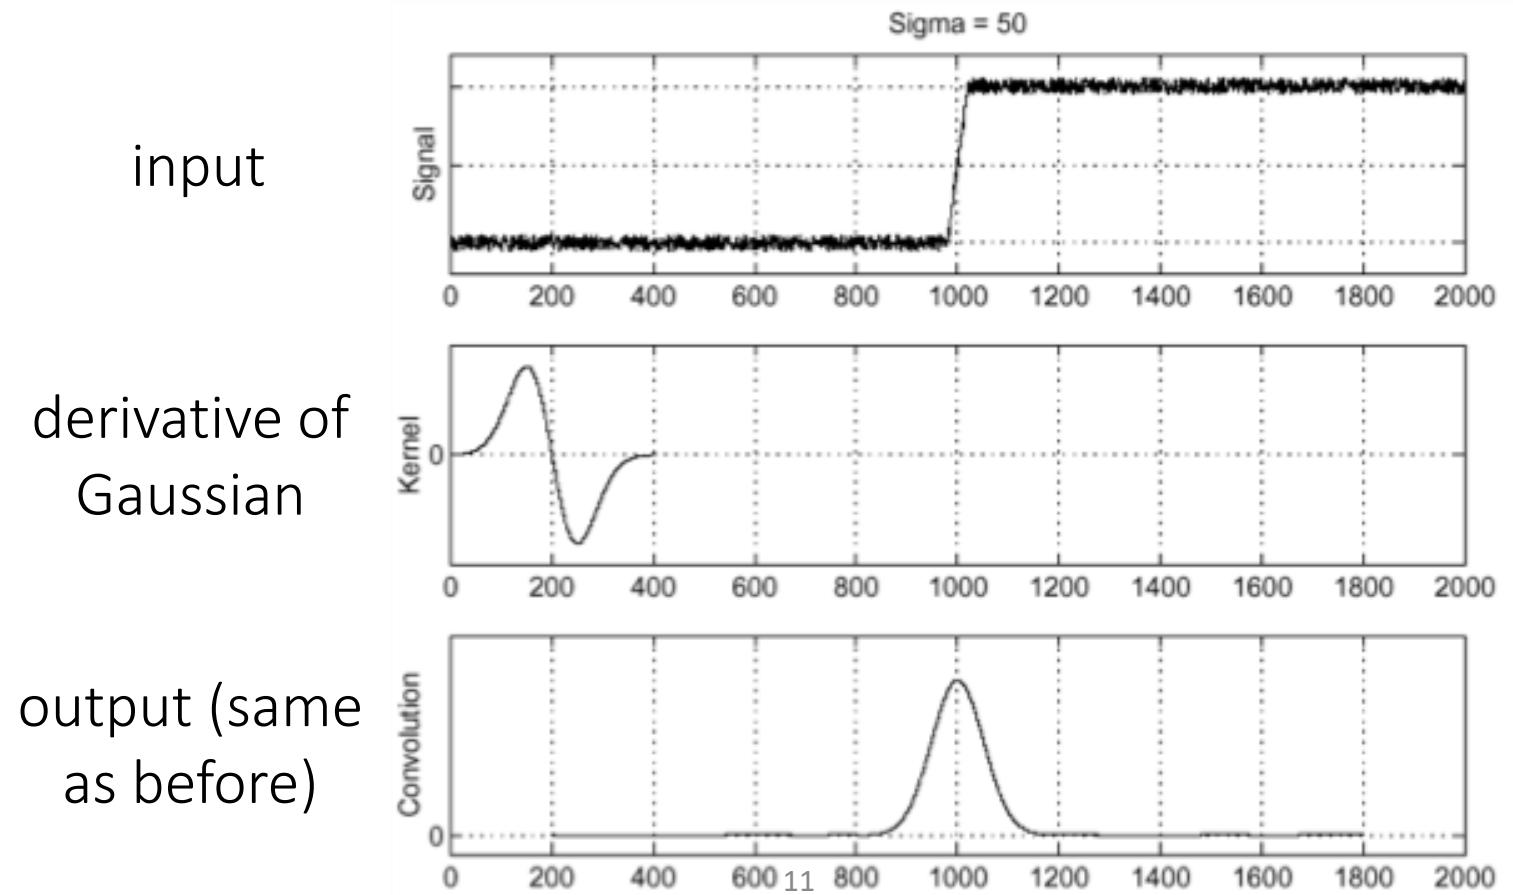
- When using derivative filters, it is critical to blur first!

input  
Gaussian  
blurred  
derivative of  
blurred

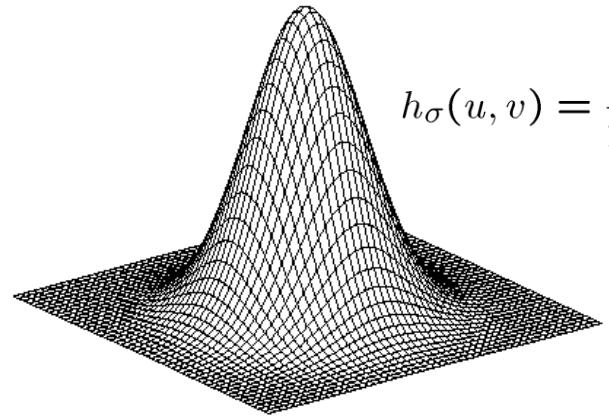


# Derivative of Gaussian (DoG) filter

- Derivative theorem of convolution:  $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$
- Applying DoG = Applying blur first and then taking derivative

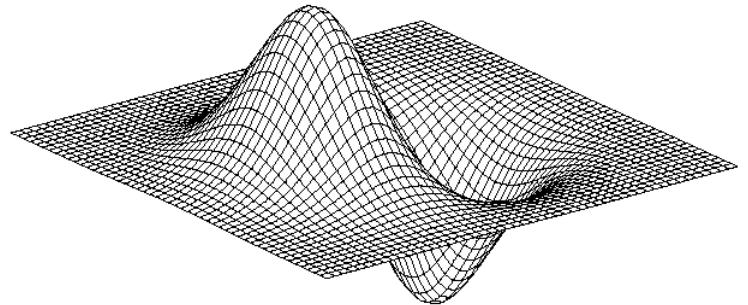


# 2D Gaussian filters



$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

Gaussian



Derivative of Gaussian



<sub>12</sub> Derivative of Gaussian filtering



# The Sobel filter

Horizontal Sobel filter:

1	0	-1
2	0	-2
1	0	-1

Sobel filter

$$= \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

Blurring

1	0	-1
---	---	----

1D derivative  
filter

Vertical Sobel filter:

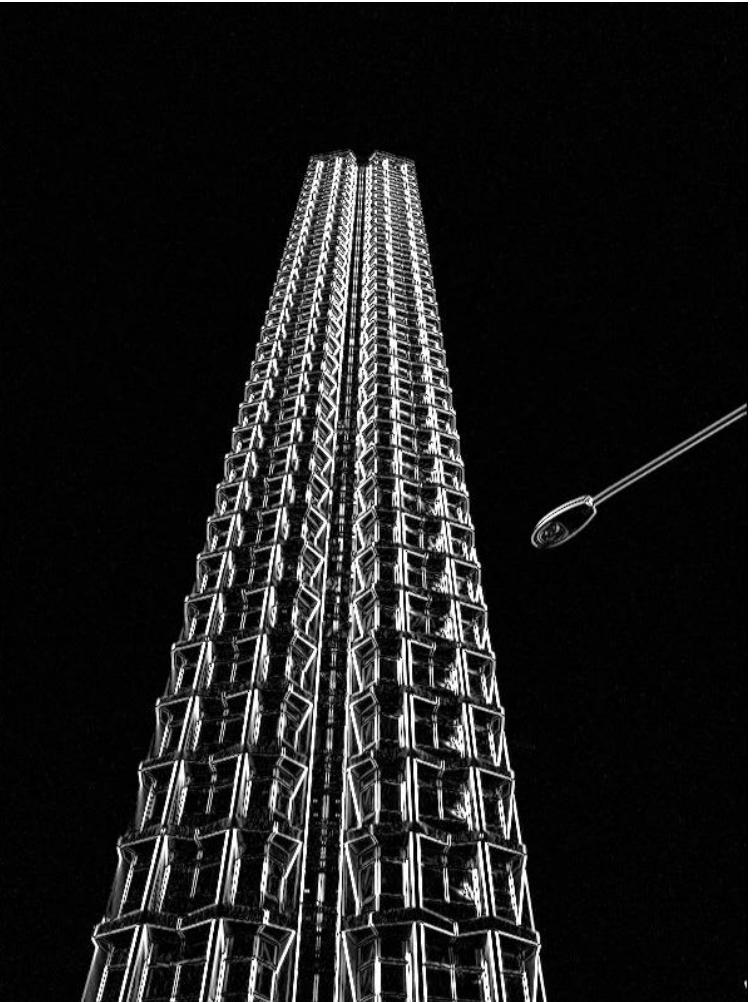
1	2	1
0	0	0
-1	-2	-1

$$= \begin{array}{|c|c|} \hline 1 \\ \hline 0 \\ \hline -1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

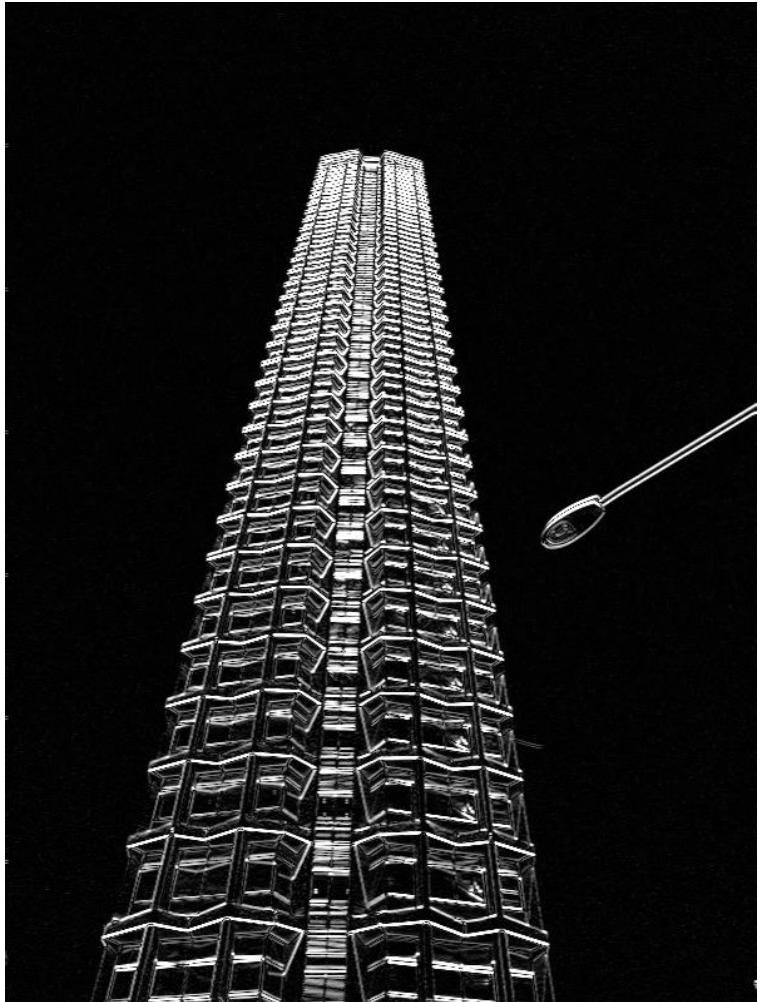
# Sobel filter example



original

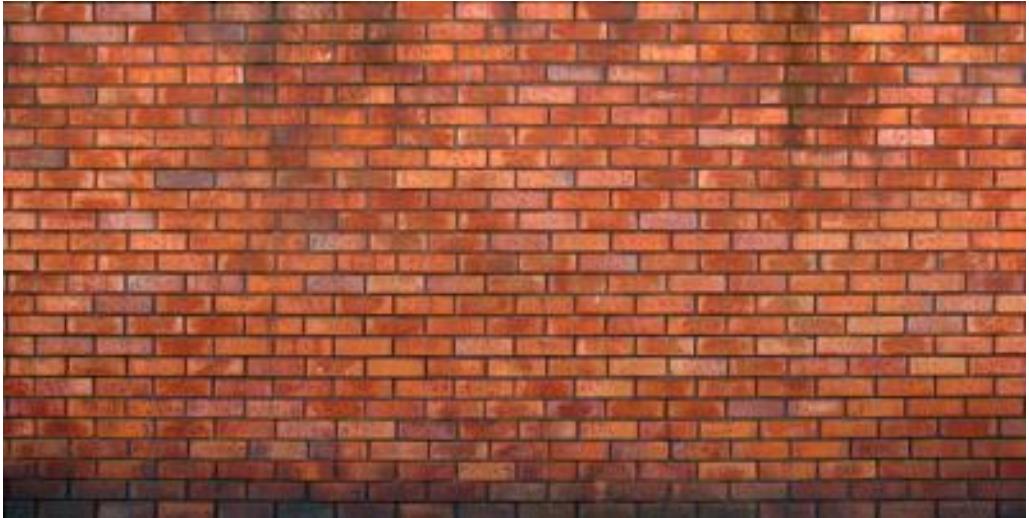


horizontal Sobel filter



vertical Sobel filter

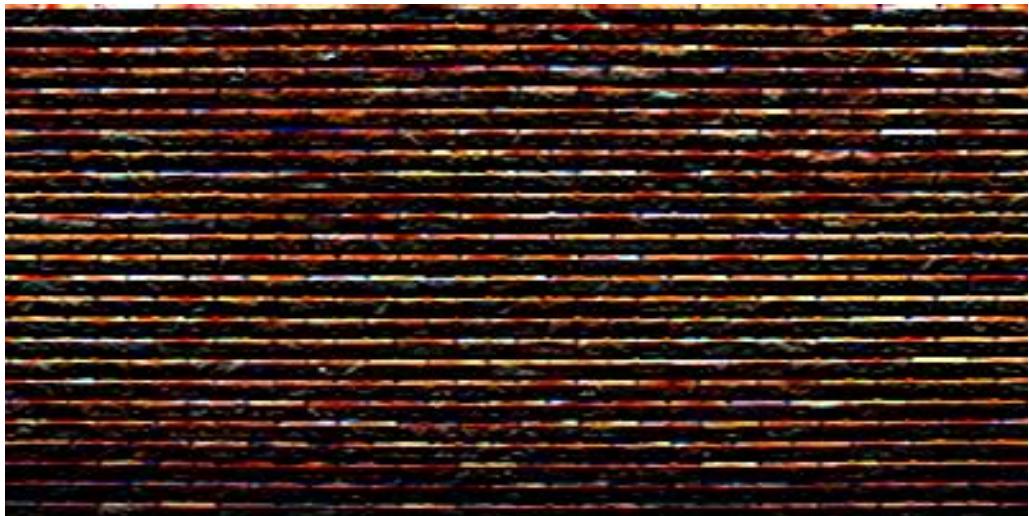
# Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter



# Computing Image Gradients

1. Select a derivative filters (there are other similar filters, e.g. Scharr)

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial f}{\partial x} = \mathbf{S}_x \otimes f$$

$$\frac{\partial f}{\partial y} = \mathbf{S}_y \otimes f$$

3. Form the image gradient, and compute its direction and amplitude.

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

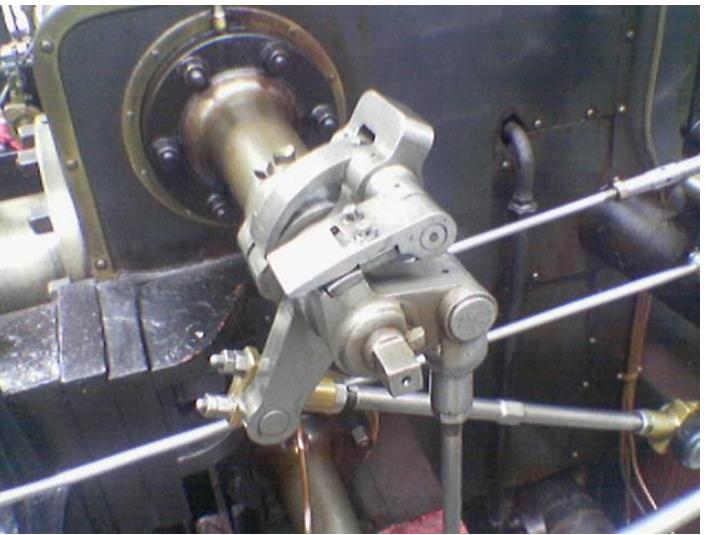
direction

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

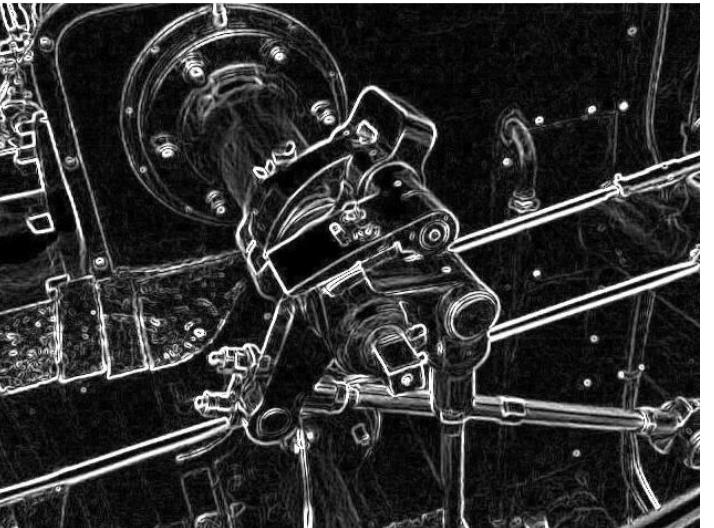
amplitude

# Image gradient example

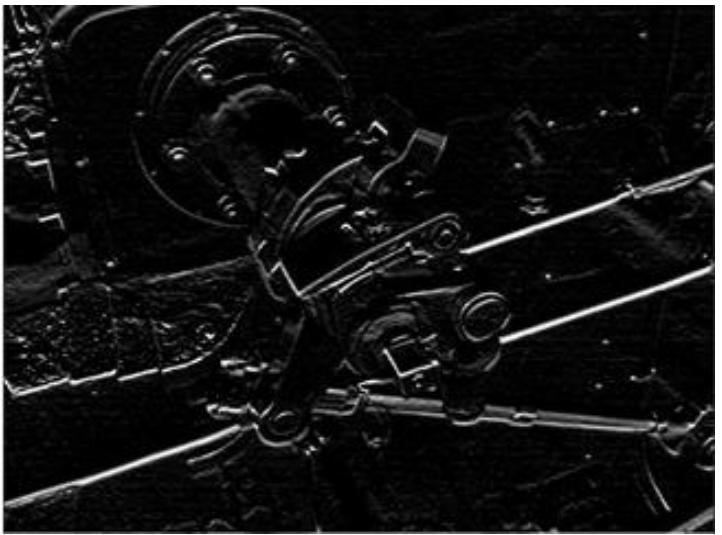
original



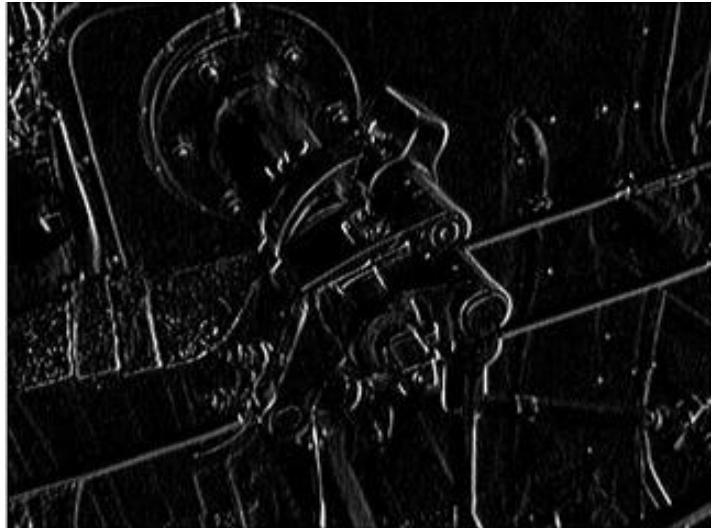
gradient  
amplitude



vertical derivative



horizontal  
derivative



# Questions?



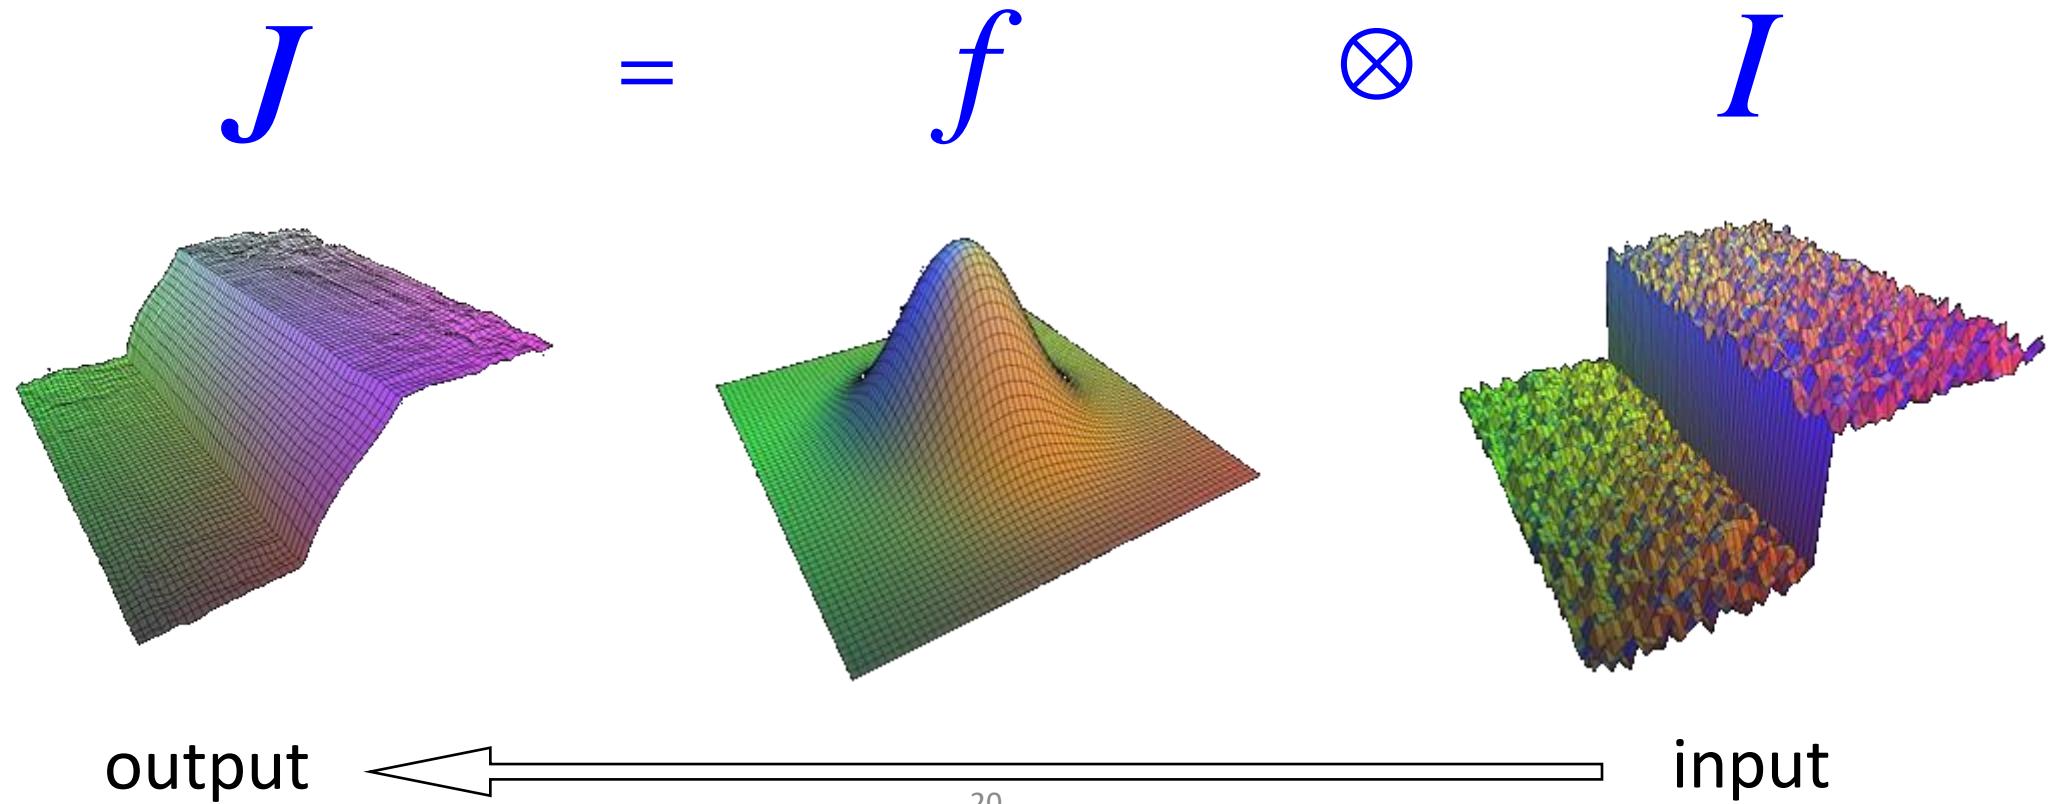


# Bilateral filter

- Proposed by Tomasi and Manduci in ICCV 1998
  - “Bilateral Filtering for Gray and Color Images”
- A very good survey by Sylvain Paris et al. 2009
  - Published at Foundations and Trends in Computer Graphics and Vision
  - “Bilateral Filtering: Theory and Applications”
- Related to
  - SUSAN filter  
[Smith and Brady 95] <http://citeseer.ist.psu.edu/smith95susan.html>
  - Digital-TV [Chan, Osher and Chen 2001]  
<http://citeseer.ist.psu.edu/ch01digital.html>
  - sigma filter <http://www.geogr.ku.dk/CHIPS/Manual/f187.htm>

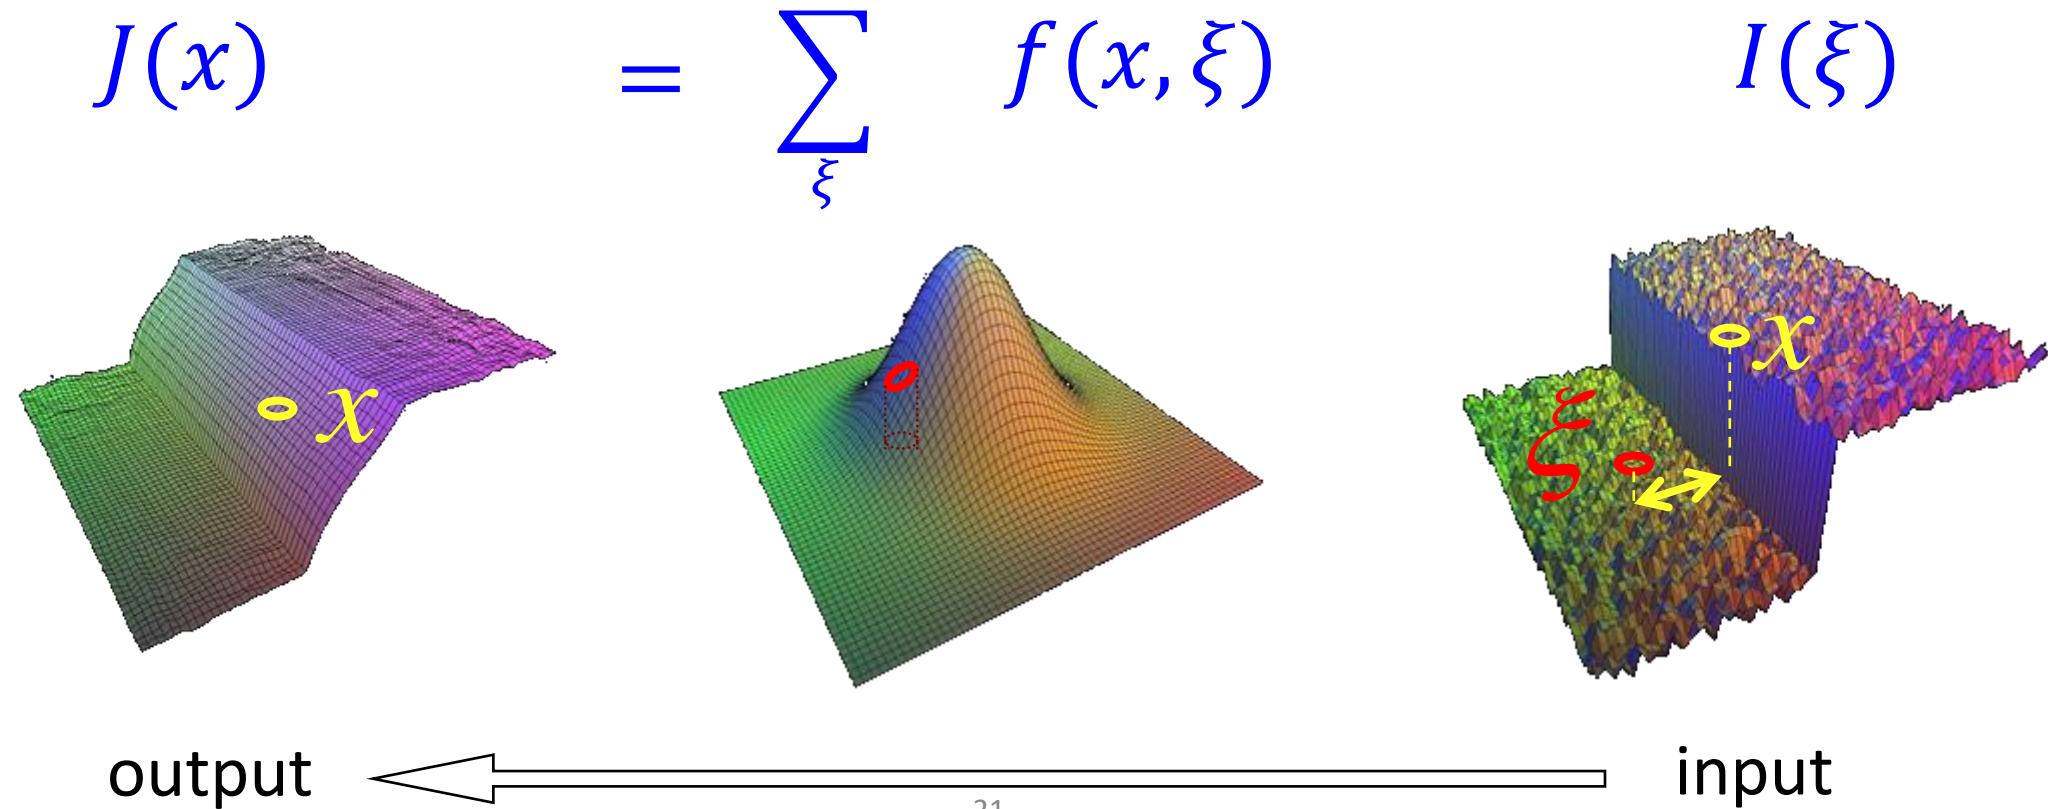
# Start with Gaussian Filter

- Here, input is a step function + noise
- Spatial Gaussian filter  $f$
- Output is blurred



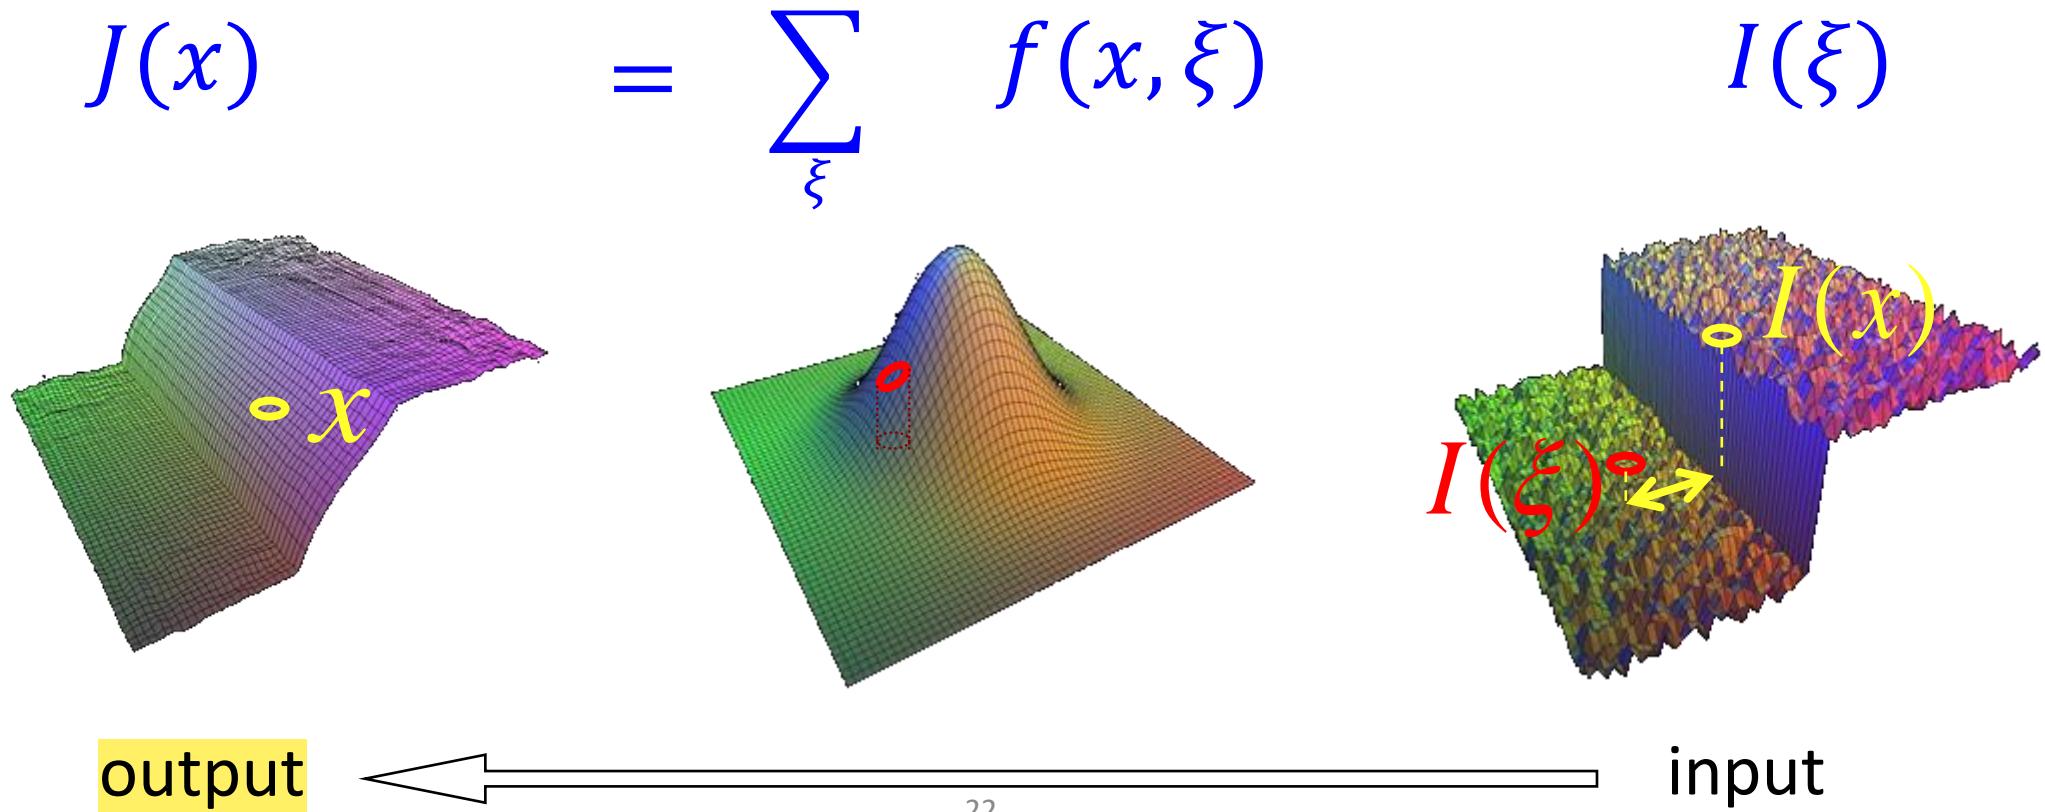
# Gaussian Filter as Weighted Average

- Weight of  $\xi$  depends on its distance to  $x$



# The Problem of Edges

- Here,  $I(\xi)$  “pollutes” our estimated  $J(x)$
- It is too different to be averaged together

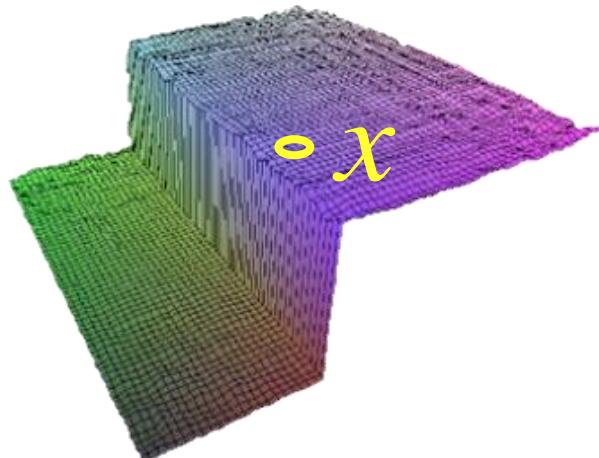


# Principle of Bilateral Filter

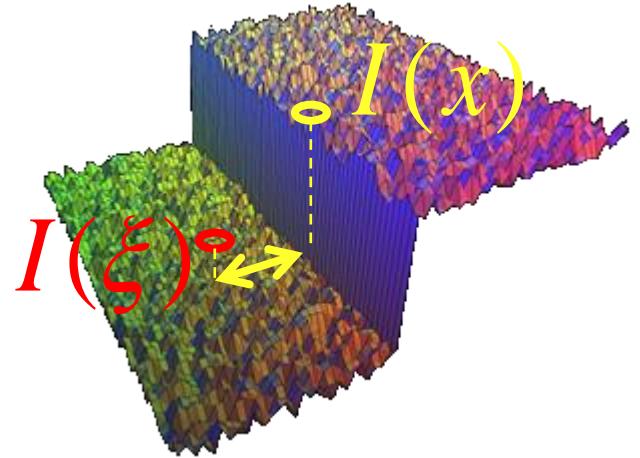
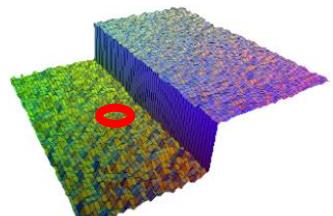
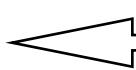
[Tomasi and Manduchi 1998]

- Penalty Gaussian  $g$  on the intensity difference

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$



output



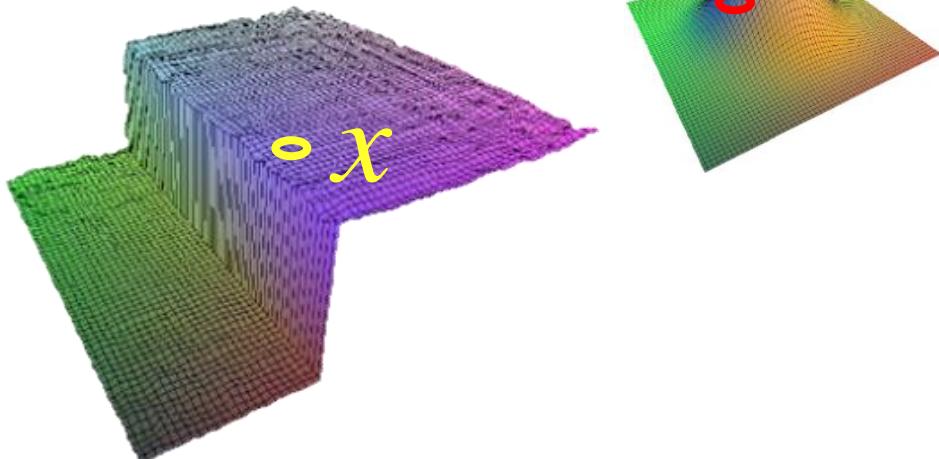
input

# Bilateral Filter

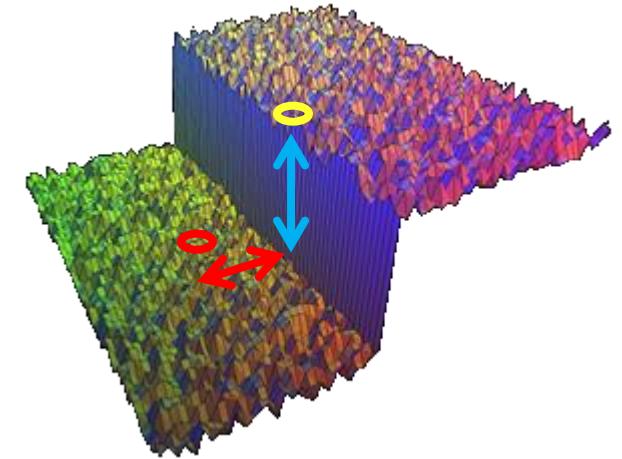
[Tomasi and Manduchi 1998]

- Spatial Gaussian  $f$

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$



output



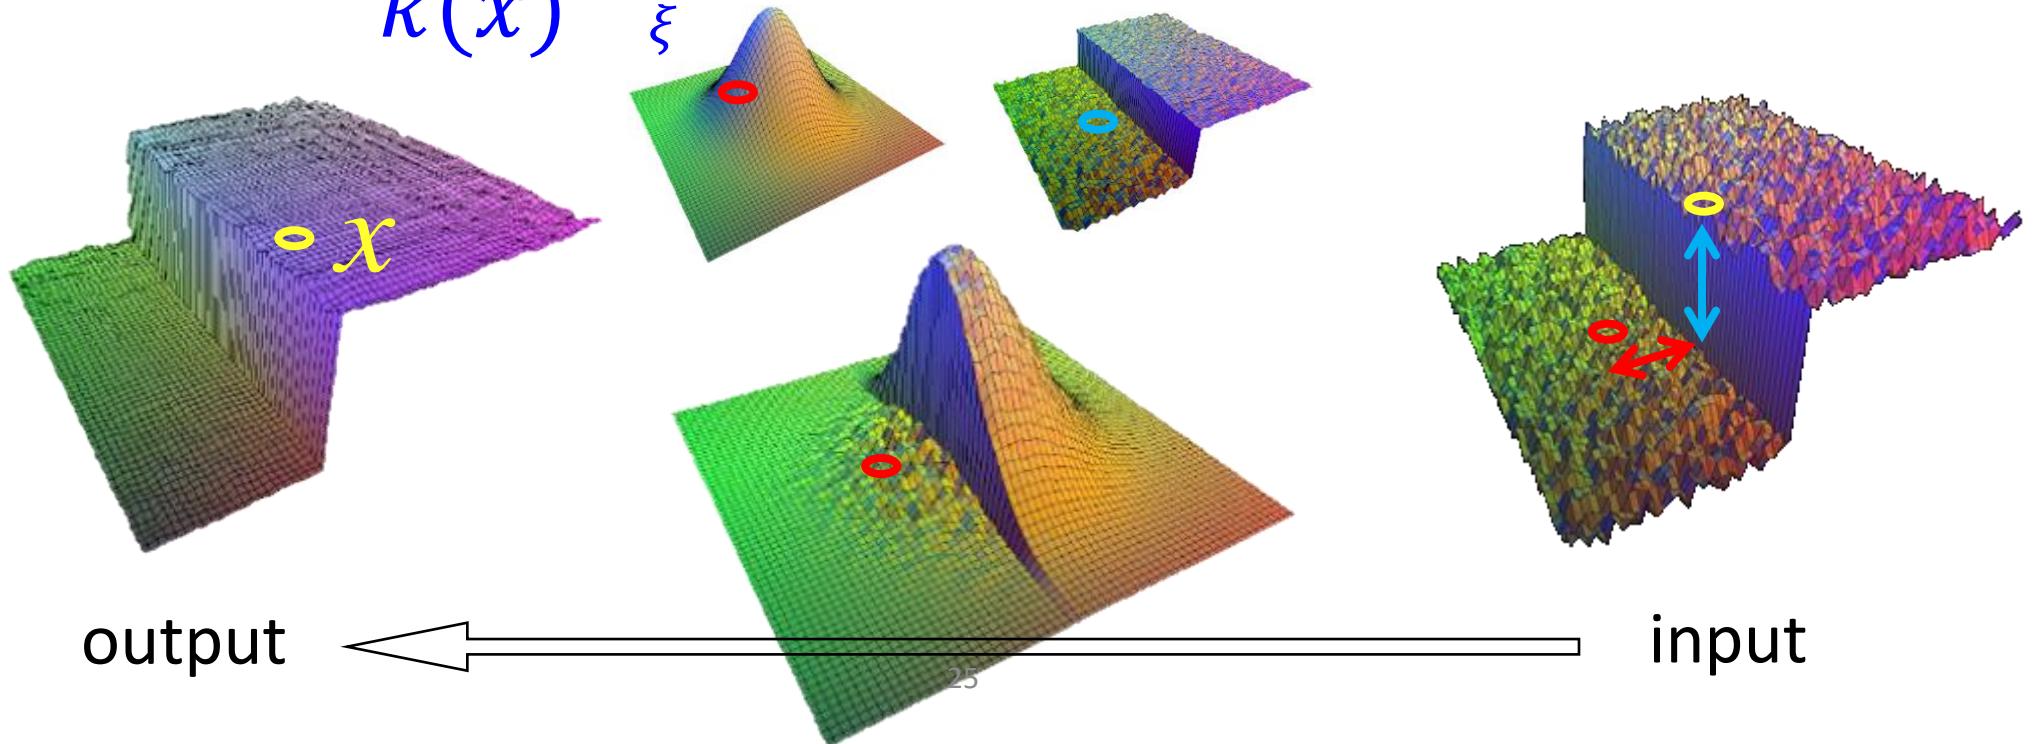
input

# Bilateral Filter

[Tomasi and Manduchi 1998]

- Combined weight

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$

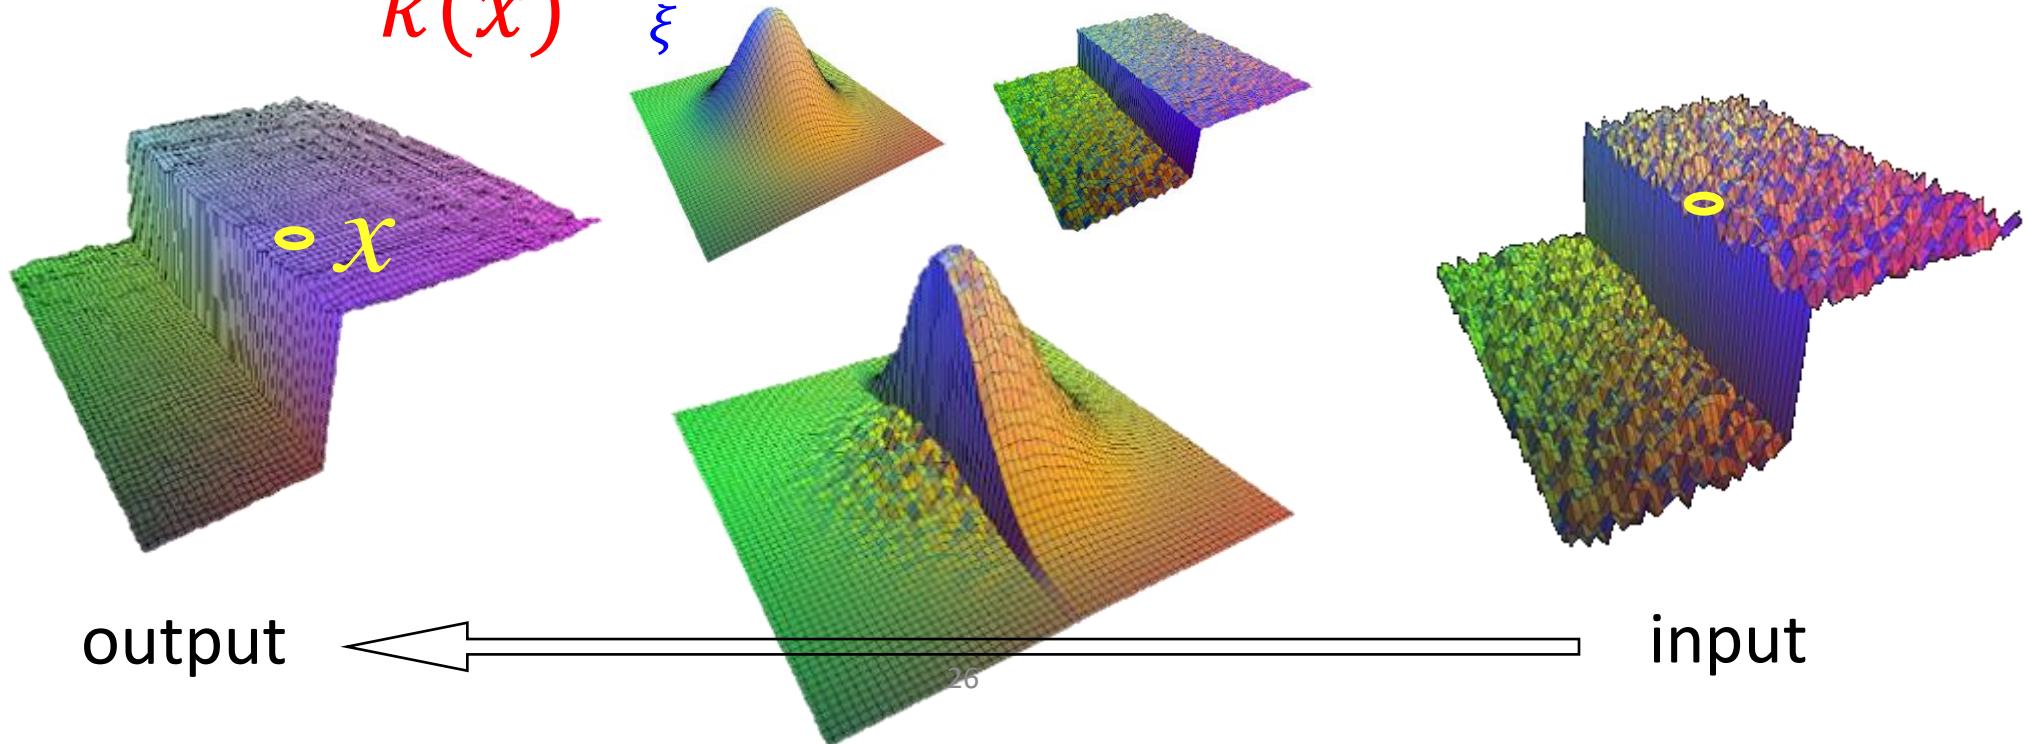


# Normalization Factor

[Tomasi and Manduchi 1998]

$$k(x) = \sum_{\xi} f(x, \xi) g(I(\xi) - I(x))$$

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) \cdot I(\xi)$$

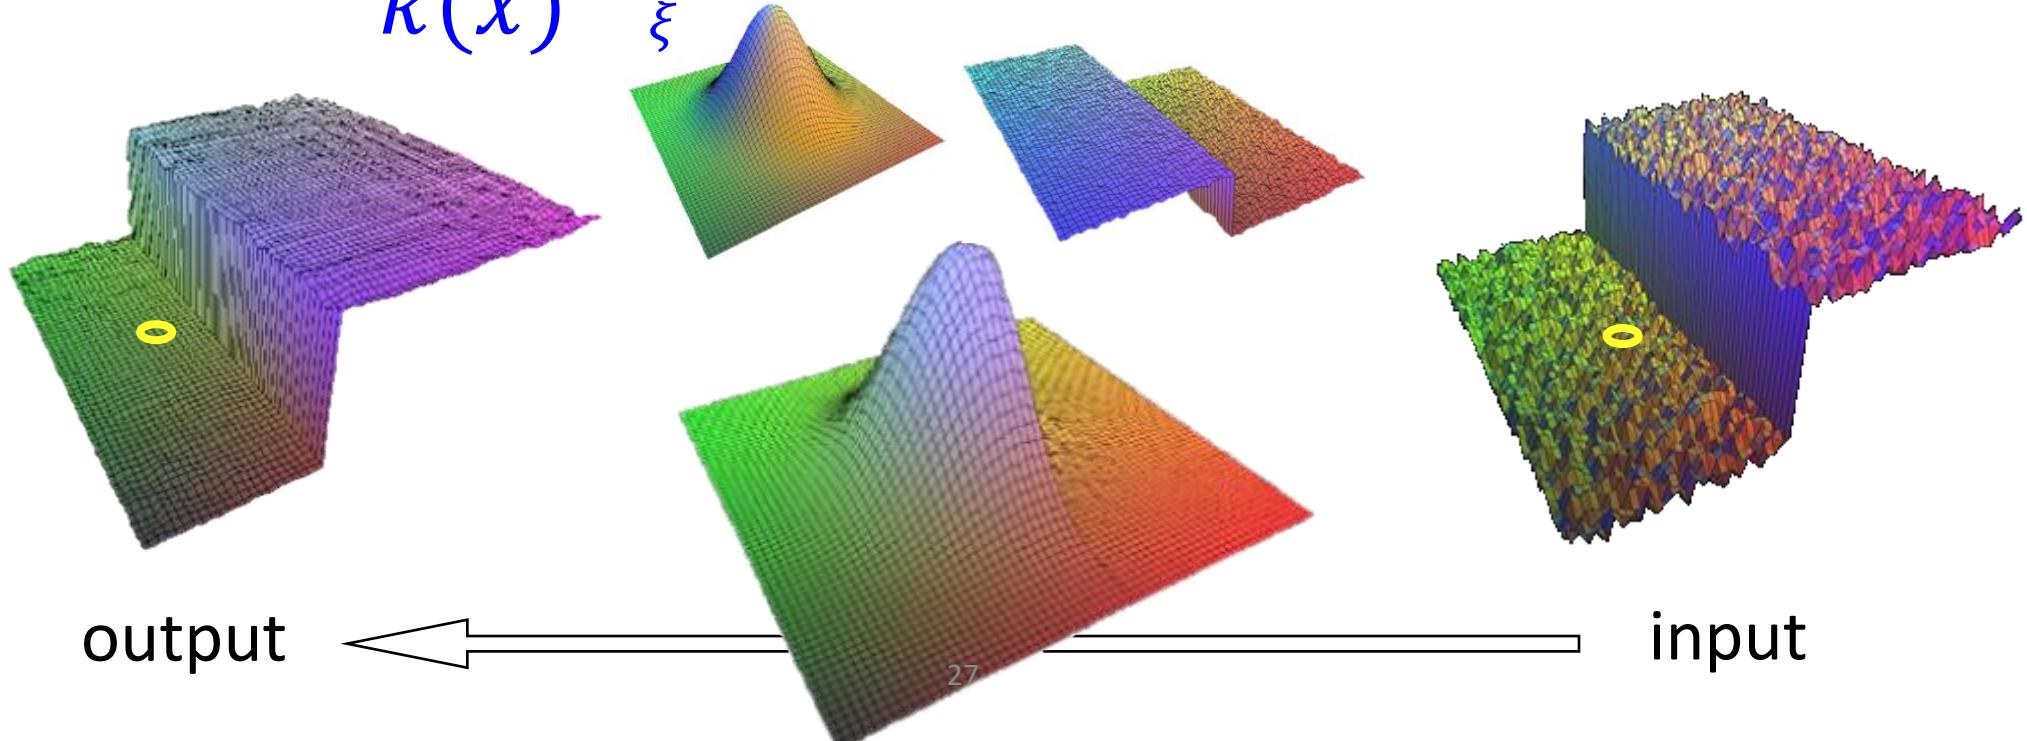


# Bilateral Filter is Spatially-Variant

[Tomasi and Manduchi 1998]

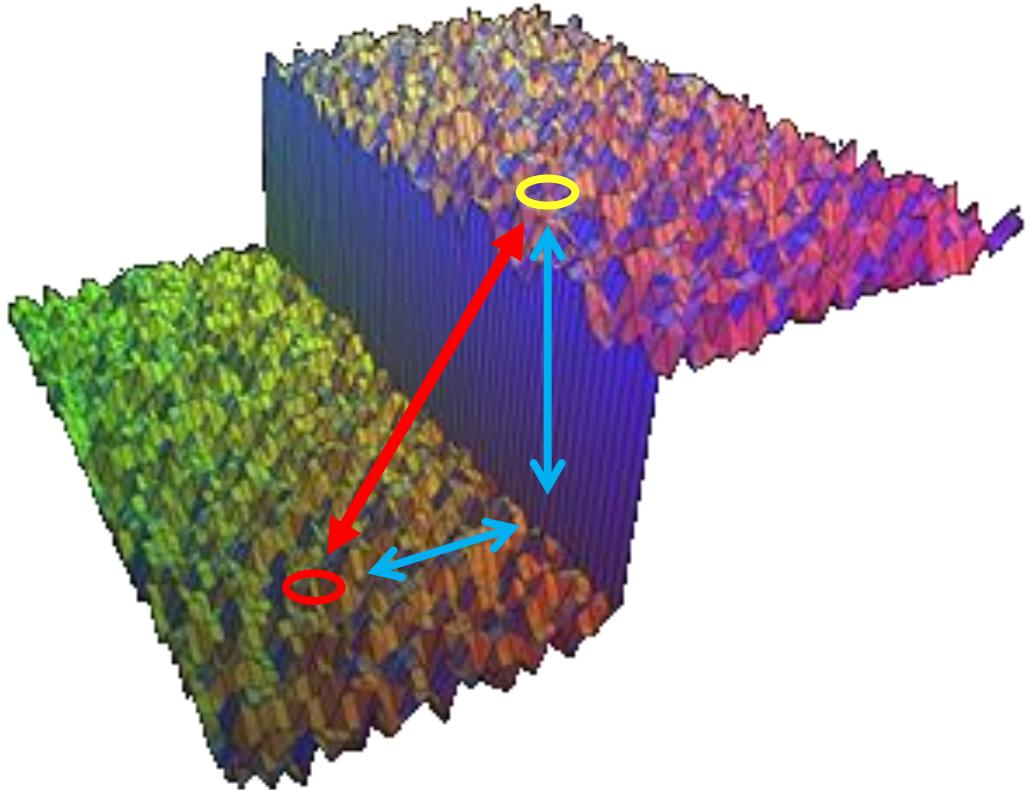
- The weights are different for each output pixel

$$J(x) = \frac{1}{k(x)} \sum_{\xi} f(x, \xi) g(I(\xi) - I(x)) I(\xi)$$



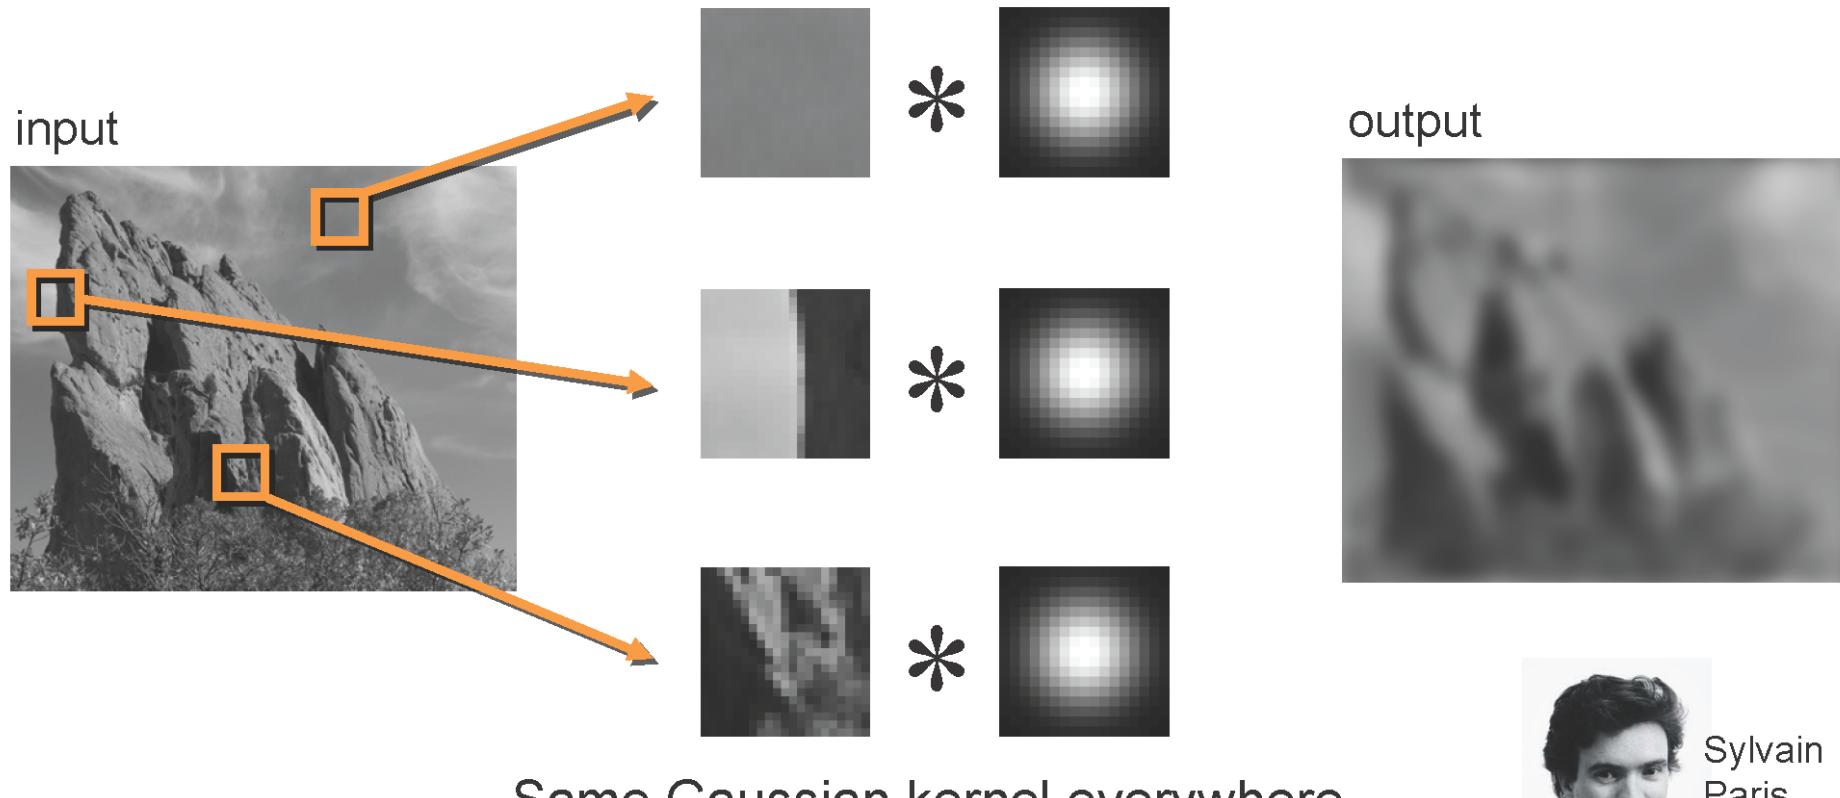
# Other Explanation

- The bilateral filter uses the **3D distance** to compute the weight



# Bilateral Filter vs Gaussian Filter

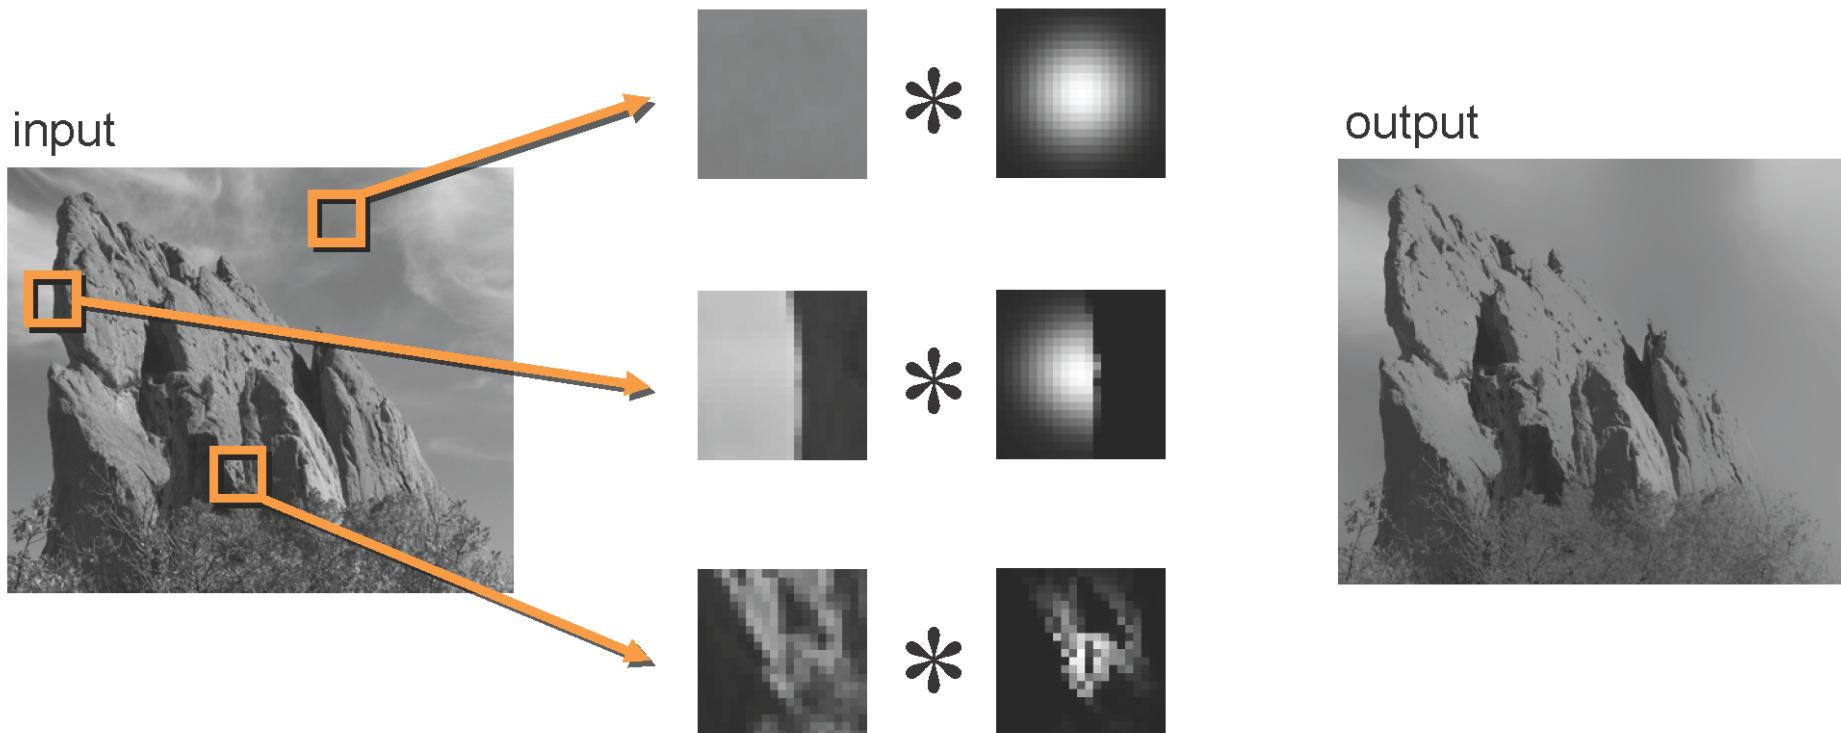
- Bilateral Filter – fix to the Gaussian filter
- Gaussian filtering applies the same filter everywhere



Sylvain  
Paris

# Bilateral Filter vs Gaussian Filter

- Adjust kernel based on image content



The kernel shape depends on the image content.

# Results: Denoise



noisy image



naïve denoising  
Gaussian blur



better denoising  
edge-preserving filter

Smoothing an image without blurring its edges.

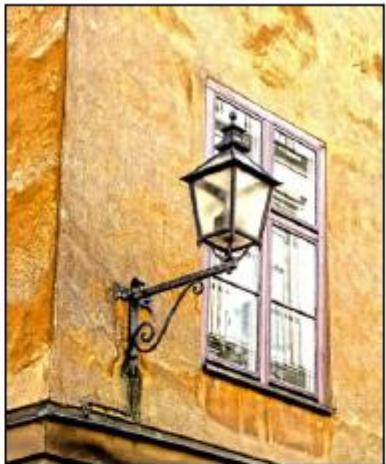
# More Examples



(a) *Photograph*



(b) *Edge-aware smoothing*



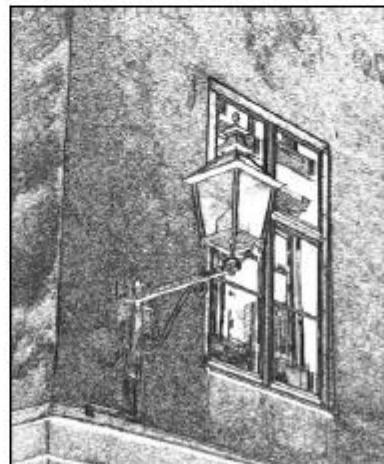
(c) *Detail enhancement*



(d) *Stylization*



(e) *Recoloring*



(f) *Pencil drawing*



(g) *Depth-of-field*

# Questions?





# Fast Bilateral Filtering for the Display of High-Dynamic-Range Images

Frédo Durand  
Julie Dorsey

MIT

SIGGRAPH 2002

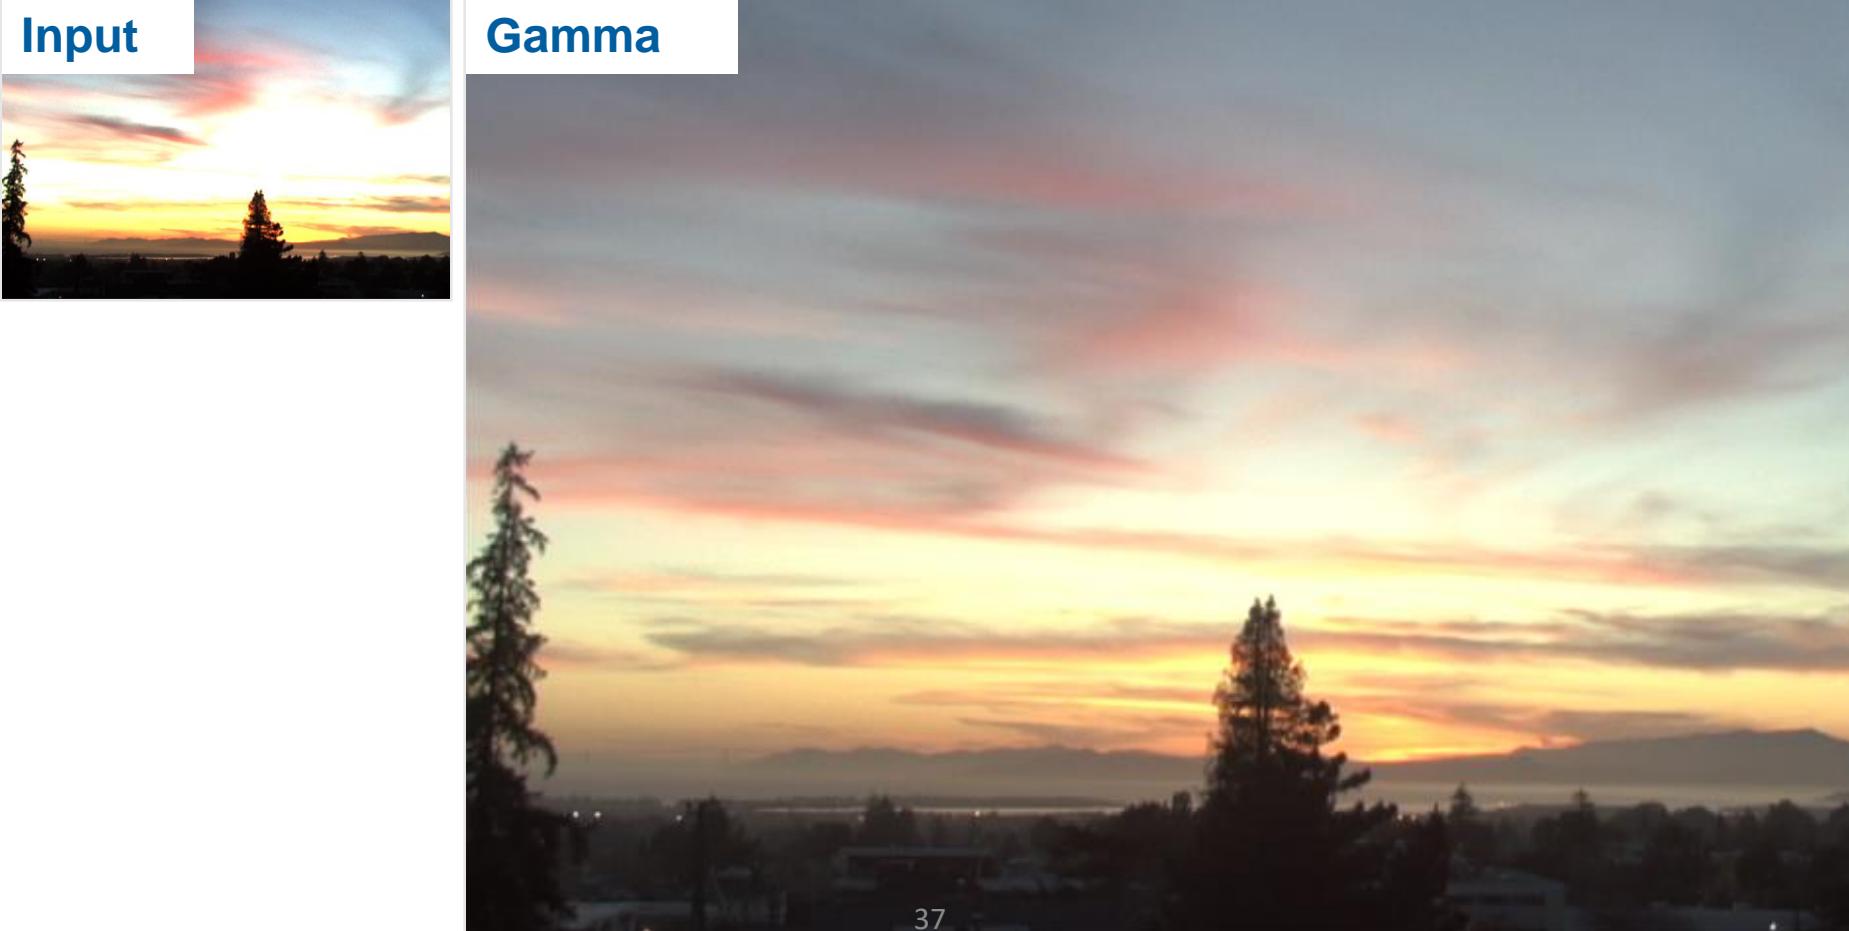
# Show HDR Images in LDR Displayers

- Input: high-dynamic-range image
  - (floating point per pixel)



# Naïve Method: Gamma Compression

- $X \rightarrow X^\gamma$  (where  $\gamma=0.5$  in our case)
- But... colors are washed-out.



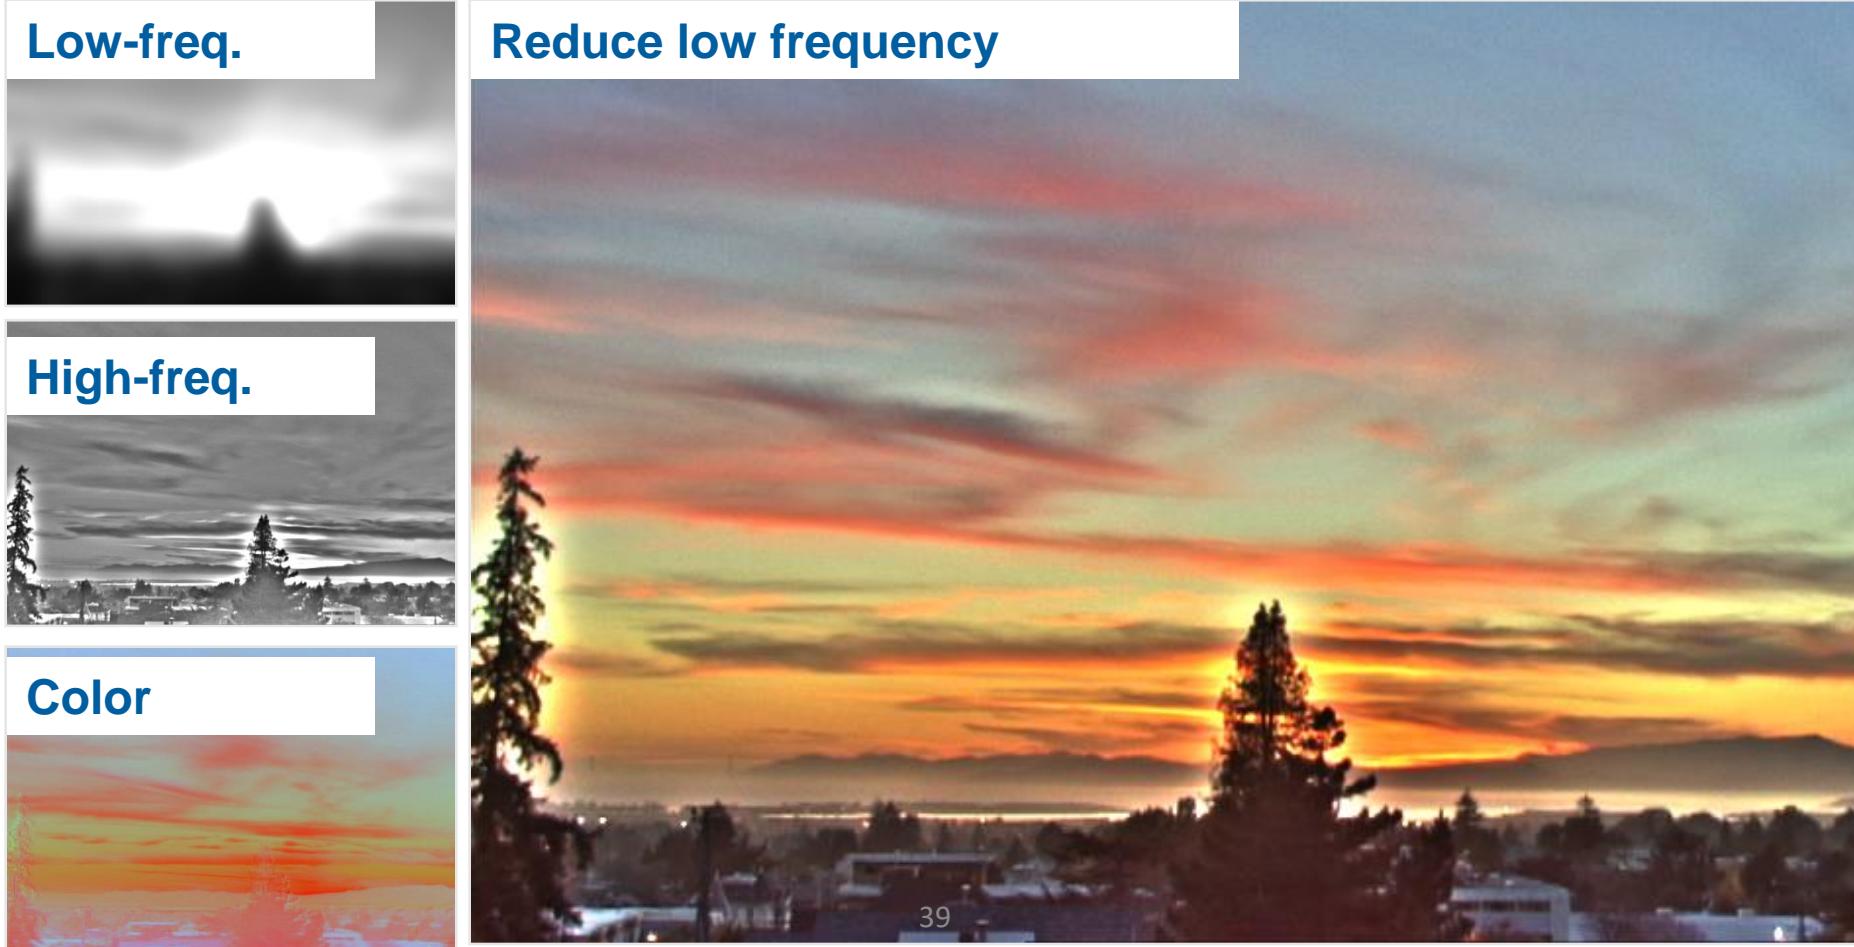
# Gamma Compression on Intensity

- Compress the L channel and keep UV unchanged
- Colors are OK, but details are blurred



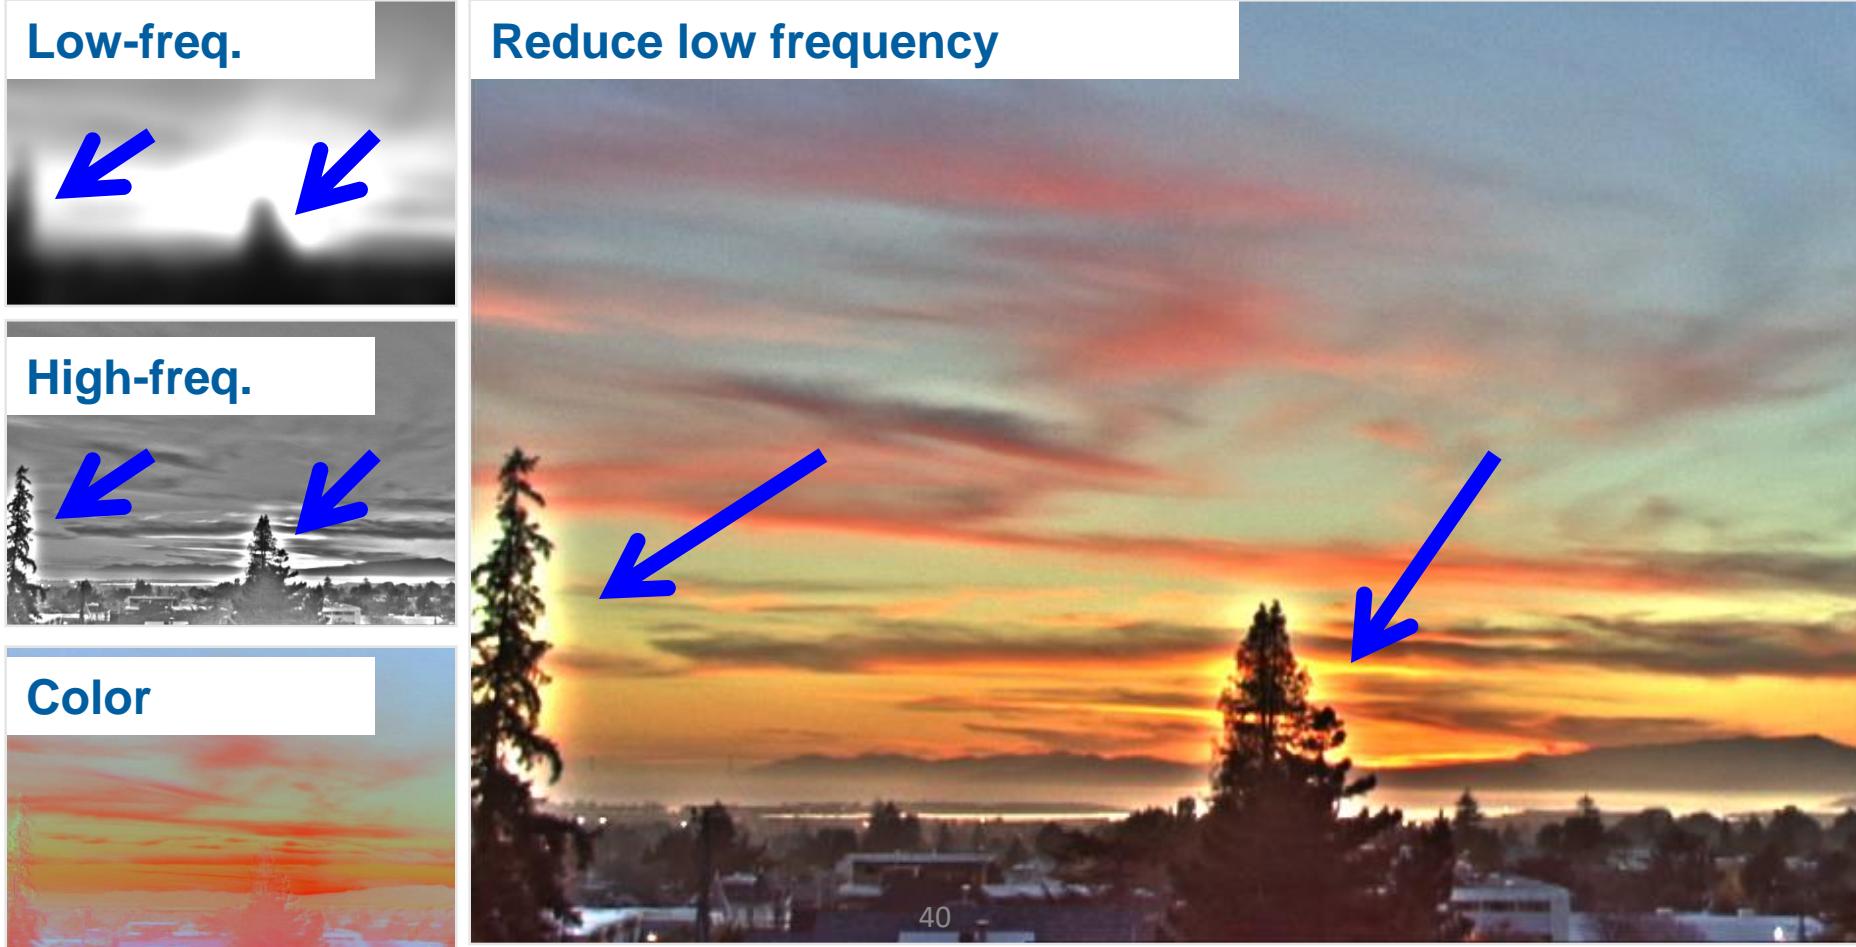
# Oppenheim 1968, Chiu et al. 1993

- Separate frequencies in L, compress low frequency
- Keep high frequencies



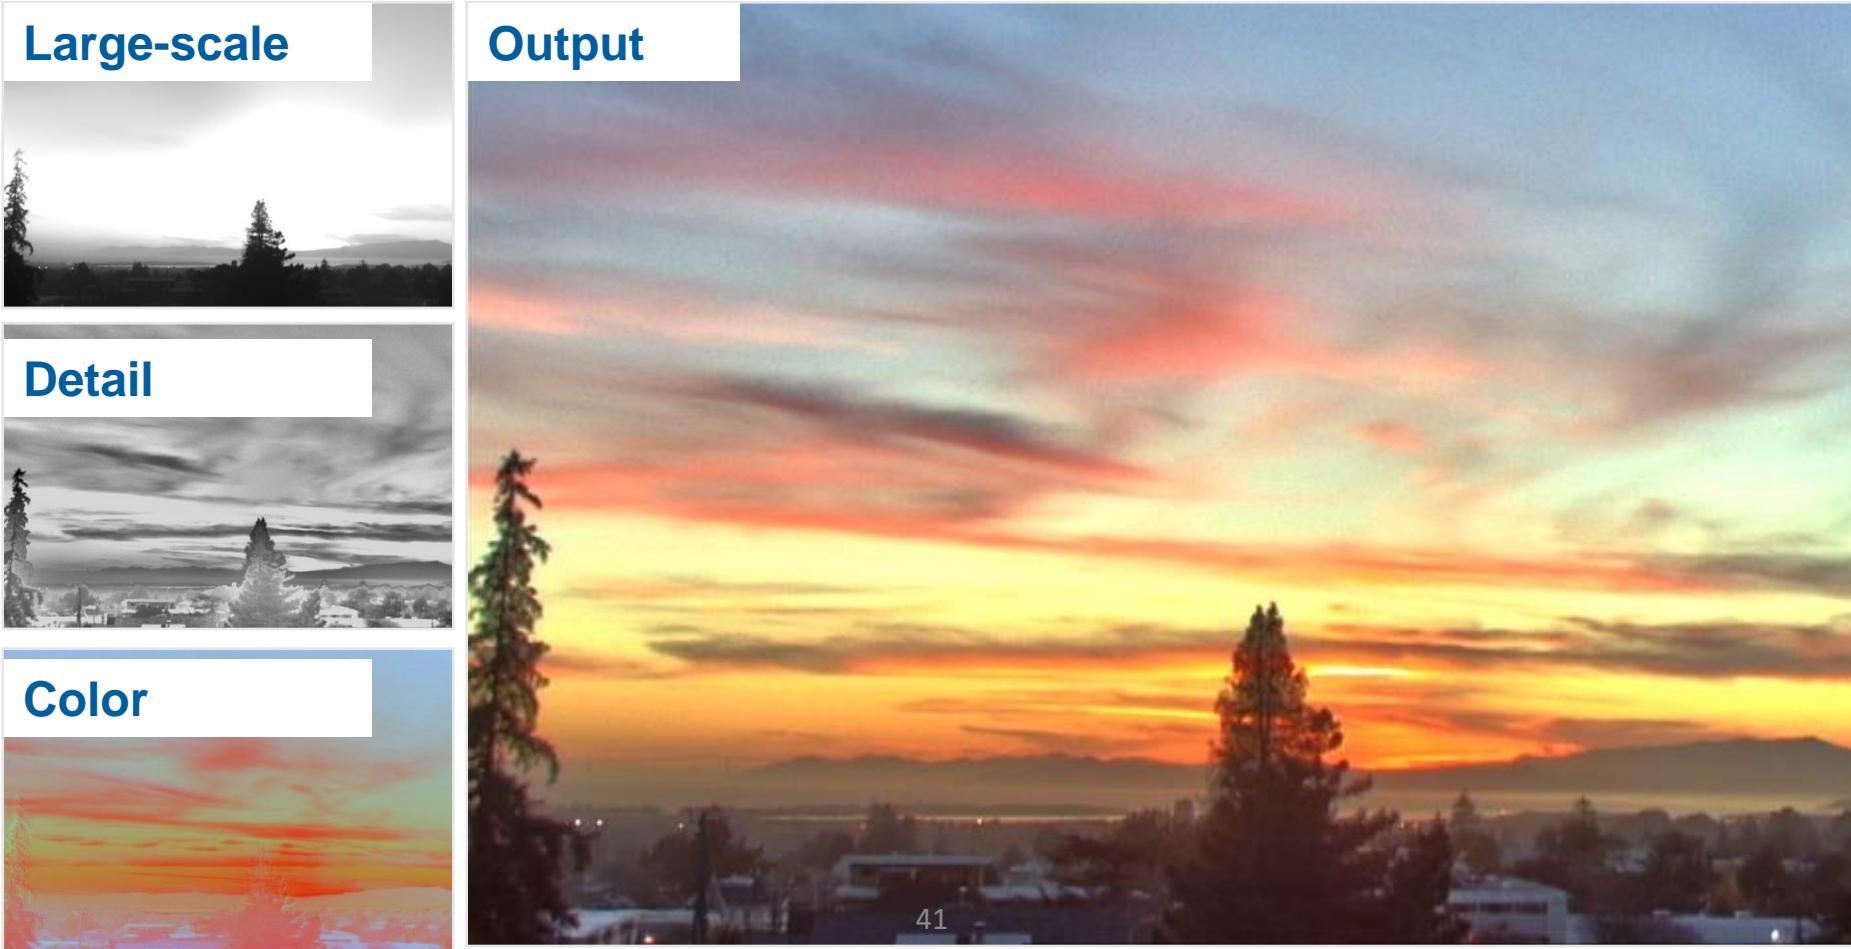
# The Halo Nightmare

- Halo appears at strong edges
- Because high frequency is not well computed



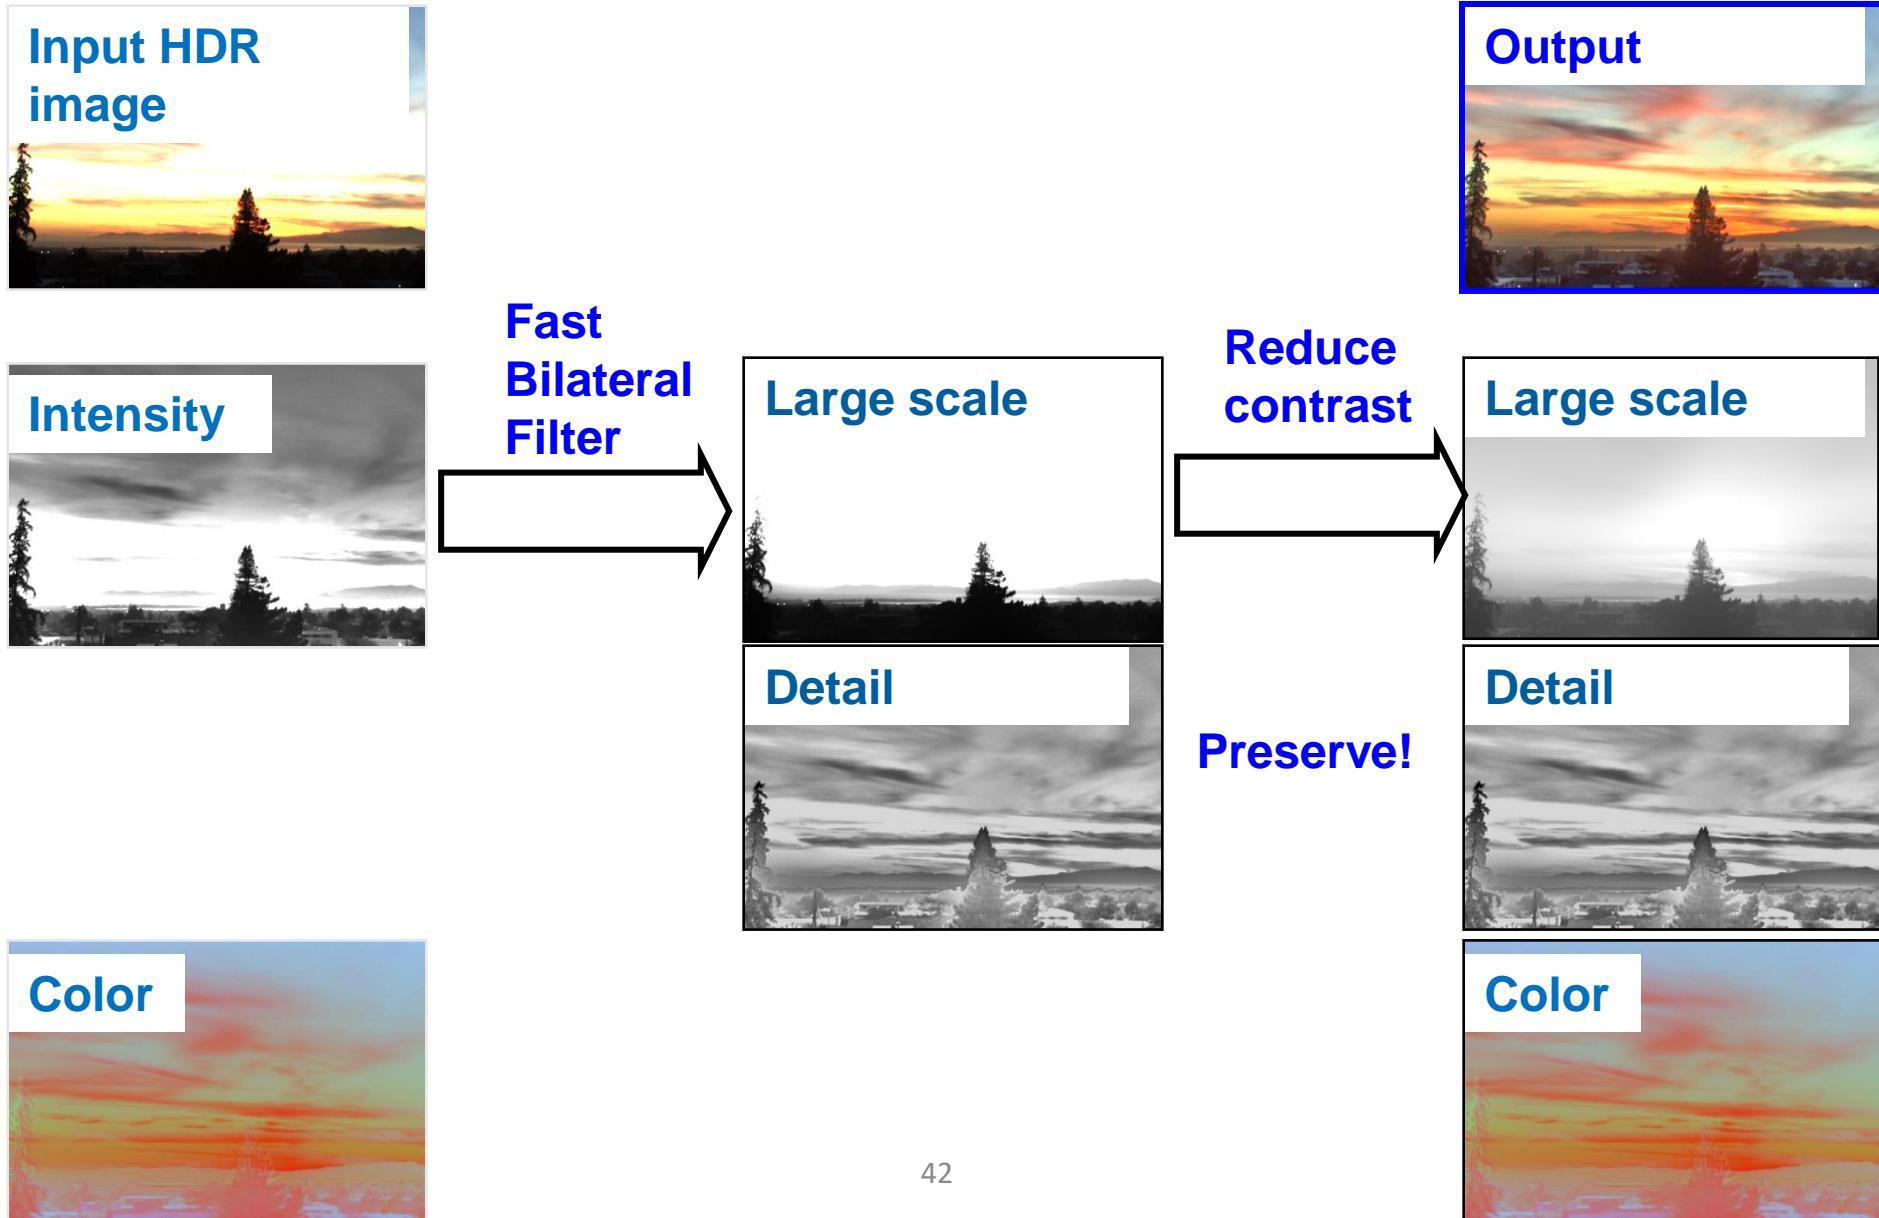
# Separate Frequencies by Bilateral Filter

- Compute high frequency in a better way
- Such that there is no halo





# Overall Pipeline



# Questions?





# Digital Photography with Flash and No-Flash Image Pairs

Georg Petschnigg  
Richard Szeliski

Maneesh Agrawala  
Michael Cohen  
Microsoft Corporation

Hugues Hoppe  
Kentaro Toyama

SIGGRAPH 2004

# Idea



## No-flash

Lighting is “warm and cozy”,  
but image is noisy from sensor

## Flash

Sharp and detailed,  
no noise, but ambience is lost

Can the information be combined?

# Goal



## No-flash

Can you see the noise?  
It is noticeable  
(see original paper)

## Flash

No noise, but color/lighting is  
completely different.

## RESULT

Flash details with  
no flash lighting!



# Bilateral and Joint Bilateral Filter

- Input:  $A$  is the ambient light image (no flash),  $F$  is the flash image
- Key to this paper, use two type of filters

**Bilateral Filter** 
$$A_p^{Base} = \frac{1}{k(p)} \sum_{p' \in \Omega} g_d(p' - p) g_r(A_p - A_{p'}) A_{p'}$$

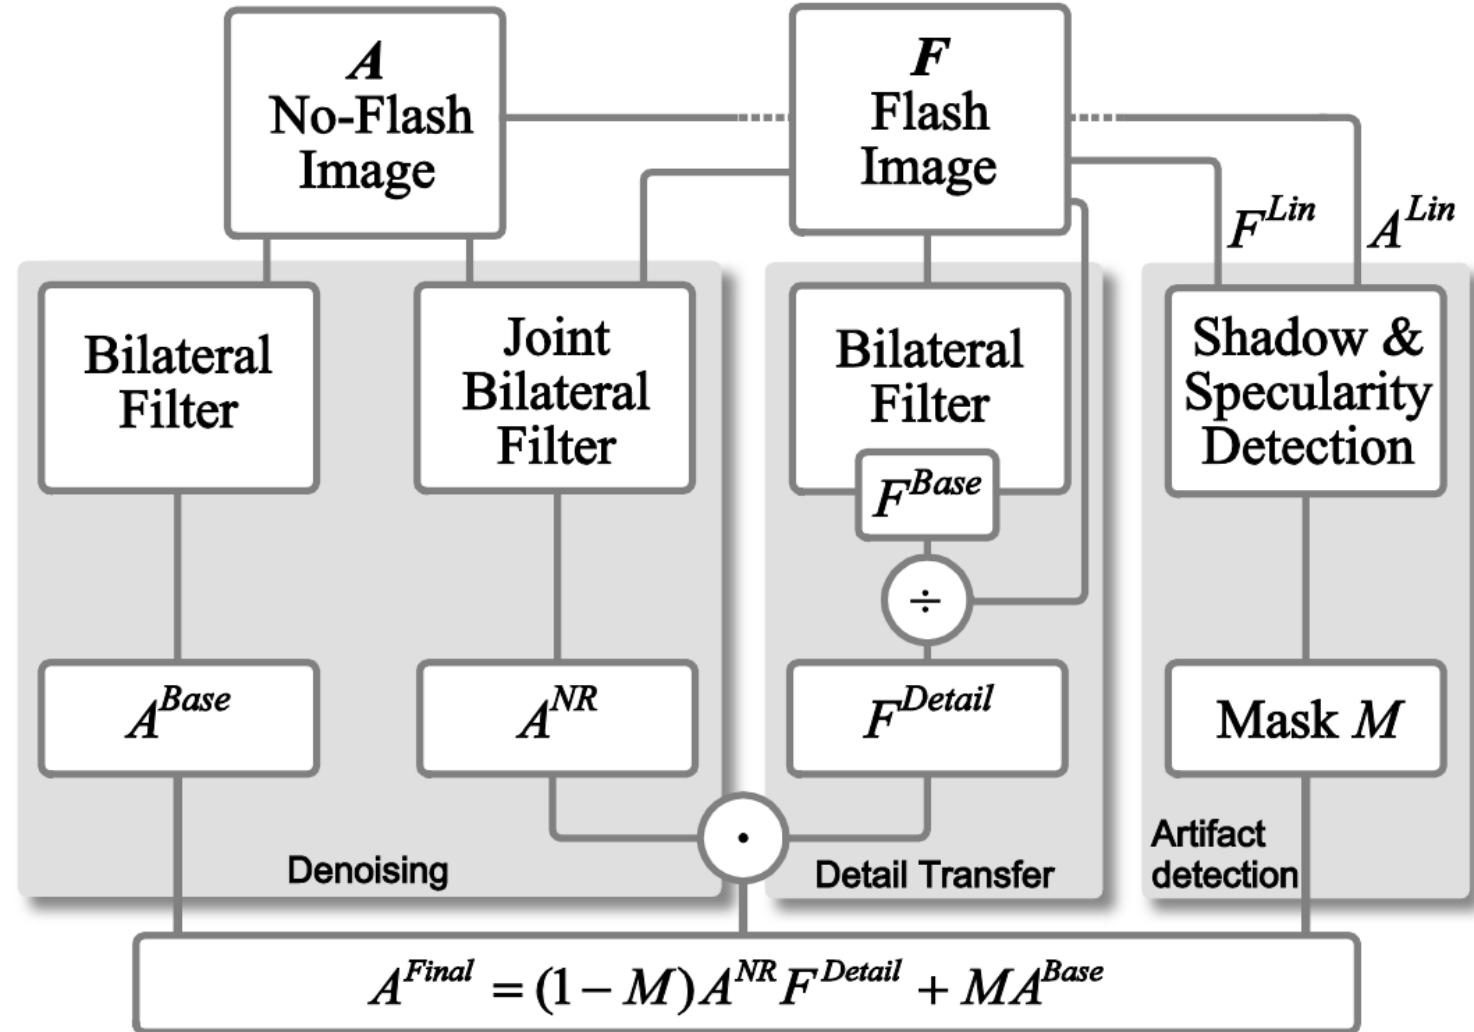
- Use flash image (with little noise) to provide range weight in the bilateral filter

## Joint Bilateral Filter

$$A_p^{NR} = \frac{1}{k(p)} \sum_{p' \in \Omega} g_d(p' - p) \boxed{g_r(F_p - F_{p'})} A_{p'}$$

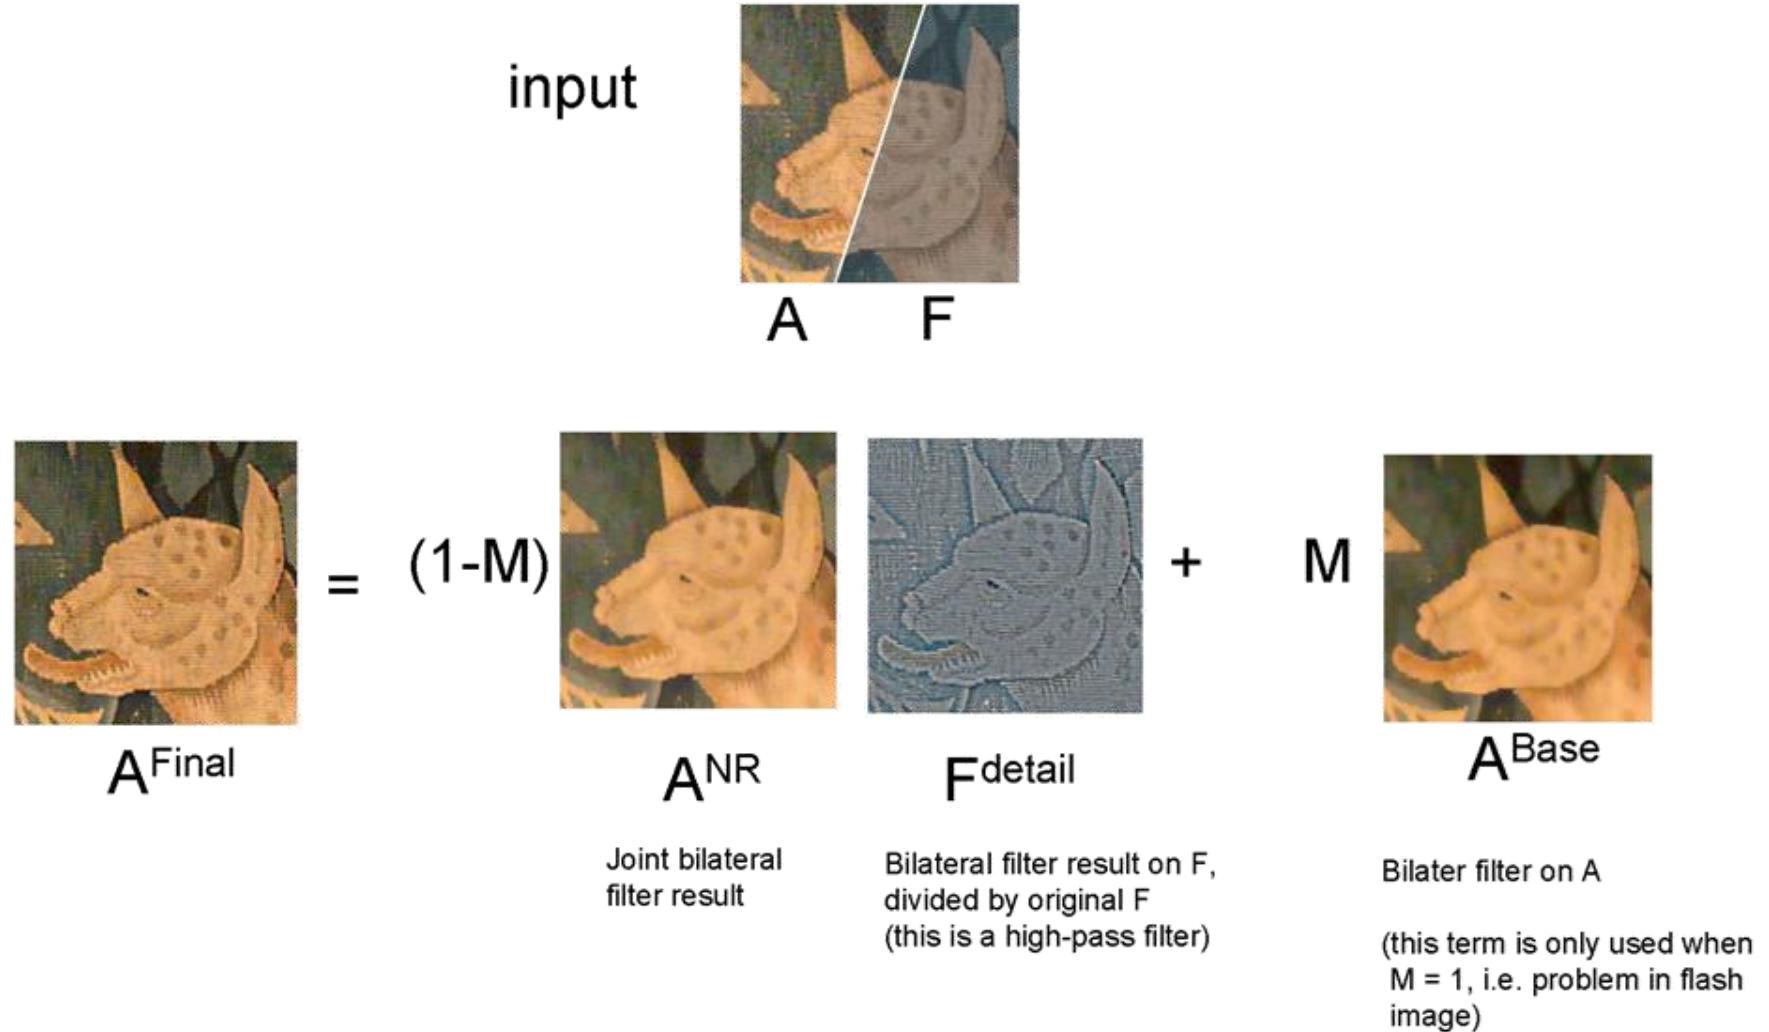
Flash image provides  
range scaling for BF for image A

# Overall Procedure

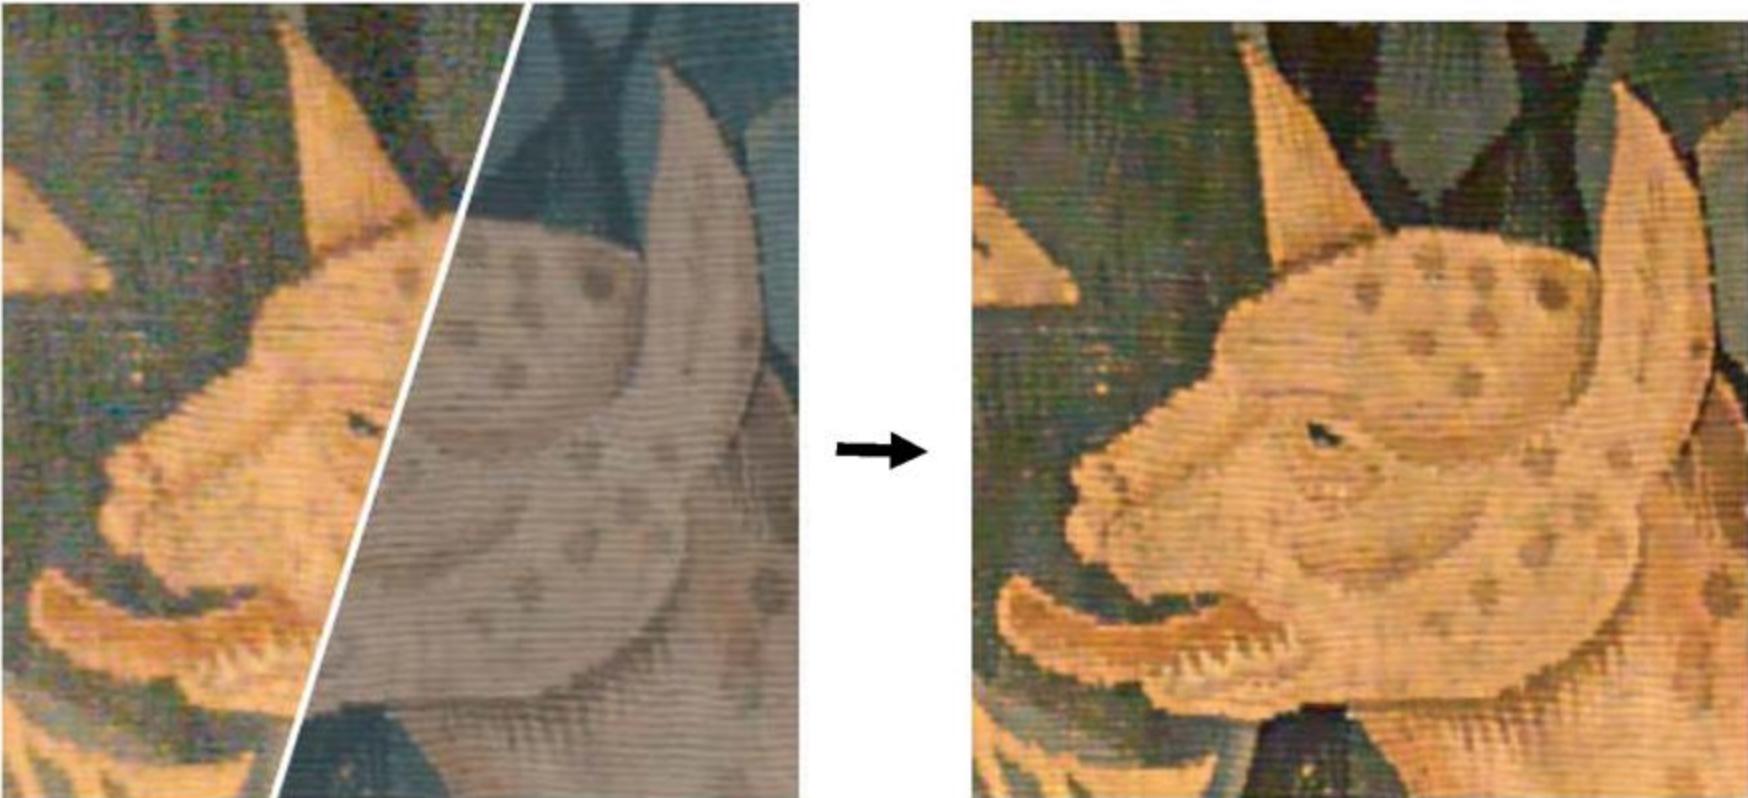


$$F^{Detail} = \frac{F + \varepsilon}{F^{Base} + \varepsilon}$$

# Putting it together



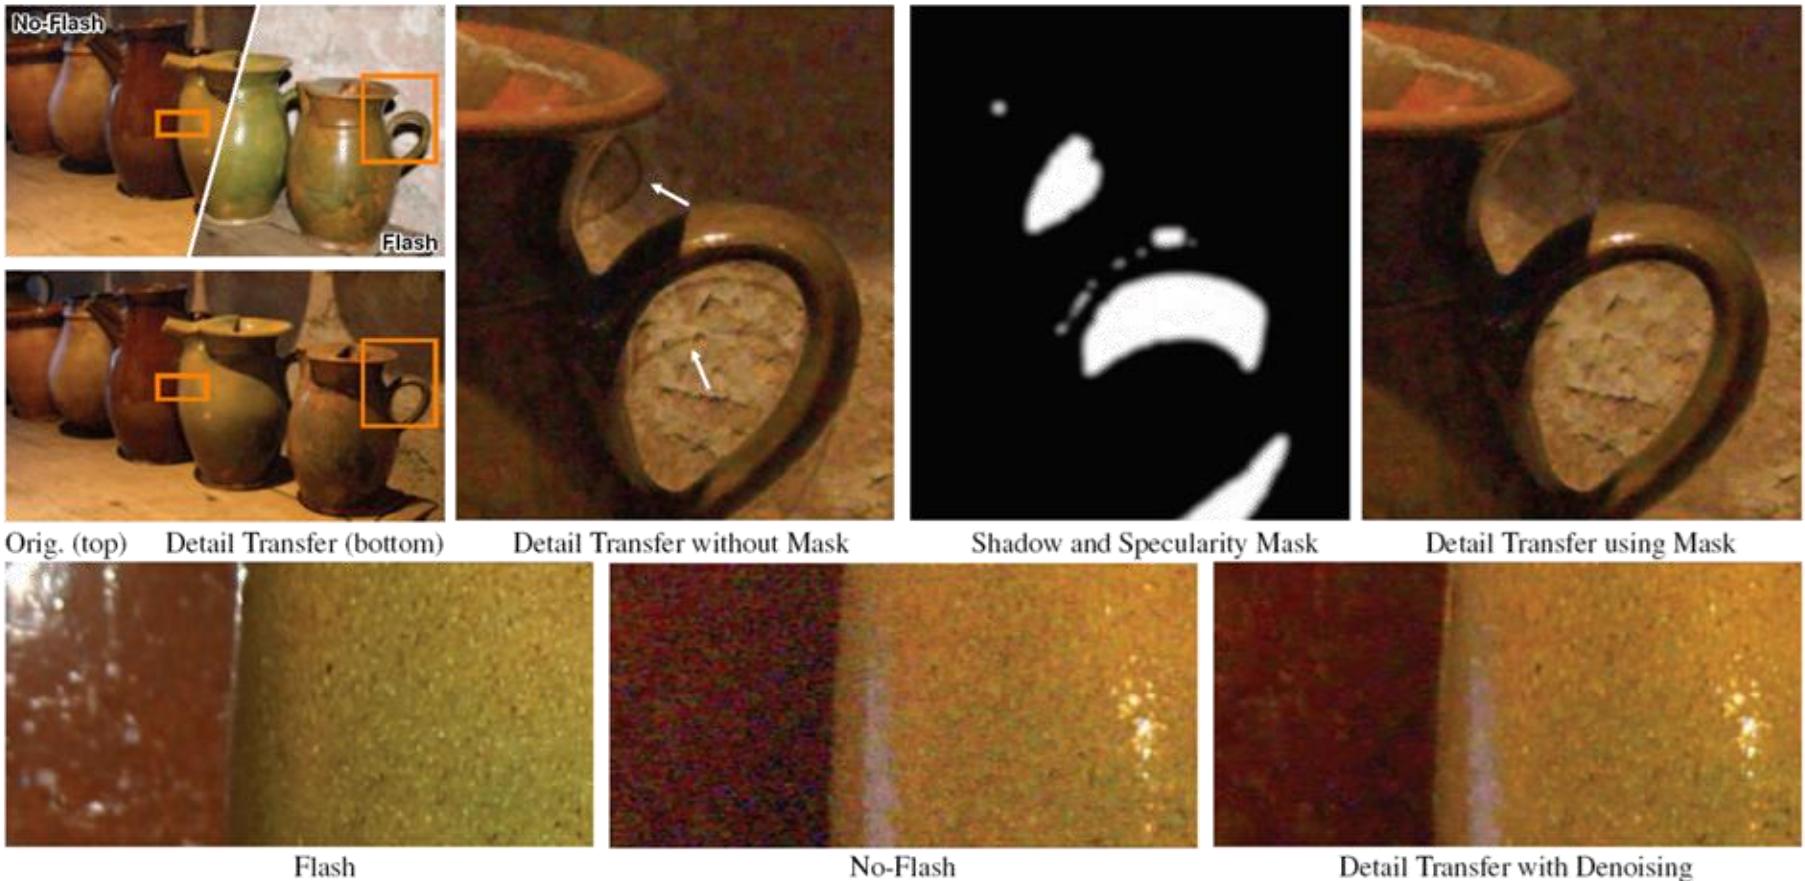
# Example



Input, notice noise in  
A image.

Output

# Mask M



Flash can result in singularities, these are avoided using a binary mask  $M$ . See equations.

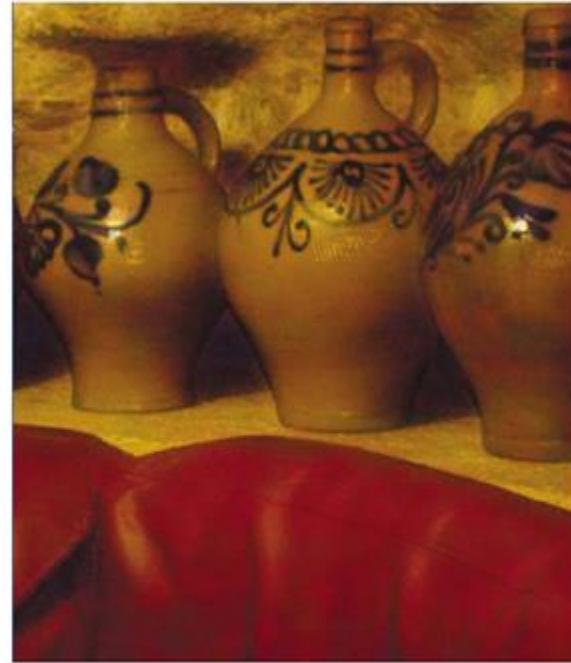
# Results



Flash



No-Flash



Result

# More Results



No-Flash

Flash



Red-Eye Corrected

# Questions?





# Guided Image Filtering

Kaiming He<sup>1</sup>, Jian Sun<sup>2</sup>, and Xiaoou Tang<sup>1,3</sup>

<sup>1</sup> Department of Information Engineering, The Chinese University of Hong Kong

<sup>2</sup> Microsoft Research Asia

<sup>3</sup> Shenzhen Institutes of Advanced Technology, Chinese Academy of Sciences, China

ECCV 2010



# Formulation

- Given an input image  $P$ , the output image  $Q$  is computed as,

$$Q_i = \sum_j W_{ij}(I)P_j$$

- The weights are computed from the guide image  $I$ ,

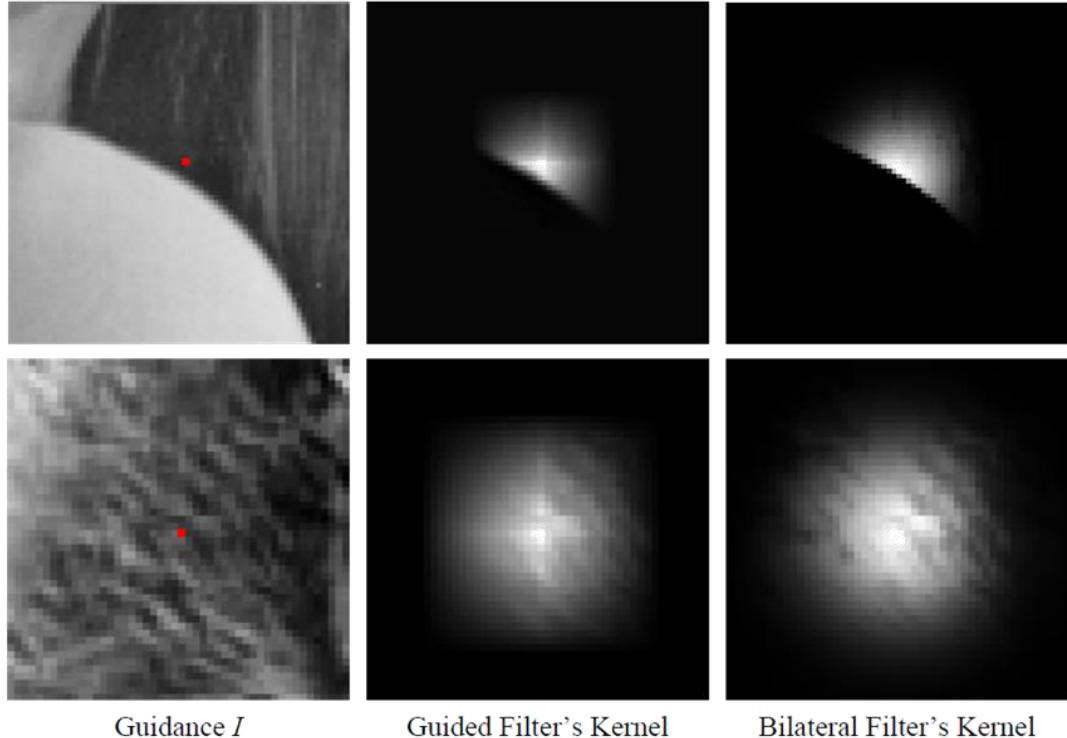
$$W_{ij}(I) = \frac{1}{|\omega|^2} \sum_{k:(i,j) \in \omega_k} \left( 1 + \frac{(I_i - \mu_k)(I_j - \mu_k)}{\sigma_k^2 + \epsilon} \right)$$

the total number of such windows containing both  $i$  and  $j$

the  $k$ -th local window containing both  $i$  and  $j$

$\mu_k, \sigma_k$  are the mean and variance of pixel values in the window  $\omega_k$

# Spatially-Variant Kernals



**Fig. 3.** Filter kernels. Top: a step edge (guided filter:  $r = 7, \epsilon = 0.1^2$ , bilateral filter:  $\sigma_s = 7, \sigma_r = 0.1$ ). Bottom: a textured patch (guided filter:  $r = 8, \epsilon = 0.2^2$ , bilateral filter:  $\sigma_s = 8, \sigma_r = 0.2$ ). The kernels are centered at the pixels denote by the red dots.

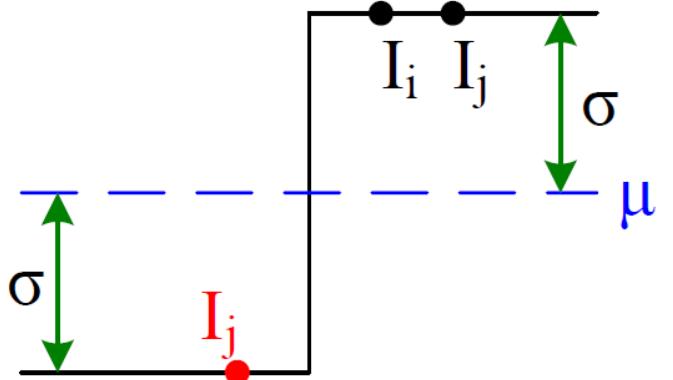
# Explanation

- When  $i$  and  $j$  are from different sides of an edge,

$$\frac{(I_i - \mu_k)(I_j - \mu_k)}{\sigma_k^2} \sim -1$$

- When  $i$  and  $j$  are from the same side,

$$\frac{(I_i - \mu_k)(I_j - \mu_k)}{\sigma_k^2} \sim 1$$



$$W_{ij}(I) = \frac{1}{|\omega|^2} \sum_{k:(i,j) \in \omega_k} \left( 1 + \frac{(I_i - \mu_k)(I_j - \mu_k)}{\sigma_k^2 + \epsilon} \right)$$



# Justification

- How is the filter kernel designed??
- Inspired by the Matting Laplacian
  - Assuming the output  $Q$  is a linear scale of  $I$  within a local window  
i.e.  $Q_i = a_k I_i + b_k, \forall i \in \omega_k$
  - Or, more intuitively,  $q$  has an edge only if  $I$  has an edge,  
i.e.  $\nabla Q = \nabla I$
- More in the paper (one Jacobian iteration of solving the Matting Laplacian)



# Better Than Bilateral Filter

- $O(N)$  time,  $N$  is the number of pixels
  - Bilateral filter is  $O(Nr^2)$ ,  $r$  is the local window size

$$Q_i = \sum_j W_{ij}(I)P_j$$

$$W_{ij}(I) = \frac{1}{|\omega|^2} \sum_{k:(i,j) \in \omega_k} \left( 1 + \frac{(I_i - \mu_k)(I_j - \mu_k)}{\sigma_k^2 + \epsilon} \right)$$

$\mu_k, \sigma_k$  can be computed with box filter ( $O(N)$ ).



# Homework 1: Photometric Stereo

- Deadline: 2023.12.17 23:59

Homework 1: Photometric Stereo

编辑

作业属性

作业批改

成绩统计

占成绩比例

10.0%

公布成绩时间

不公布

作业交付截止

2023.12.20 23:59

作业形式

个人作业 (全部学生: 120人)

计分规则

最后一次得分

完成指标

提交作业

评分方式 (教师评阅 100.0%)

教师评阅

占成绩比例

100.0%

Please refer to the supplementary files (附件) for the detailed illustration of the homework (Homework 1.pdf) and the test data (Homework 1.zip).

Note that the deadline for the submission is 23:59, December 17th (2023). You can still upload it until December 20th (2023), but the penalty of 20% score reduction per 12 hours after ddl will be applied.



# Submission & Late Policy

- Assignments are submitted via Learning in ZJU (学在浙大) .



- For a late submission, penalty of 20% score reduction per 12 hours is applied. If one submits an assignment within 12 hours after the deadline, the score becomes 80%. Within 24 hours, 60%, and so on.