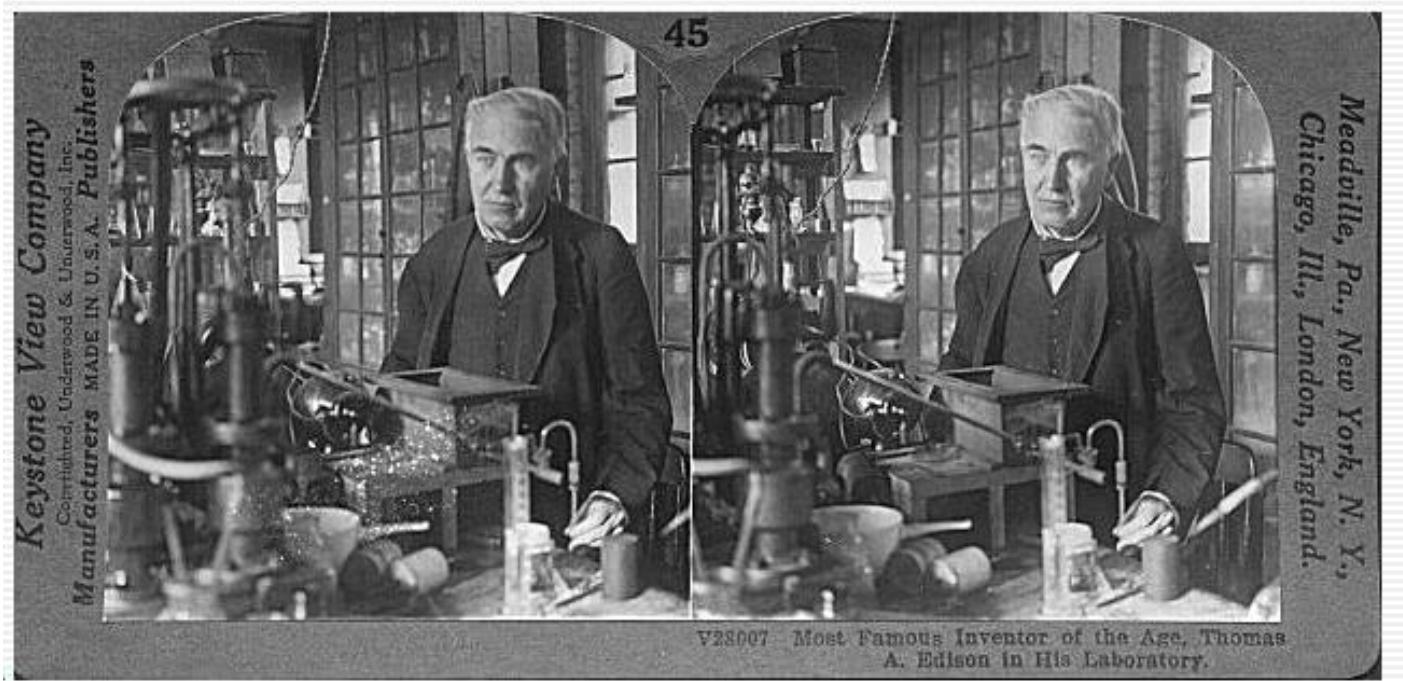


13. Stereo & Optical Flow



Stereo

- Given two calibrated cameras, decide the depth of every pixel in the two images
 - Based on how much each pixel moves between the two images



Invented by Sir Charles Wheatstone, 1838

Stereoscopes: a 19th century pastime



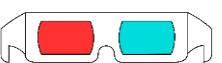
Anaglyph 3D



- Encoding each image using different colors (e.g. red and cyan)



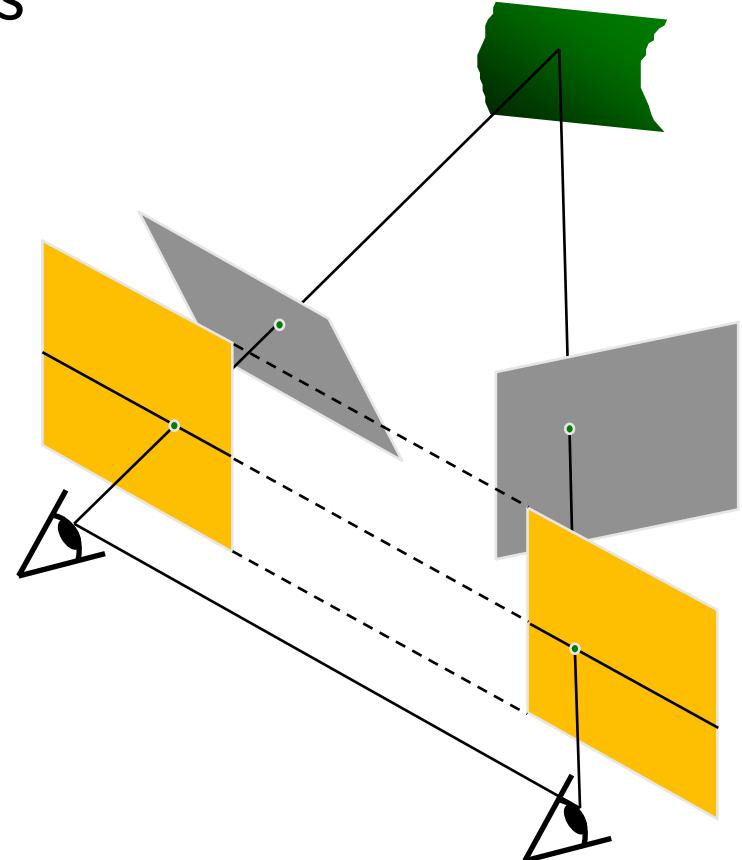
Piero della Francesca, Ideal City in an Anaglyph version



From wikipedia

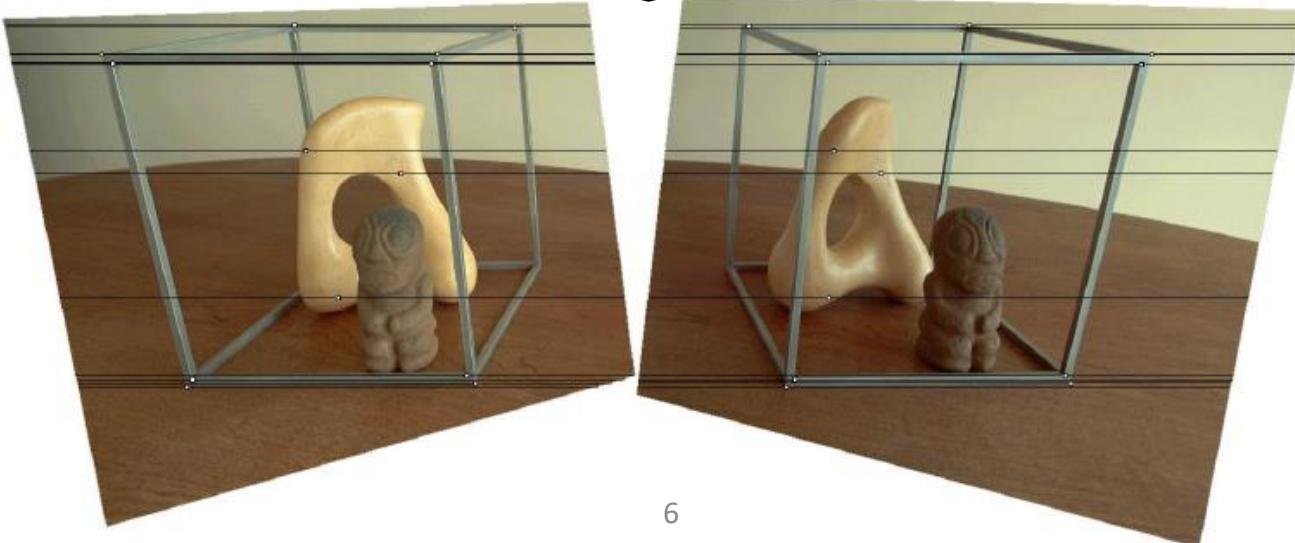
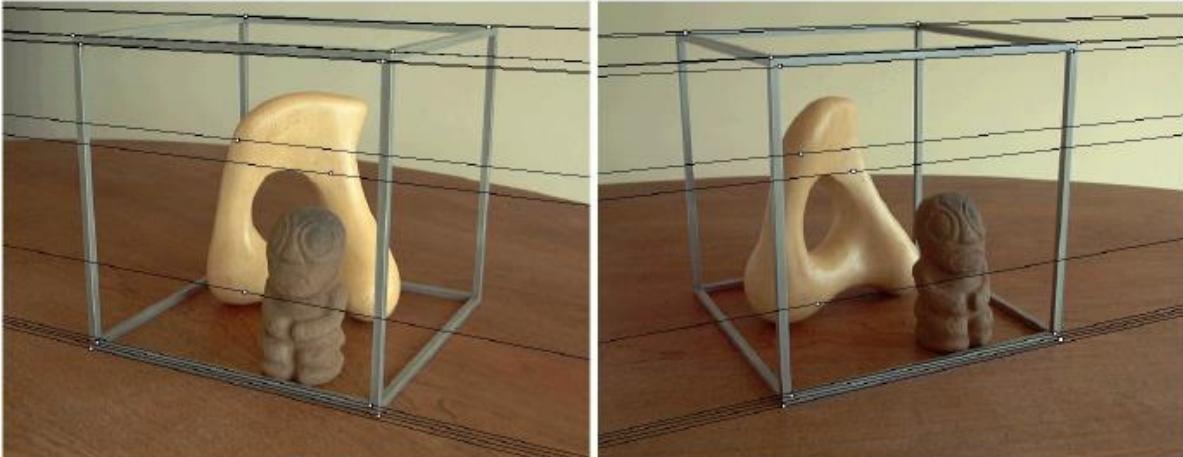
Stereo Image Rectification

- In practice, it is convenient if image scanlines (rows) are the epipolar lines.
- Reproject image planes onto a common plane parallel to the line between optical centers
 - Rotating both cameras without translation
 - Applying a homography to the image
- Pixel motion is horizontal after this transformation



C. Loop and Z. Zhang. Computing Rectifying Homographies for Stereo Vision. CVPR 1999.

Stereo Image Rectification: example

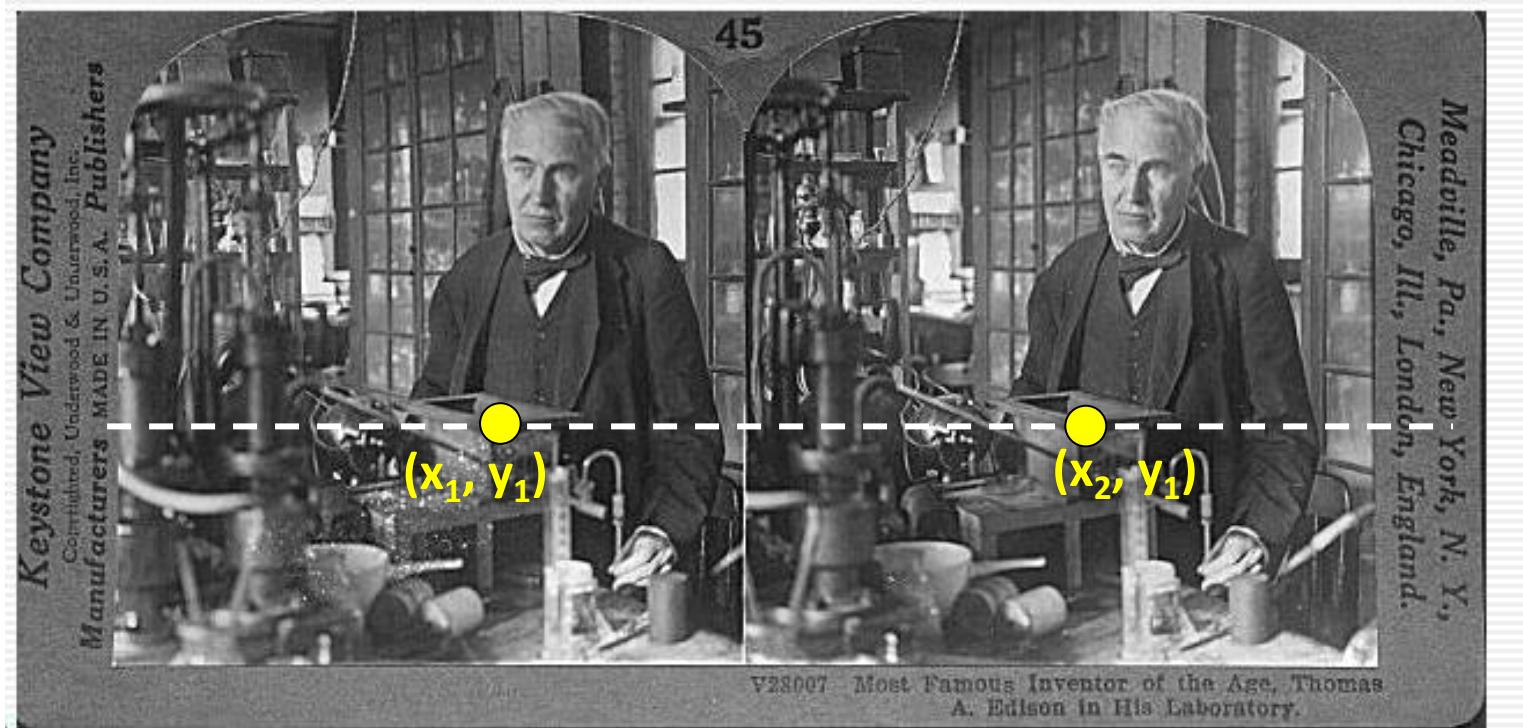




Epipolar Geometry

The correspondence point lies on the epipolar line

epipolar
lines

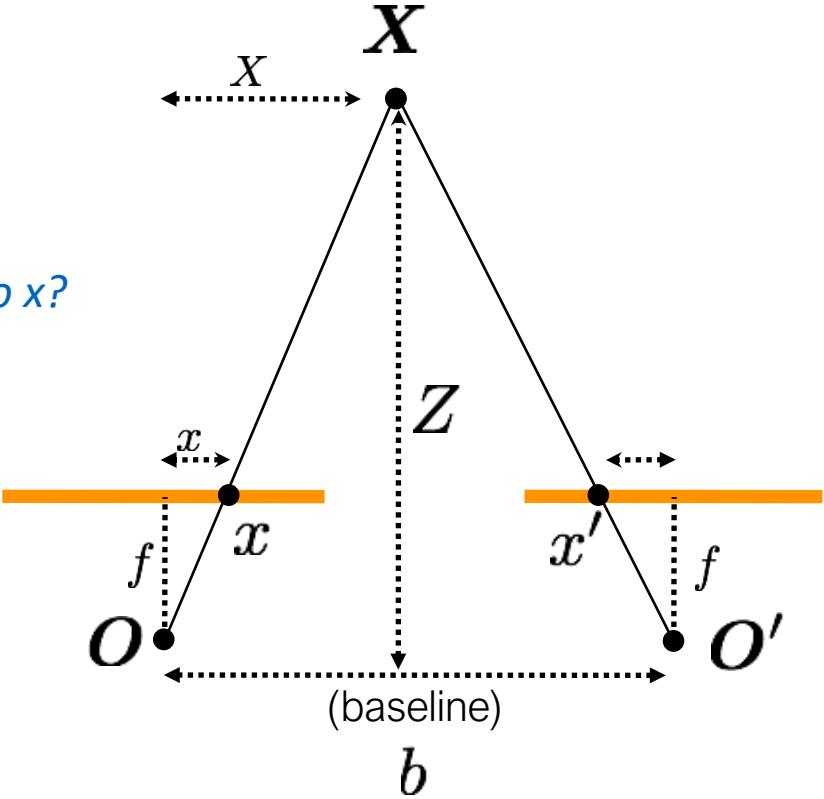


Two images captured by a purely horizontal translating camera

$x_2 - x_1$ = the *disparity* of pixel (x_1, y_1)

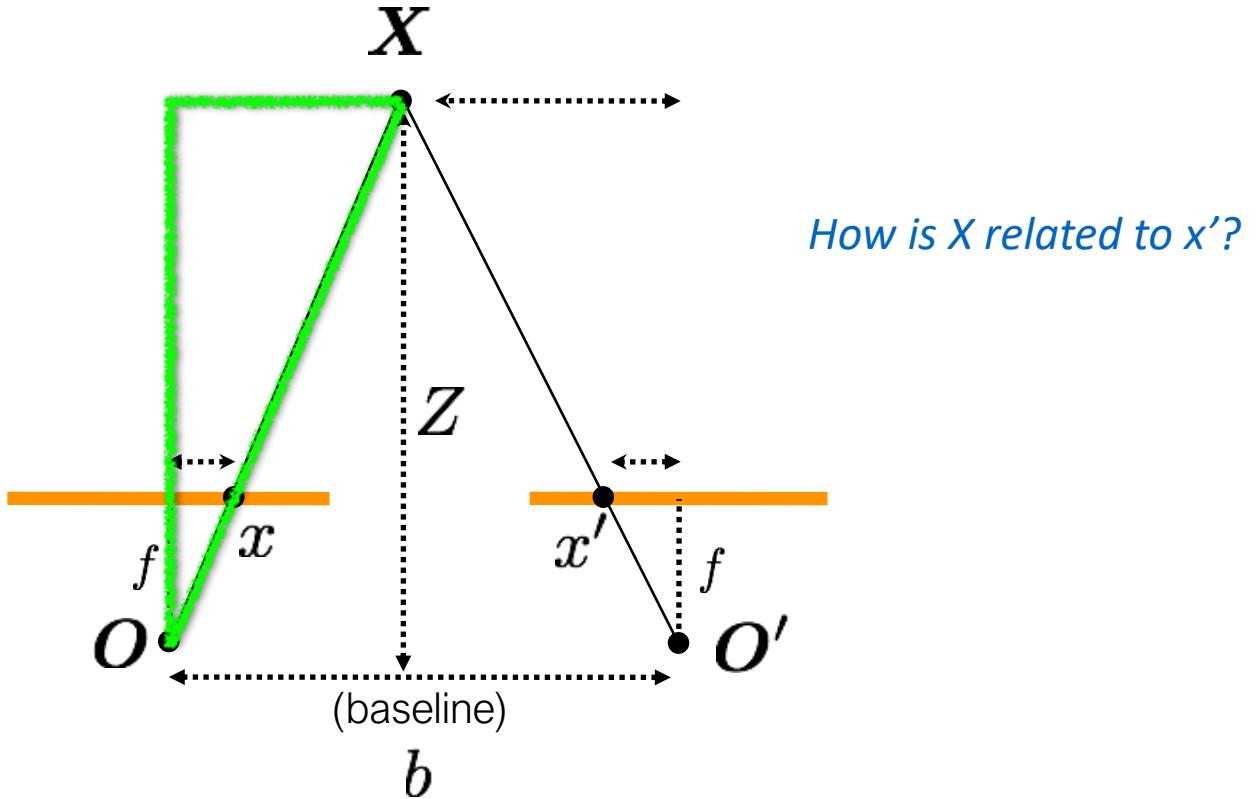
Disparity & Depth

How is X related to x ?



Disparity & Depth

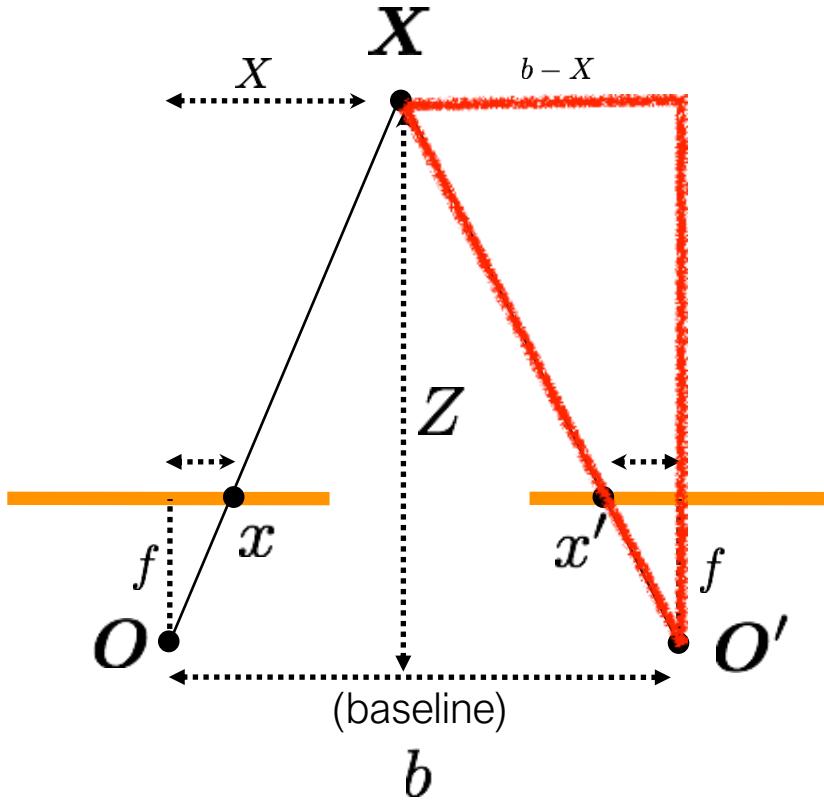
$$\frac{X}{Z} = \frac{x}{f}$$





Disparity & Depth

$$\frac{X}{Z} = \frac{x}{f}$$

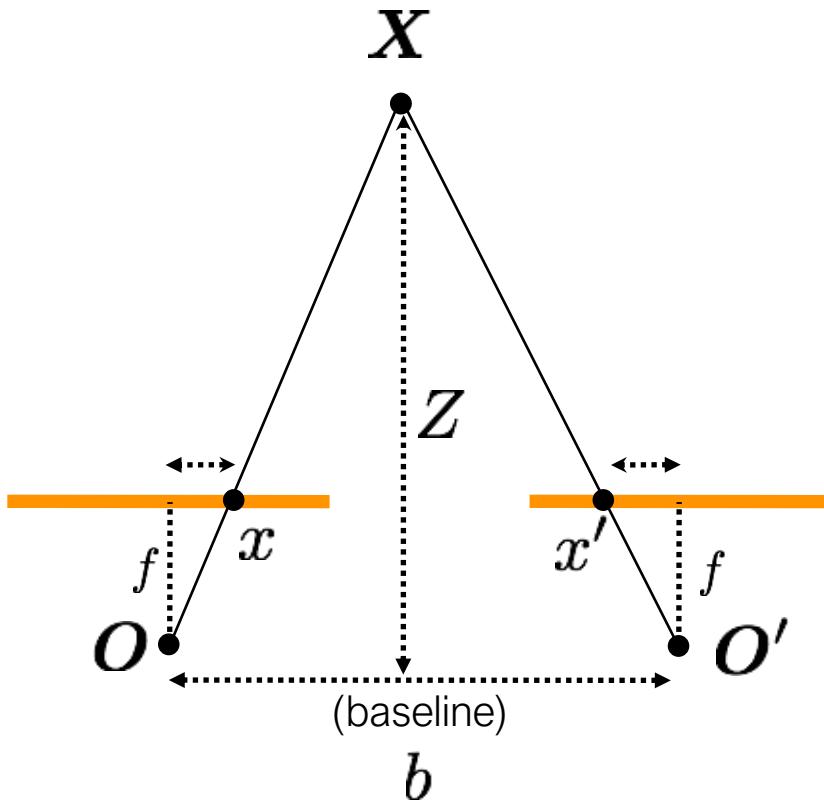


$$\frac{b - X}{Z} = \frac{-x'}{f}$$



Disparity & Depth

$$\frac{X}{Z} = \frac{x}{f}$$



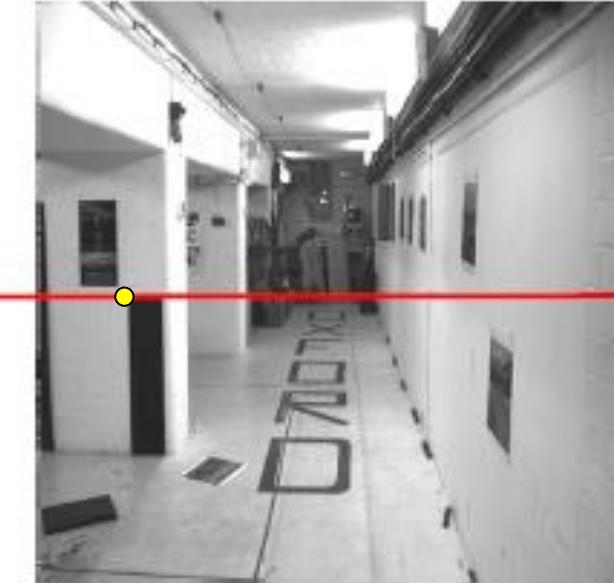
$$\frac{b - X}{Z} = \frac{-x'}{f}$$

Disparity $d = x - x' = \frac{bf}{Z}$ inversely proportional to depth

Questions?



Local Methods



For each epipolar line

f

For each pixel in the left image

- compare with every pixel on the same epipolar line in right image, e.g.
 $|f(x, y) - g(x + d, y)|^2$
- pick pixel with minimum match cost

Disparity Space Image (DSI)

- At each pixel (x, y) , a cost can be evaluated for each disparity d
- This defines a *cost volume* $C(x, y, d)$, also known as disparity space image (DSI)
- It is often helpful to look at a 2D slice of this cost volume, e.g. by fixing y



$$f(x, y)$$

$$g(x, y)$$



Winner Take All

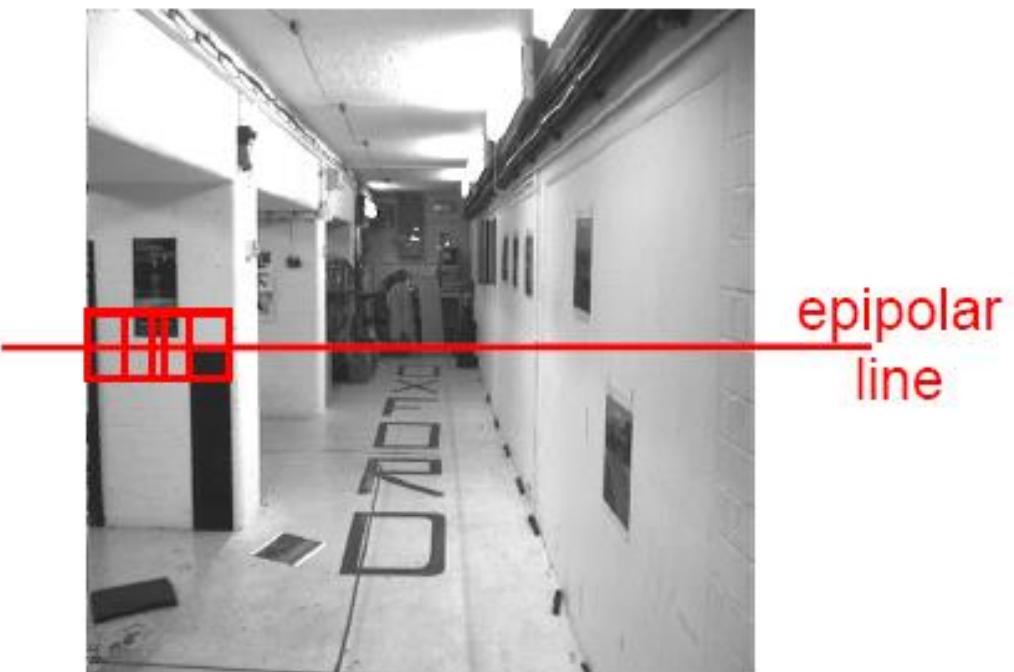
- Choose the minimum of each column in the DSI independently:

$$d(x, y) = \arg \min_{d'} C(x, y, d')$$



Cost Aggregation

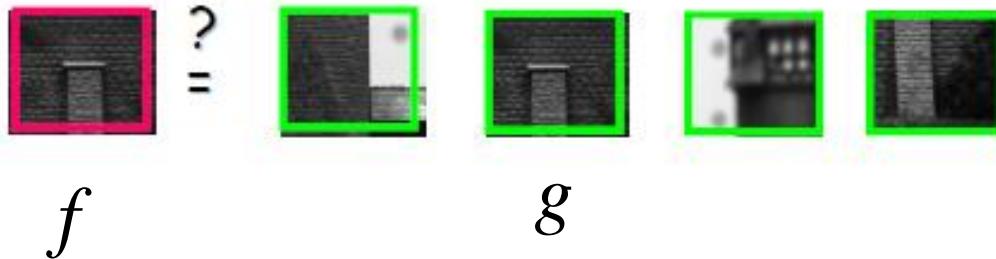
- Some pixels are textureless, matching in the second image is ambiguous
- Evaluate correspondence by comparing a local neighborhood to reduce ambiguity/noise





Cost Aggregation

Questions: Which is the corresponding window in the second image?



$$SSD = \sum_{(x,y)} |f(x,y) - g(x,y)|^2 \quad \text{Sum over costs within a local window}$$

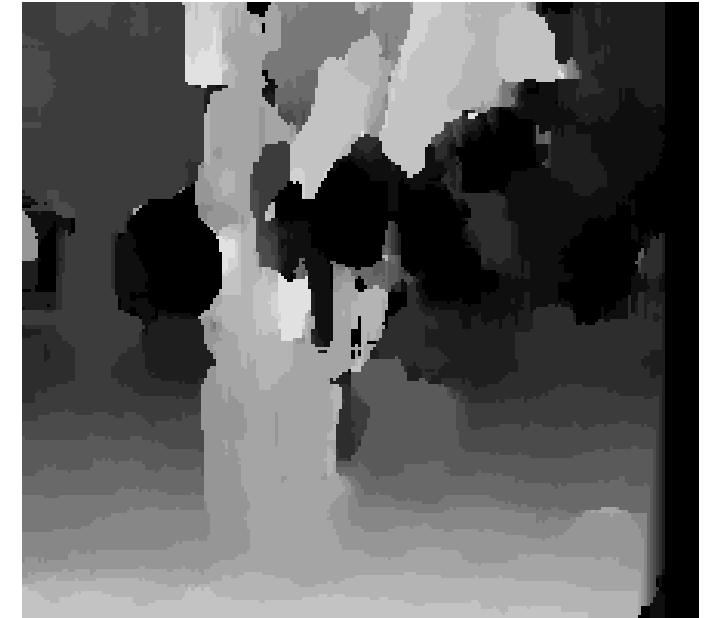
$$\text{ZNCC} = \sum_{(x,y)} \hat{f}(x,y) \hat{g}(x,y) \quad \text{Zero-mean Normalized Cross Correlation}$$

$$\hat{f} = \frac{f - \bar{f}}{\sqrt{\sum |f - \bar{f}|^2}} \quad \hat{g} = \frac{g - \bar{g}}{\sqrt{\sum |g - \bar{g}|^2}}$$

Drawback of Using a Window



$W = 3$



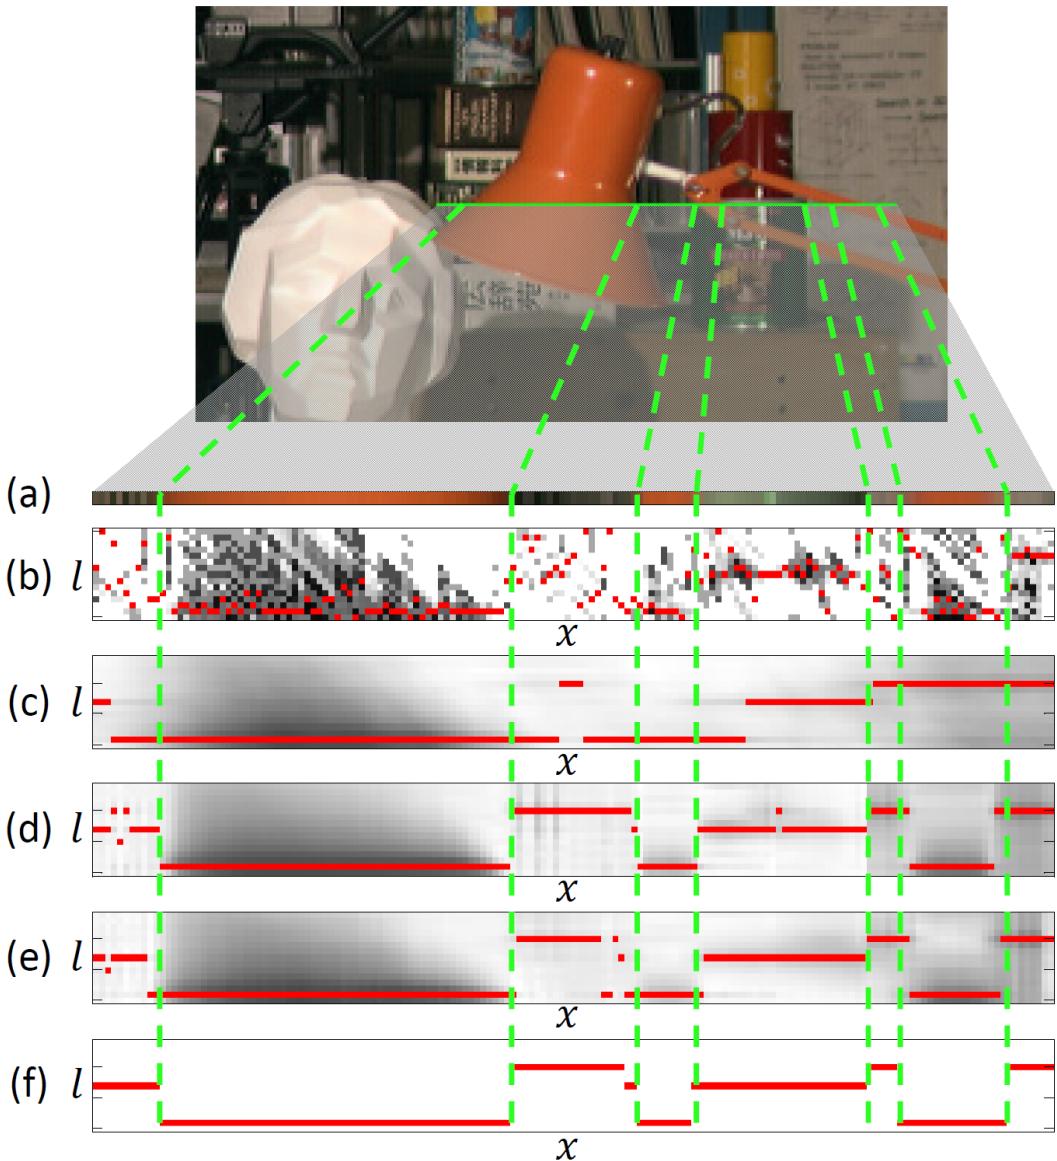
$W = 20$

- The entire window is assumed to have the same depth (planar and facing the camera)
- Want window large enough to have sufficient intensity variation, yet small enough to contain only pixels with about the same depth



Aggregation by Filtering

- Aggregating with SSD is equivalent to filter the DSI with a box filter
 - (b) the original DSI (red marks the disparity with minimum cost)
 - (c) the DSI filtered by a box filter
- Bilateral filters can generate better results
 - (d) the DSI filtered by a bilateral filter
- Guided filters are faster when large windows are used
 - (e) the DSI filtered by a Guided filter



Spatially Variant Kernels

- Bilateral filter and guided filters adaptively change kernel weights according to image content
 - Assuming pixels on the same object to have the same depth

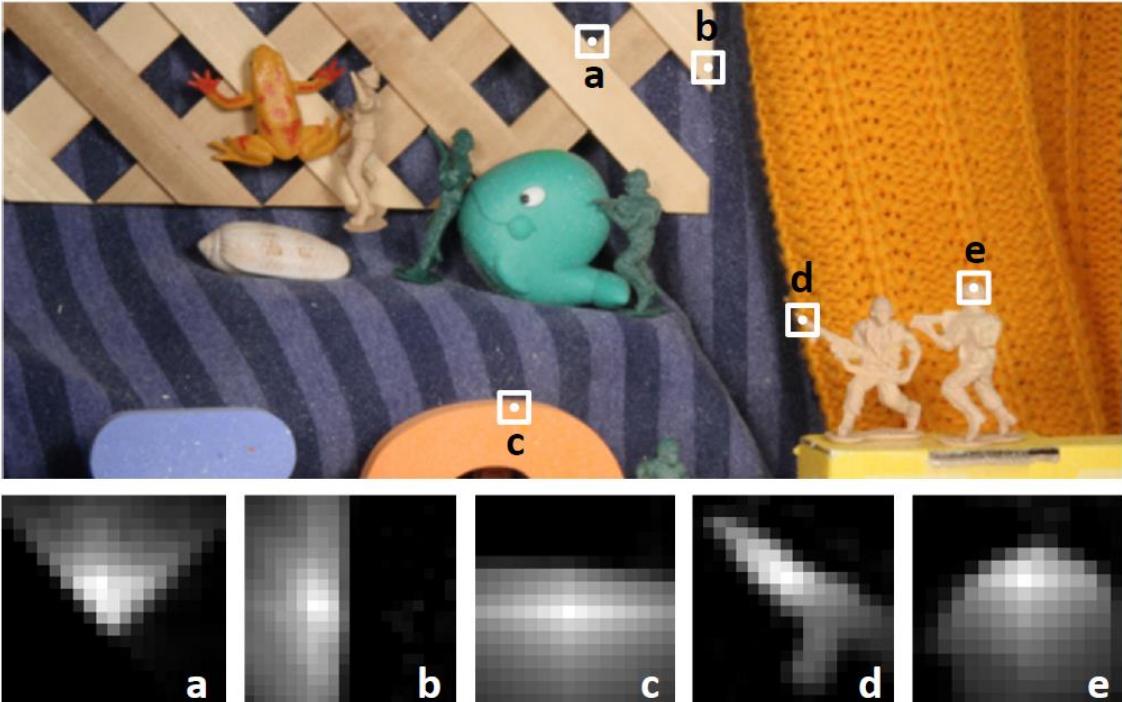


Figure 3. **Filter kernels.** We show kernels of the guided filter with $r = 9$ and $\epsilon = 0.01^2$, at different locations in an image of [1].

Local Methods



Adaptive Support-Weight Approach for Correspondence Search

Kuk-Jin Yoon, *Student Member, IEEE*, and
In So Kweon, *Member, IEEE*

[PAMI 2006]

Fast Cost-Volume Filtering for Visual Correspondence and Beyond

Christoph Rhemann¹, Asmaa Hosni¹, Michael Bleyer¹, Carsten Rother², Margrit Gelautz¹

¹Vienna University of Technology, Vienna, Austria ²Microsoft Research Cambridge, Cambridge, UK

[CVPR 2011]

Questions?



Global Methods



result by global methods

Boykov et al., [Fast Approximate Energy Minimization via Graph Cuts](#),
International Conference on Computer Vision, September 1999.

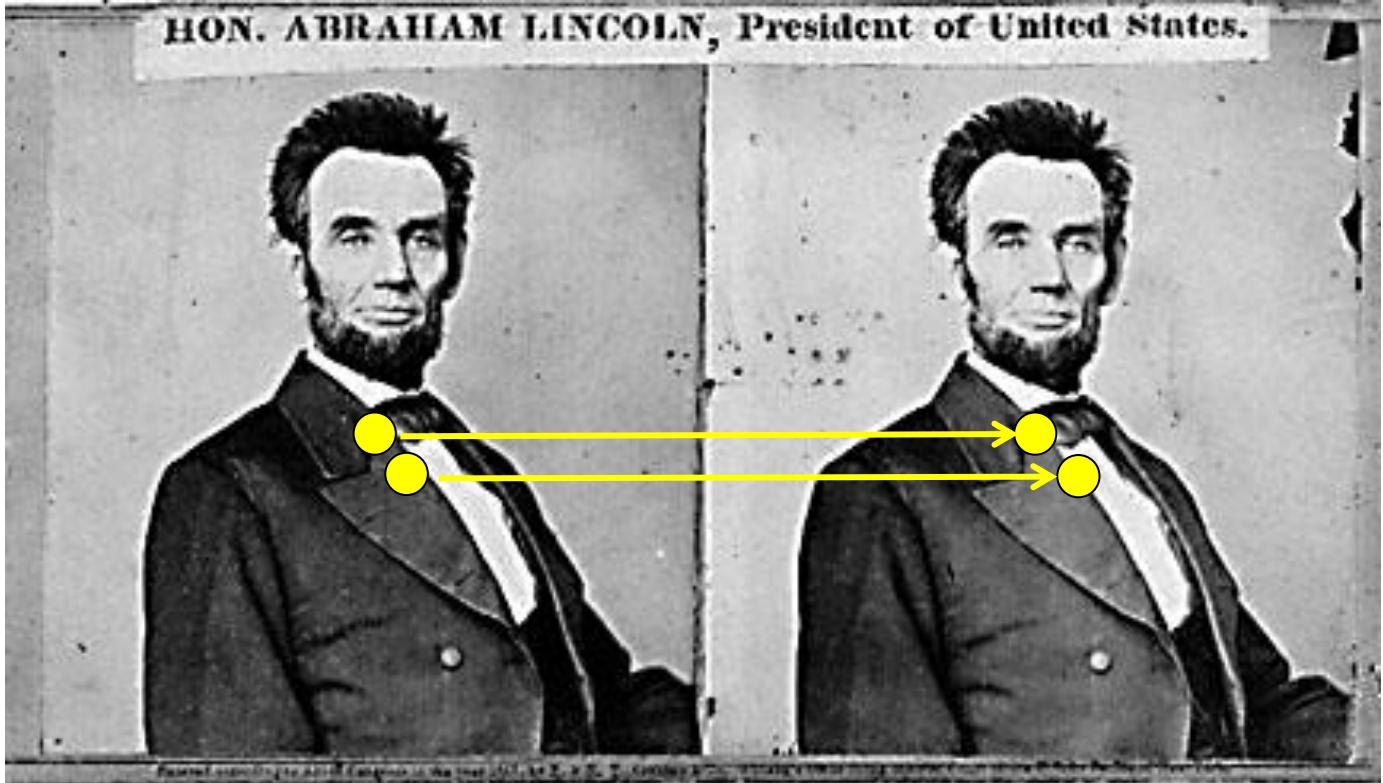


ground truth

the latest and greatest results at: <http://www.middlebury.edu/stereo>

Smoothness Assumption

- Textureless regions are ambiguous to match
- If two pixels are adjacent, their disparities should be similar





Stereo by Energy Minimization

- An objective function with data/match cost and smoothness cost

$$E(d) = \underbrace{E_d(d)}_{\text{match cost}} + \lambda \underbrace{E_s(d)}_{\text{smoothness cost}}$$

Want each pixel to find a good match in the other image

Adjacent pixels should (usually) have the same depth



Stereo by Energy Minimization

$$E(d) = E_d(d) + \lambda E_s(d)$$

match cost:

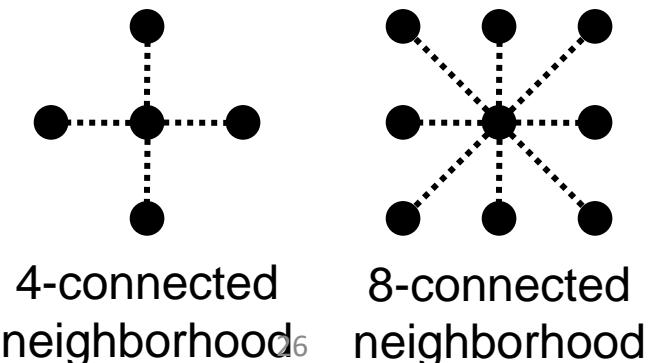
$$E_d(d) = \sum_{p \in I} D(p, d_p)$$

$D(p, d_p)$ the matching cost of the pixel p with disparity d_p

smoothness cost: $E_s(d) = \sum_{(p,q) \in E} S(p, q, d_p, d_q)$

$S(p, q, d_p, d_q)$ the smoothness cost when p, q have disparities of d_p, d_q respectively

E : set of neighboring pixels



Smoothness Cost

$$E_s(d) = \sum_{(p,q) \in E} S(p, q, d_p, d_q)$$

How do we choose S ?

1. Pixels of similar color have stronger smoothness

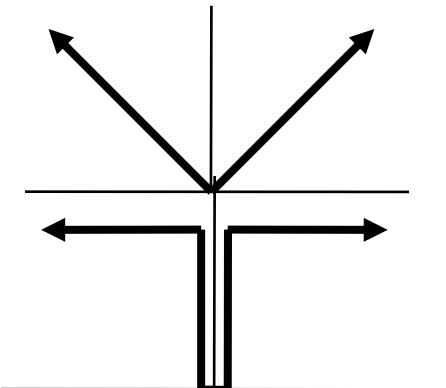
$$S \propto |\text{color}(p) - \text{color}(q)|^{-1}$$

2. Neighboring pixels have similar depth

$$S \propto |d_p - d_q|$$

or

$$S \propto \begin{cases} 0 & \text{if } d_p = d_q \\ 1 & \text{if } d_p \neq d_q \end{cases}$$



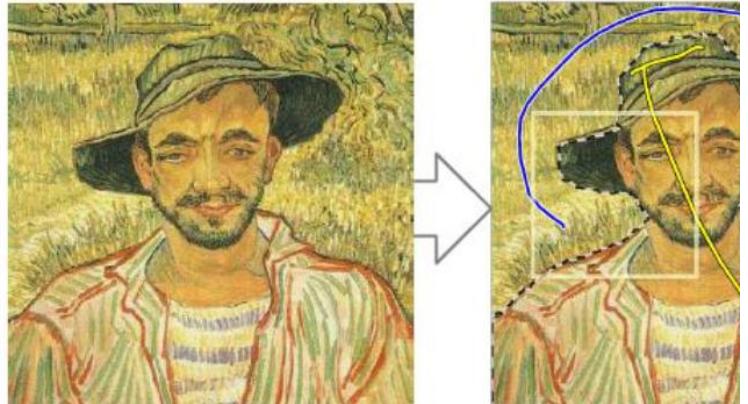
“Potts model”
The most common choice

$$S = |\text{color}(p) - \text{color}(q)|^{-1} \delta(d_p - d_q)$$

Graph-Cut Optimization



6. Interactive Segmentation & Graph-Cut



Binary Graph-Cut Optimization

- Minimize an objective function defined on a graph

$$E(X) = \sum_{p \in V} D_p(x_p) + \sum_{(p,q) \in E} S_{pq}(x_p, x_q)$$

$$X = \{x_1, x_2, \dots, x_N\}, x_p = \{0,1\},$$

V, E are the set of vertices and edges of a graph

$D_p(\cdot)$ and $S_{pq}(\cdot)$ are functions defined on vertices and edges

- We begin with the binary problem with Potts model:

$$S_{pq}(x_p, x_q) = w_{pq} \delta(|x_p - x_q|) = \begin{cases} w_{pq} & \text{if } x_p \neq x_q \\ 0 & \text{otherwise} \end{cases}$$

- It is possible to define edge weights to solve the minimization by min-cut



Graph-Cut Optimization

Binary Graph-Cut Optimization

- More general result is provided in the following paper

What Energy Functions Can Be Minimized
via Graph Cuts'

Vladimir Kolmogorov, Member, IEEE, and Ramin Z

- Given a binary function, $x_p \in \{0, 1\}$

$$E(X) = \sum_{p \in V} D_p(x_p) + \sum_{(p,q)}$$

Graph-cut can find the GLOBAL minimum
 $S_{pq}(0,0) + S_{pq}(1,1) < S_{pq}(0,1) + S_{pq}(1,0)$



What if x_p is not binary?

- What if x_p is not binary, e.g. $x_p \in \{1, 2, \dots, n\}$?
- The basic idea: convert this problem to a binary one
- Start from an initial configuration, and iteratively improve the result
 - Two possible solutions:
 - Alpha-expansion and Alpha-beta swap
 - Both methods improve the result by solving a binary graph-cut problem each iteration
 - Converge to a LOCAL minimum





Global Methods

Results by Graph-Cut Optimization



result by global methods

Boykov et al., [Fast Approximate Energy Minimization via Graph Cuts](#),
International Conference on Computer Vision, September 1999.



ground truth

the latest and greatest results at: <http://www.middlebury.edu/stereo>

Questions?





Stereo by Dynamic Programming

$$E(d) = E_d(d) + \lambda E_s(d)$$

- The optimization problem is simplified by considering one row at a time
- Global optimal can be achieved by the *Dynamic Programming* algorithm (for each row)



Stereo by Dynamic Programming

- For each row of pixels, we need to find a “smooth” path through DS_I from left to right that minimizes the cost

$$E(d) = E_d(d) + \lambda E_s(d)$$





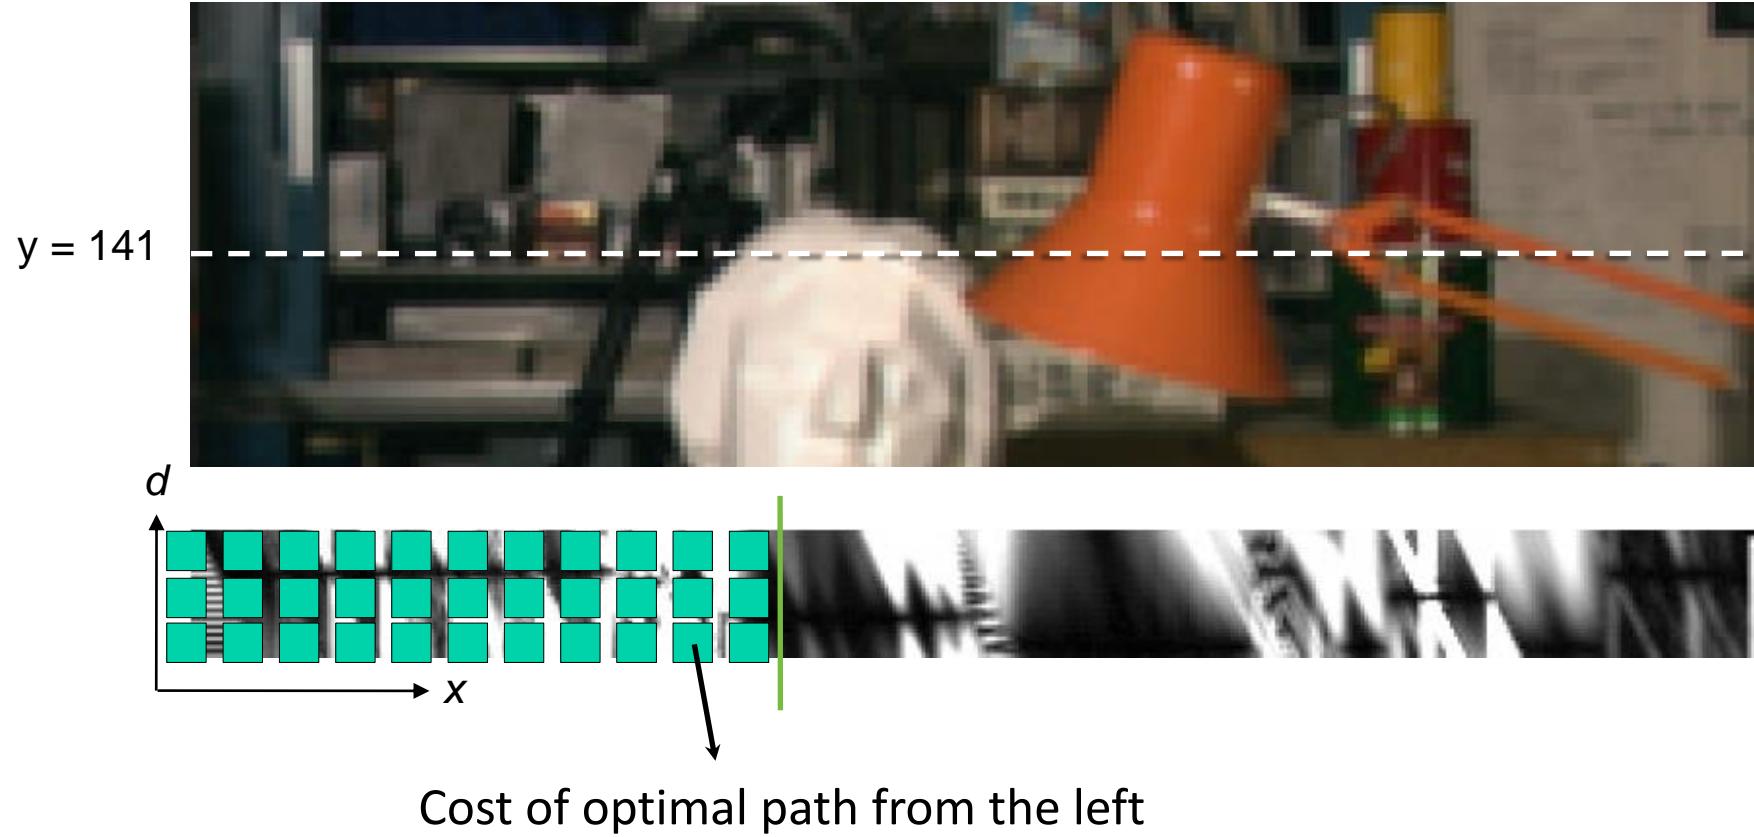
Dynamic Programming

- We define $D(x, y_0, d)$ as the minimum cost among all paths that
 - Start from $(0, y_0)$, i.e. the left border of the image
 - End at (x, y_0) with disparity d
- Dynamic programming computes D efficiently by:

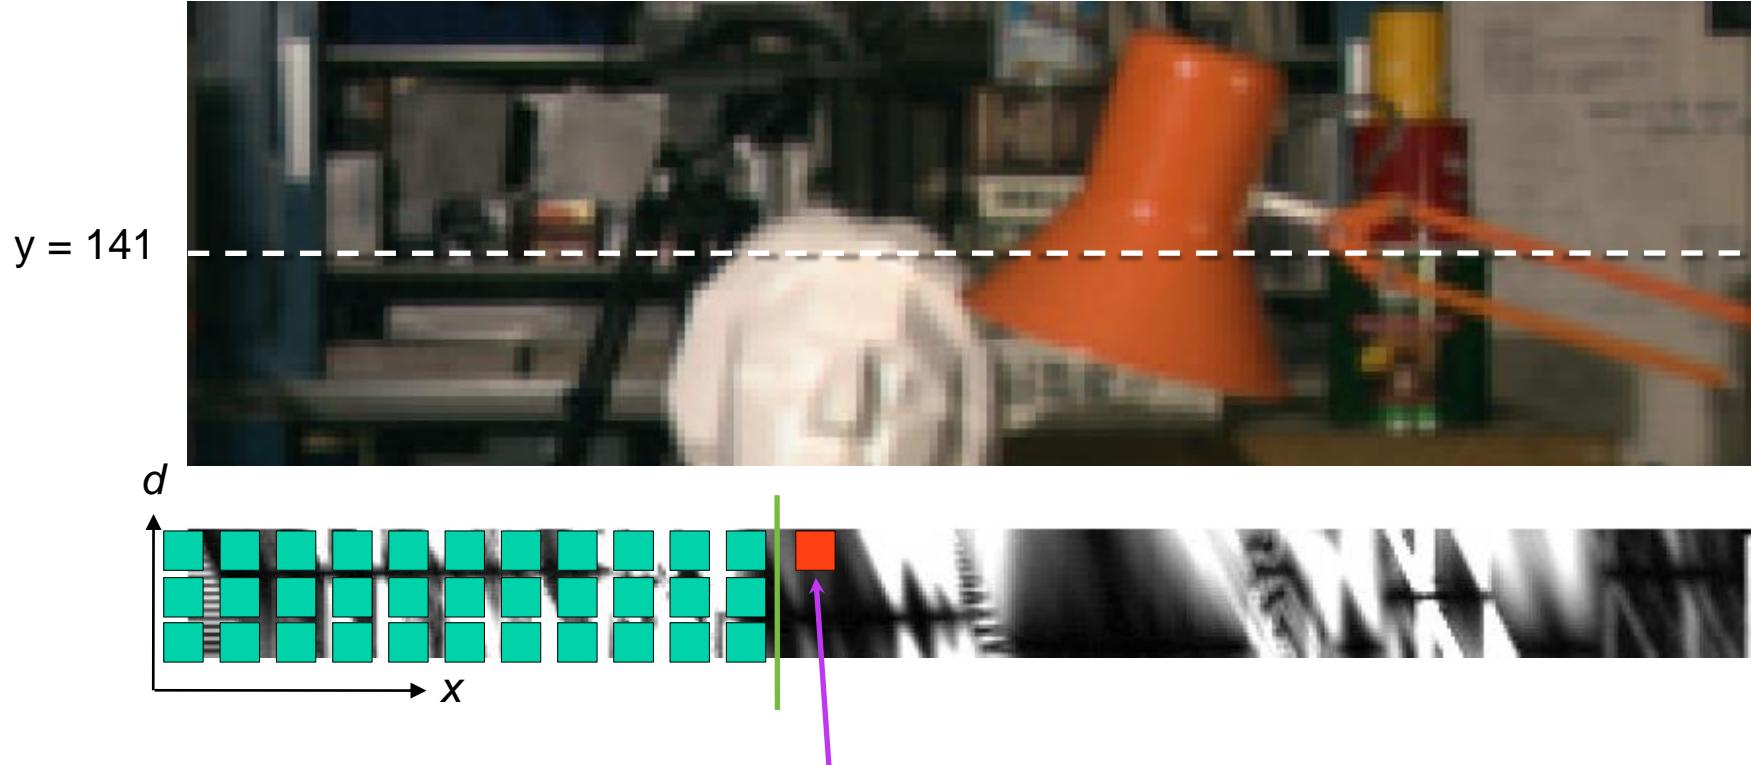
$$D(x, y_0, d) = \underbrace{C(x, y_0, d)}_{\text{Data cost at } (x, y_0)} + \underbrace{\min_{d'} \{ D(x - 1, y_0, d') + \lambda |d - d'| \}}_{\text{smooth cost between } (x - 1, y_0) \text{ and } (x, y_0)}$$

- This strategy allow us to reuse computation at previous pixels

Dynamic Programming

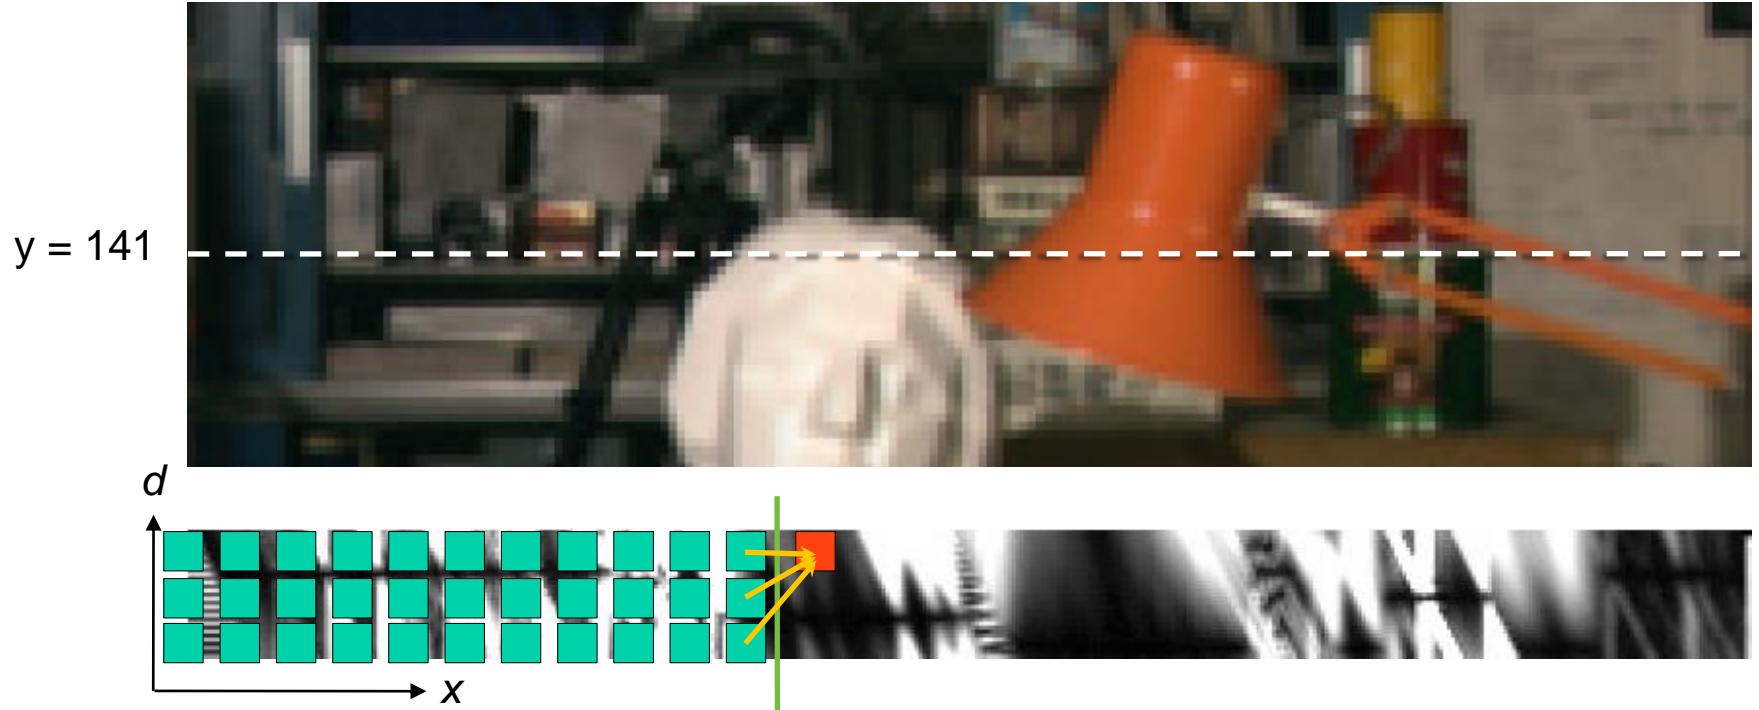


Dynamic Programming



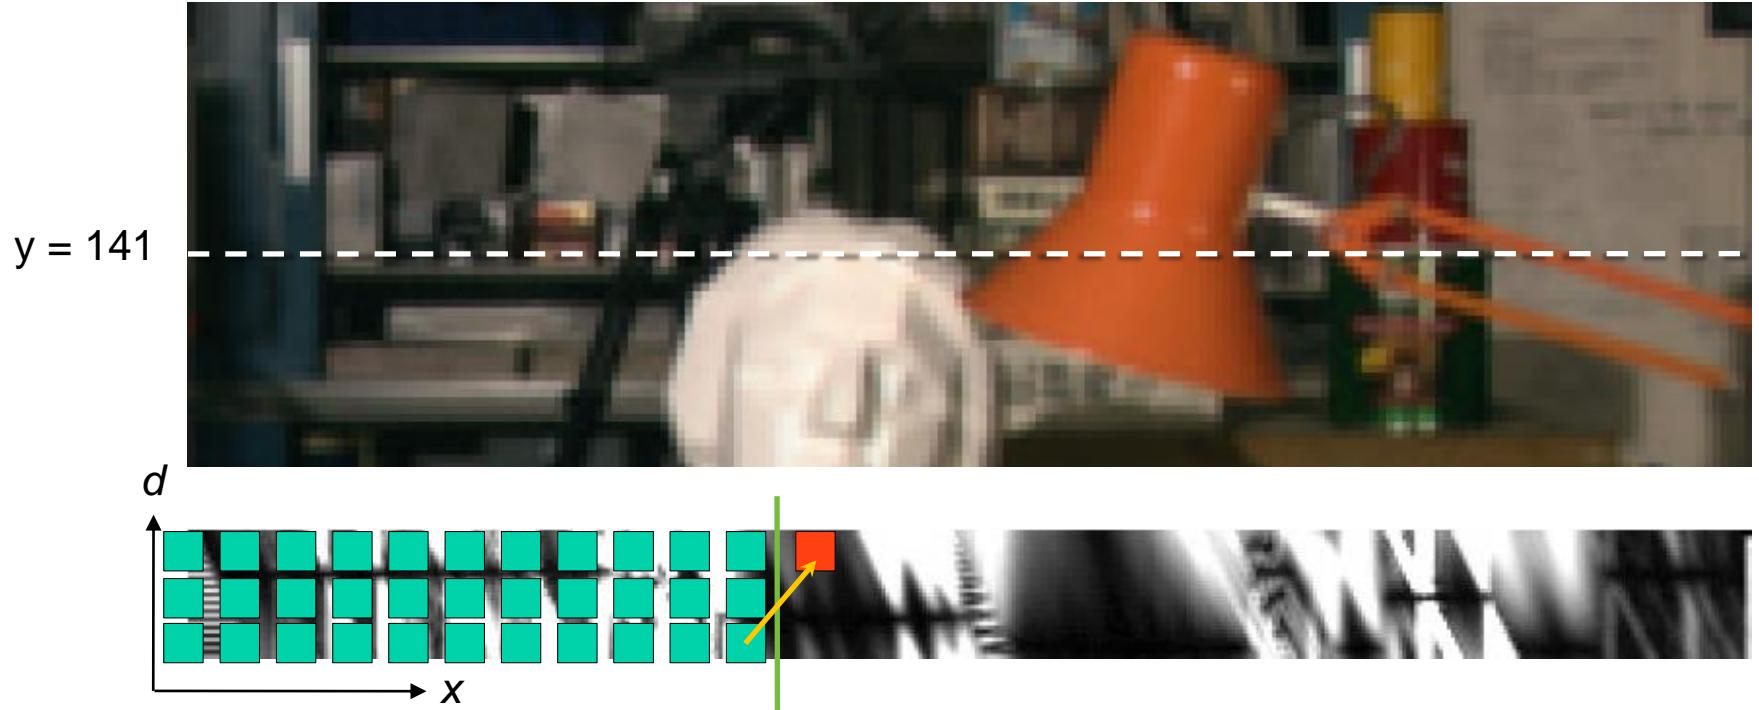
How to compute for this pixel?

Dynamic Programming



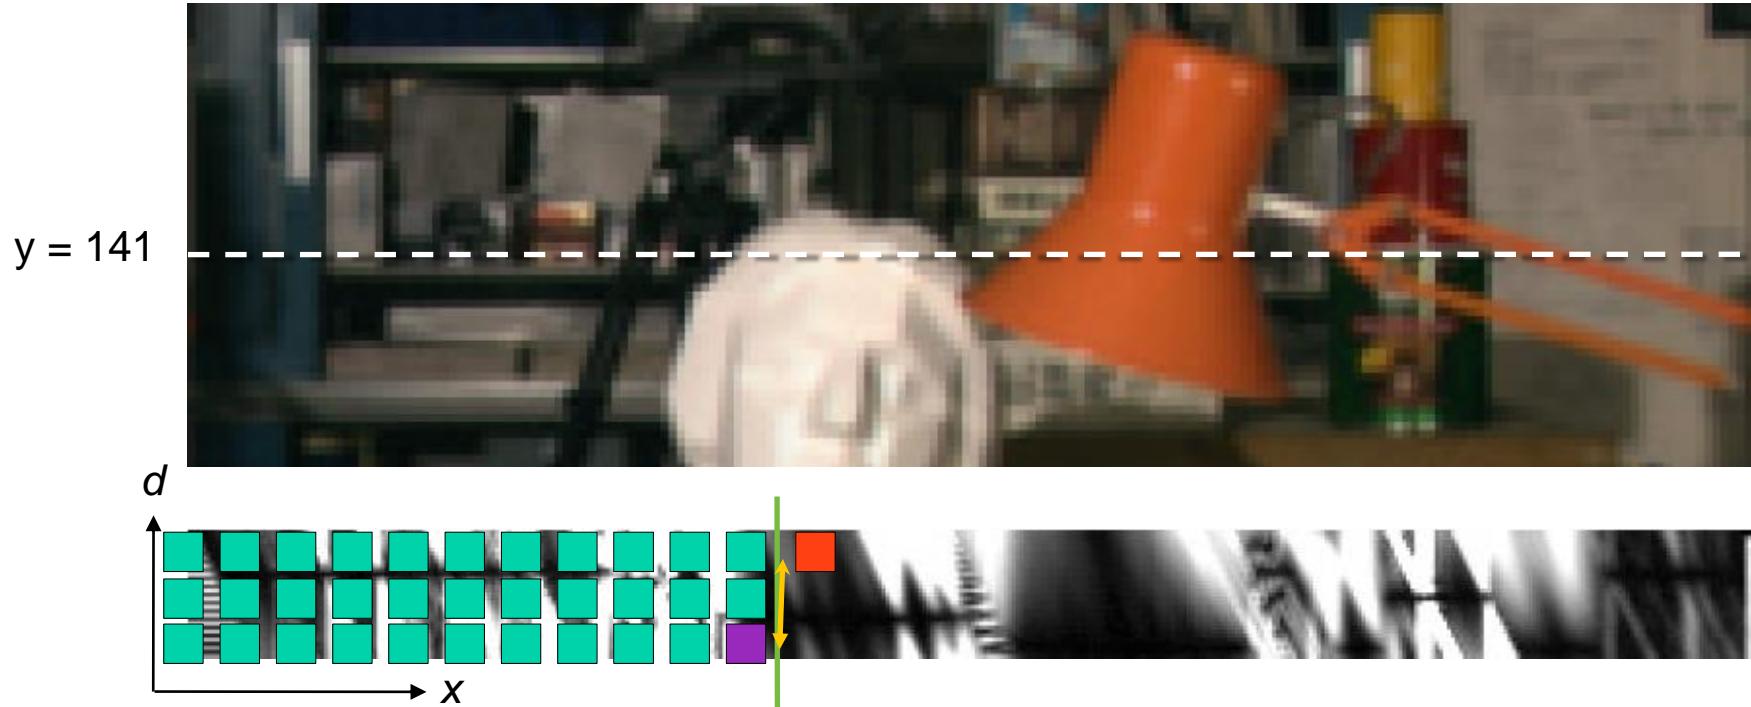
How to compute for this pixel?
Only 3 ways to get here...

Dynamic Programming



This cost is ...

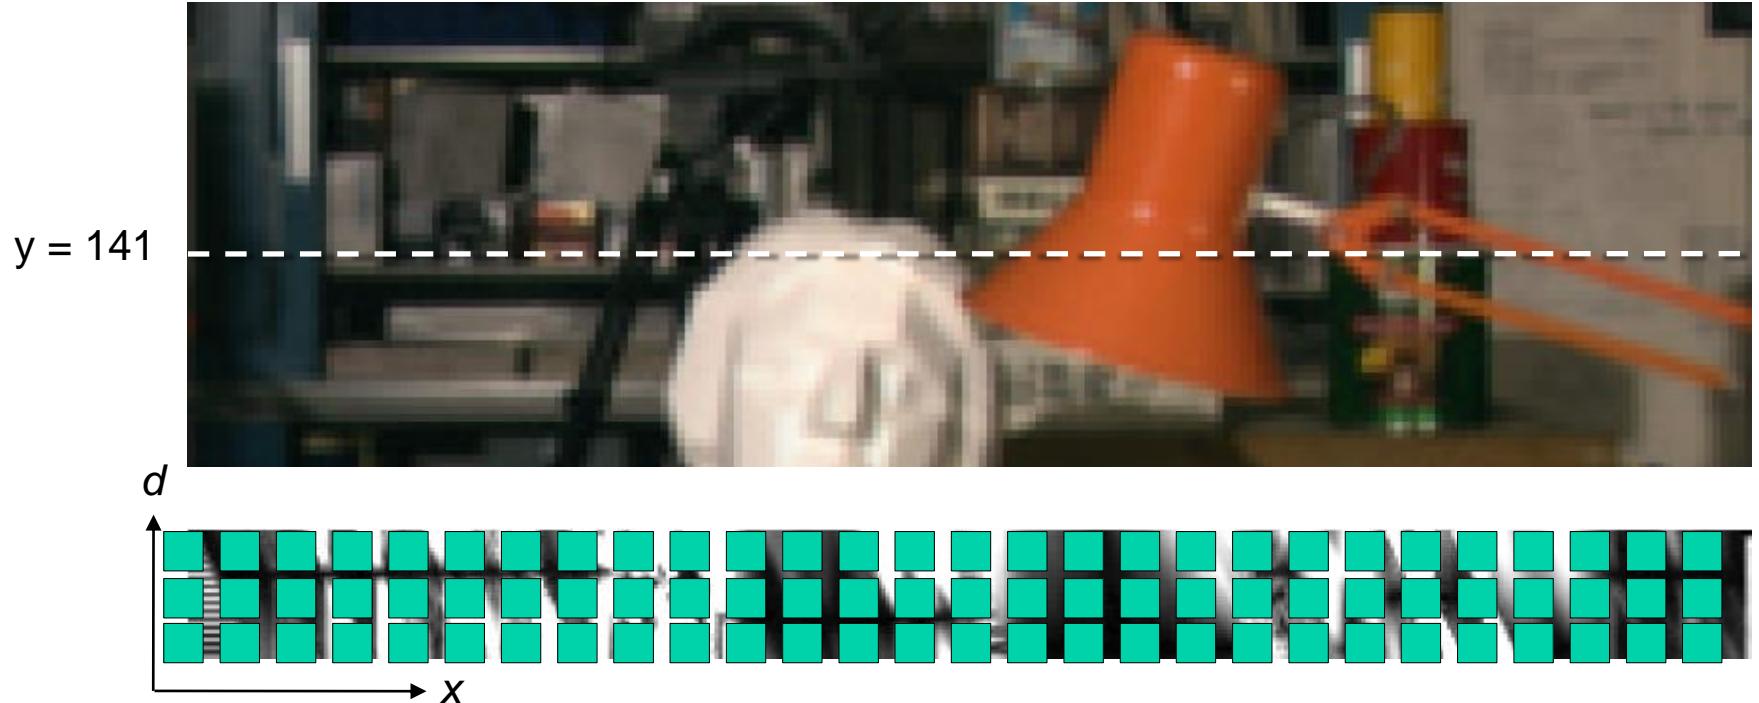
Dynamic Programming



This cost is ...

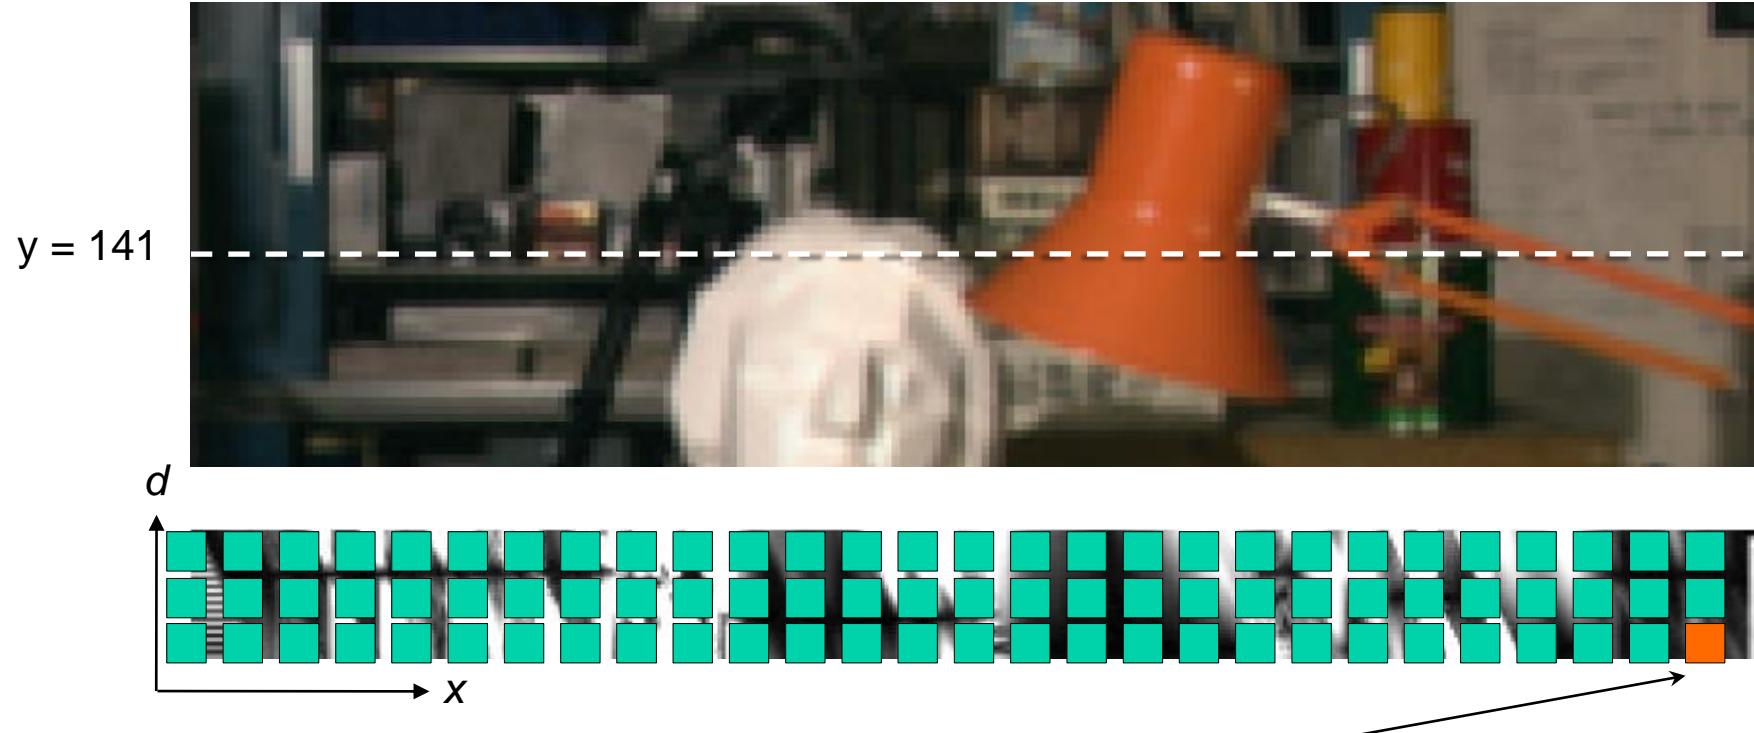
$$D + Ed + Es$$

Dynamic Programming



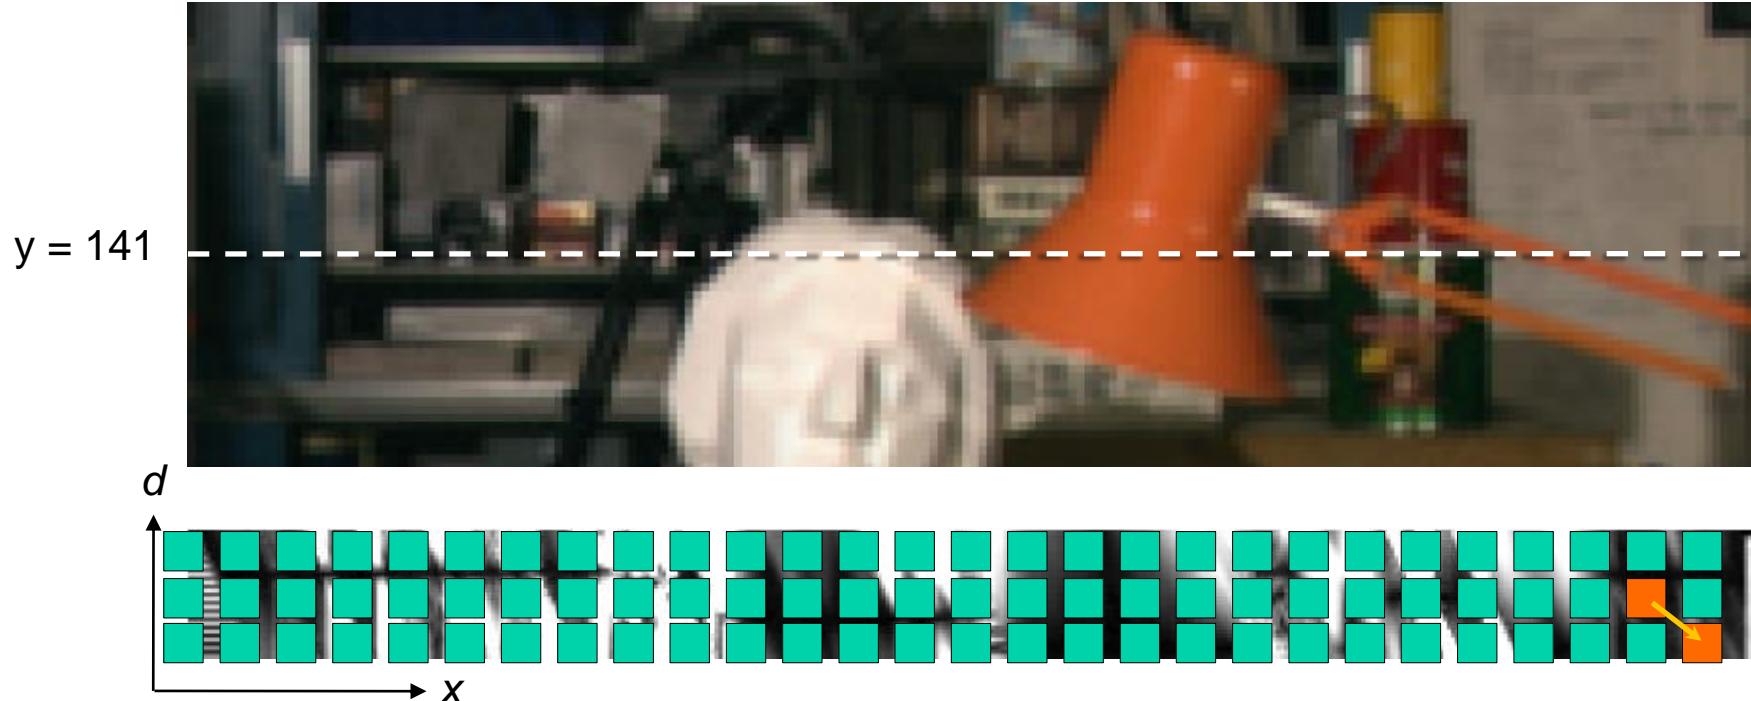
Fill D all the way to the right

Dynamic Programming



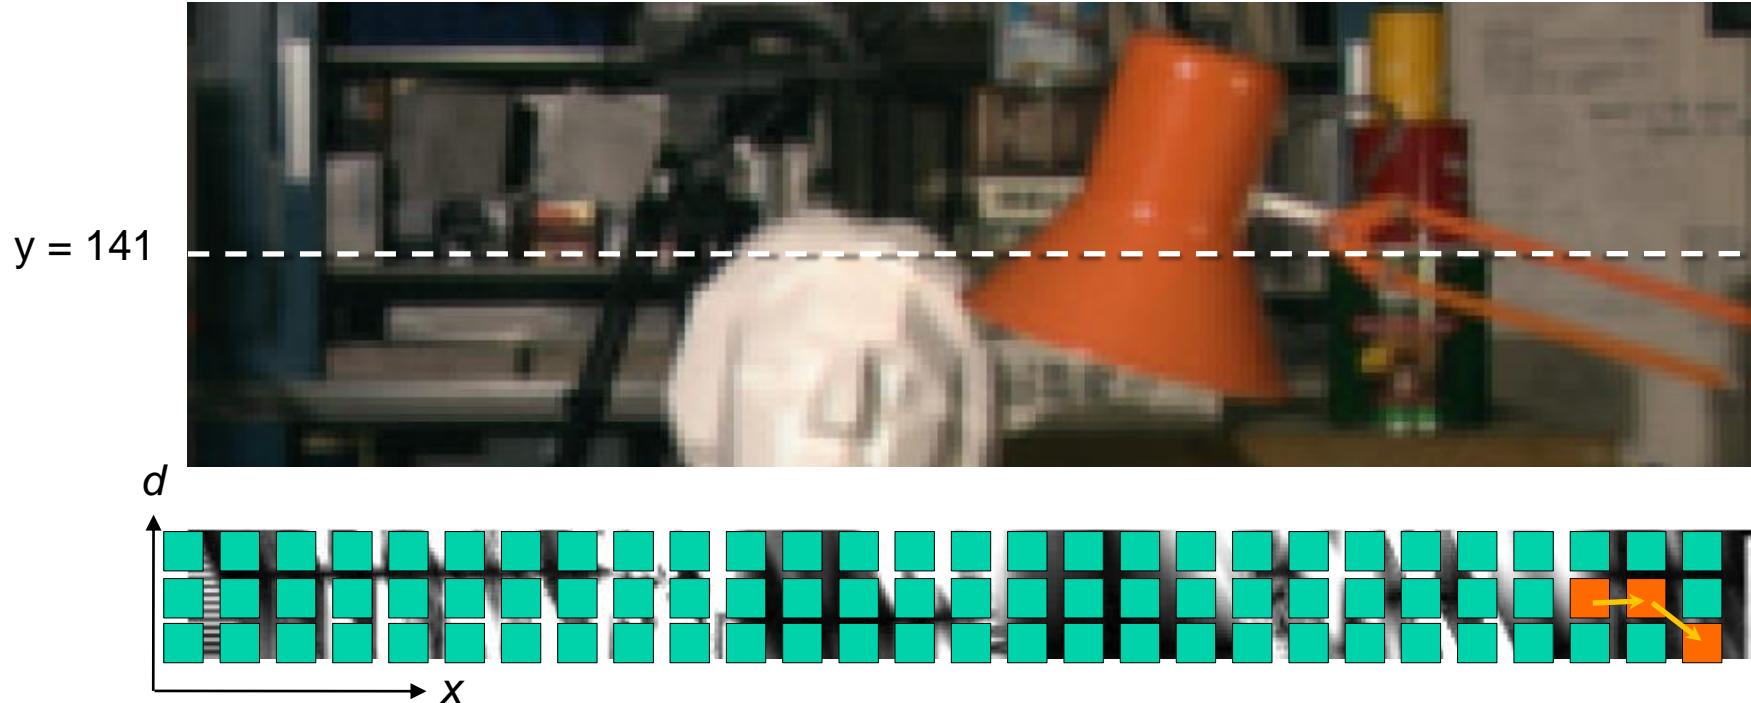
Pick the best one from the right column

Dynamic Programming



Back track....

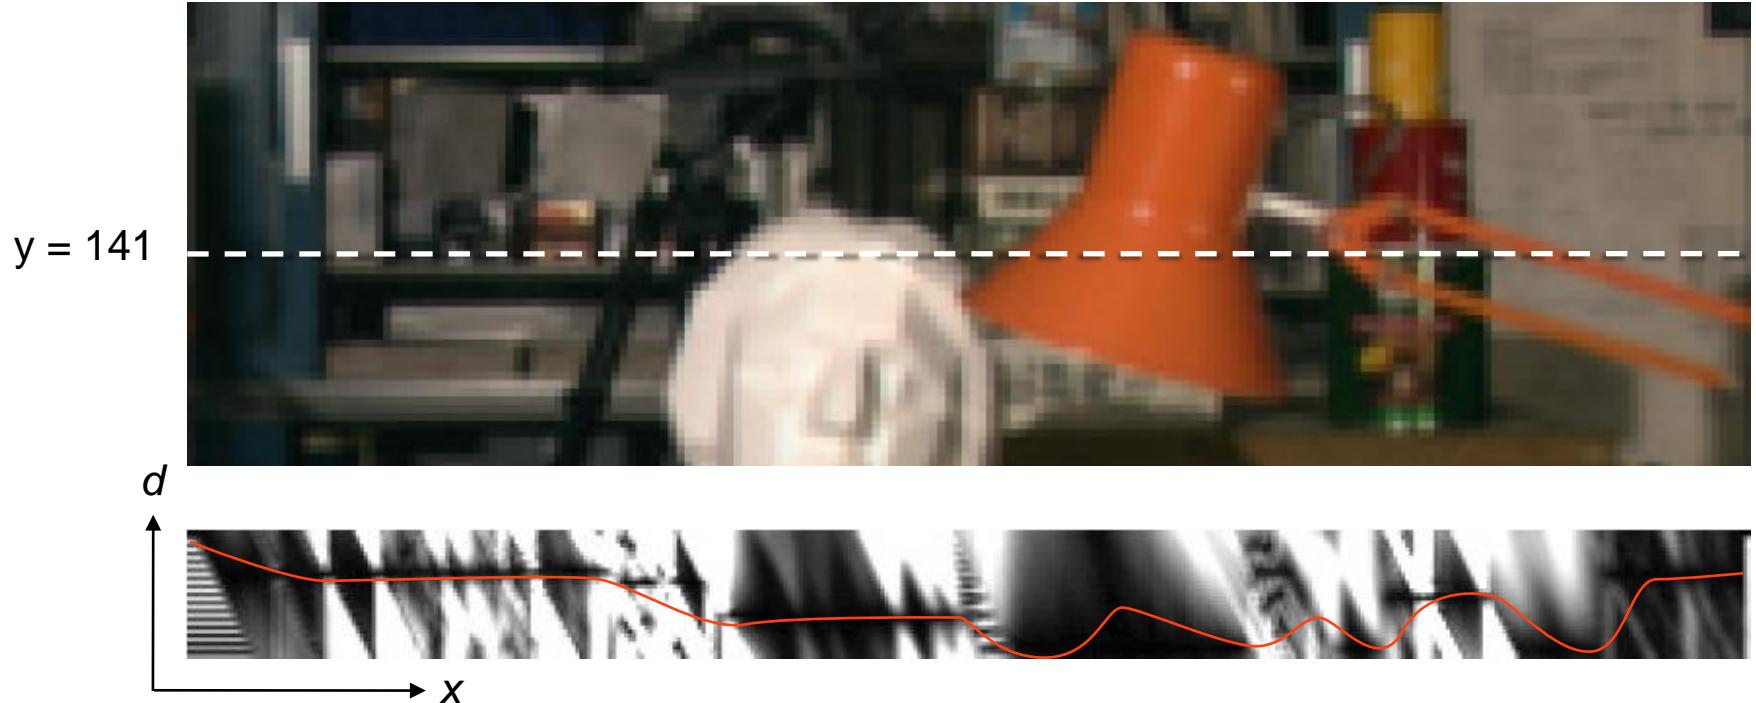
Dynamic Programming



Back track....

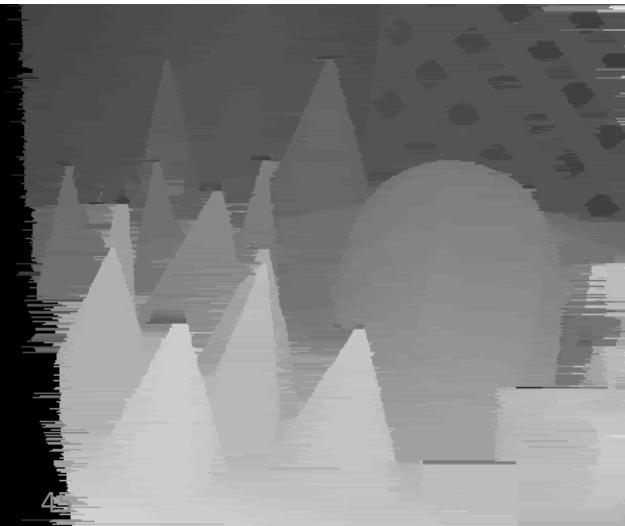
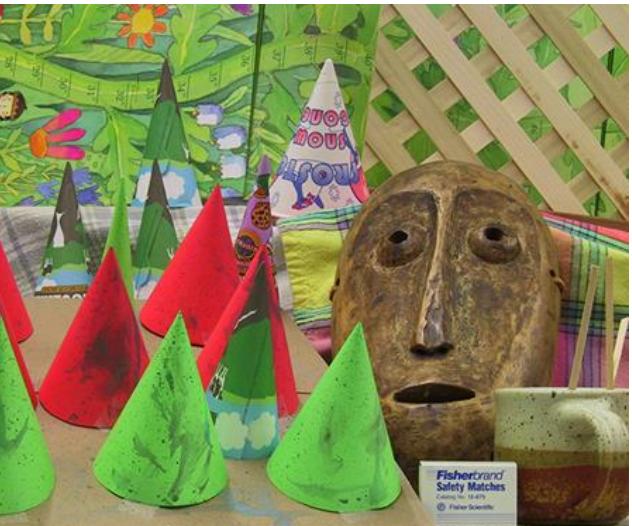
Dynamic Programming

- Final result at one row



Dynamic Programming Results

- Suffers from streak errors





Semi-Global Method

$$E(d) = E_d(d) + \lambda E_s(d)$$

- Aim to minimize this global energy function
- Between local and global optimization
- Utilize dynamic programming for *cost aggregation*
- Faster than global method, with similar accuracy

Stereo Processing by Semiglobal Matching and Mutual Information

Heiko Hirschmüller

[PAMI 2008]

Semi-Global Method

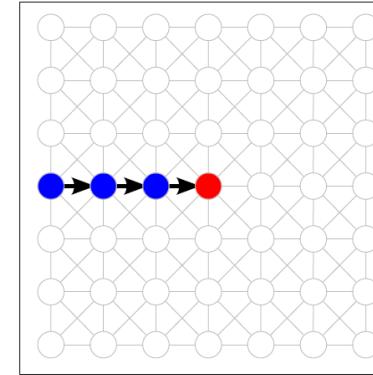
- Recall the function $D(x, y_0, d)$ defined in Page 34
- It aggregates cost from left to right
- We could aggregate cost from other directions
 - From right to left, top -> down, bottom -> up, etc
- For each direction i , define a function $D_i(x_0, y_0, d)$
 - That is the minimum cost among all paths that along the i -th direction
 - Start from the image border
 - End at the pixel (x_0, y_0) with disparity d
- D_i can be evaluated by DP as well

Dynamic Programming

- We define $D(x, y_0, d)$ as the minimum cost of all paths that
 - Start from $(0, y_0)$, i.e. the left border of the image
 - End at (x, y_0) with disparity d
- Dynamic programming computes D efficiently by:

$$D(x, y_0, d) = C(x, y_0, d) + \min_{d'} \{ D(x-1, y_0, d') + \lambda |d - d'| \}$$

- This strategy allow us to reuse computation at previous pixels

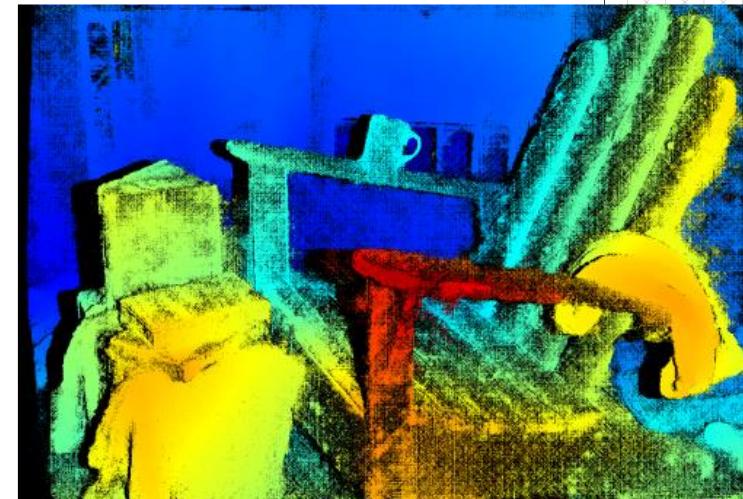
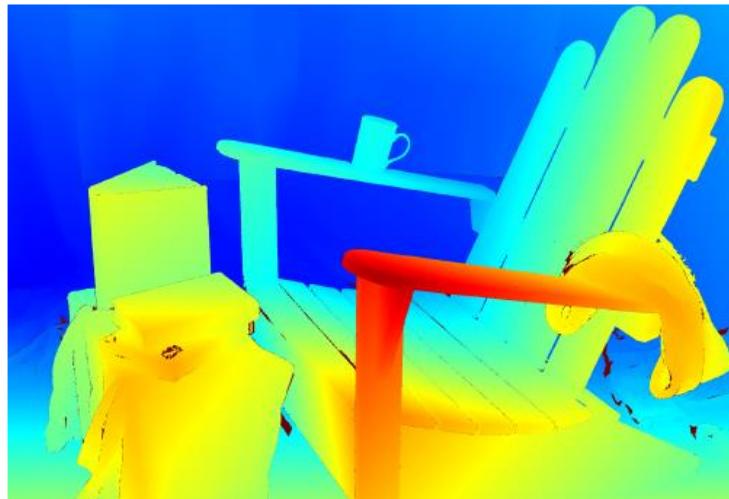


Semi-Global Method

- With D_i computed for all directions, the final cost is

$$S(x_0, y_0, d) = \sum_i D_i(x_0, y_0, d)$$

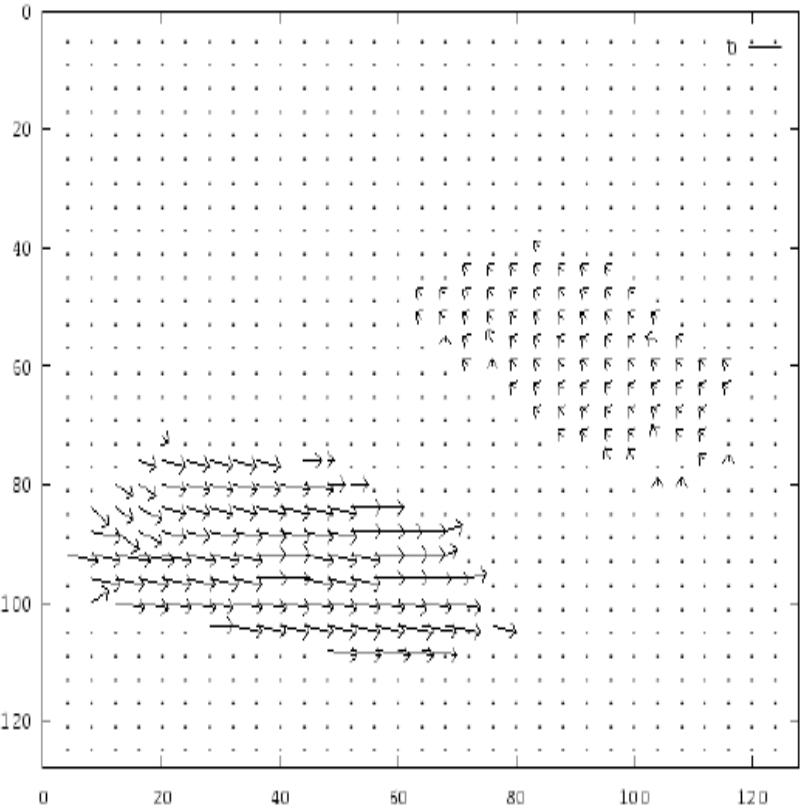
- Finally, apply Winter-Take-All at each pixel to choose the optimal d



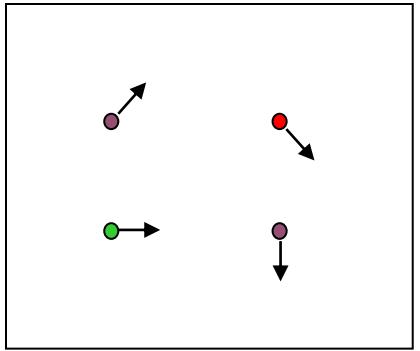
Questions?



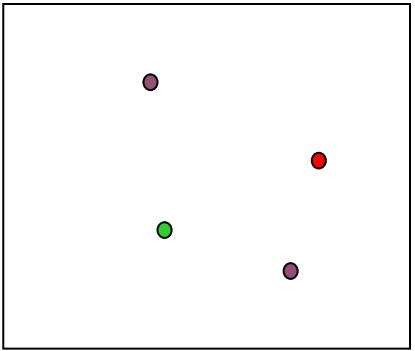
Optical Flow



Problem Definition: optical flow



$H(x, y)$

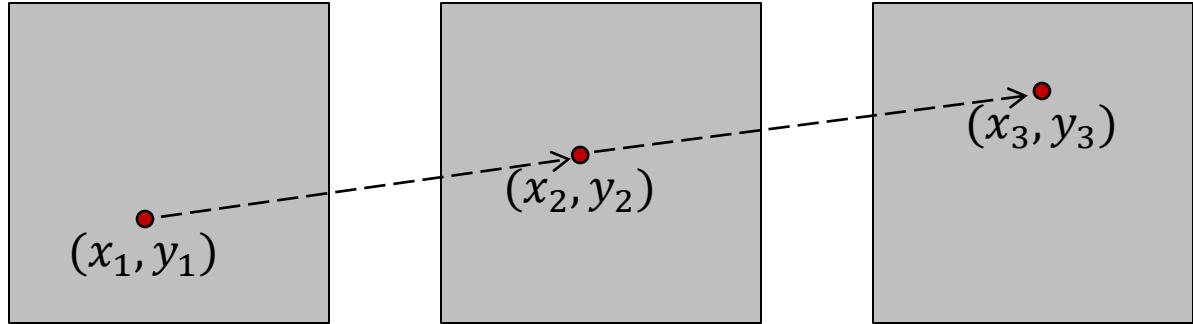


$I(x, y)$

- How to estimate pixel motion from image H to image I ?
 - Solve pixel correspondence problem
 - given a pixel in H , look for nearby pixels of the same color in I
 - Difference from stereo: no epipolar lines, dynamic scenes
 - Key assumptions
 - **color constancy**: a point in H looks the same in I
 - For grayscale images, this is brightness constancy
 - **small motion**: points do not move very far

Brightness Constancy

- A scene point moves through multiple frames
- The brightness of the point remains the same



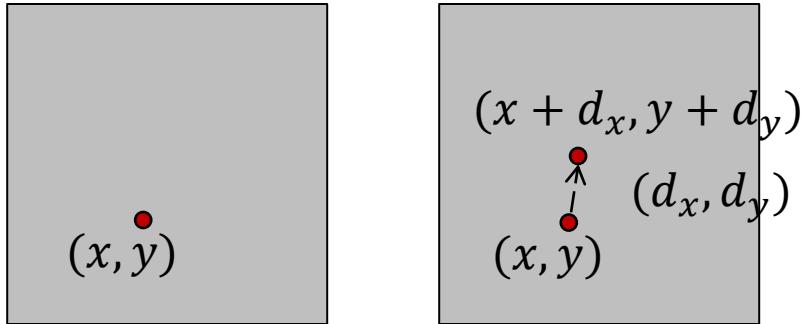
$$I(x(t), y(t), t) = c$$

constant



Small Motion

$$I(x, y, t) = I(x + d_x, y + d_y, t + d_t)$$



- Assume the motion is small

- Taylor expansion:

$$I(x, y, t) = I(x, y, t) + I_x d_x + I_y d_y + I_t d_t$$

$$\rightarrow I_x d_x + I_y d_y + I_t d_t = 0$$



The Optical Flow Equation

$$I_x d_x + I_y d_y + I_t d_t = 0$$

- Divide by d_t

$$I_x \frac{\bar{d}_x}{\bar{d}_t} + I_y \frac{\bar{d}_y}{\bar{d}_t} + I_t = 0$$

u v

(u, v) is the velocity

- The optical flow equation:

$$I_x u + I_y v + I_t = 0$$

- Derived from brightness constancy and small motion



The Optical Flow Equation

$$I_x u + I_y v + I_t = 0$$

- I_x, I_y are the image gradients, known
 - Computed by difference, e.g. Sobel filter, Scharr filter, etc
- I_t is the temporal gradient, known
 - Computed by frame difference
- u, v are the velocity of a pixel, unknown
 - Solve them by optical flow algorithm



Aperture Problem

The optical flow equation:

$$I_x u + I_y v + I_t = 0$$

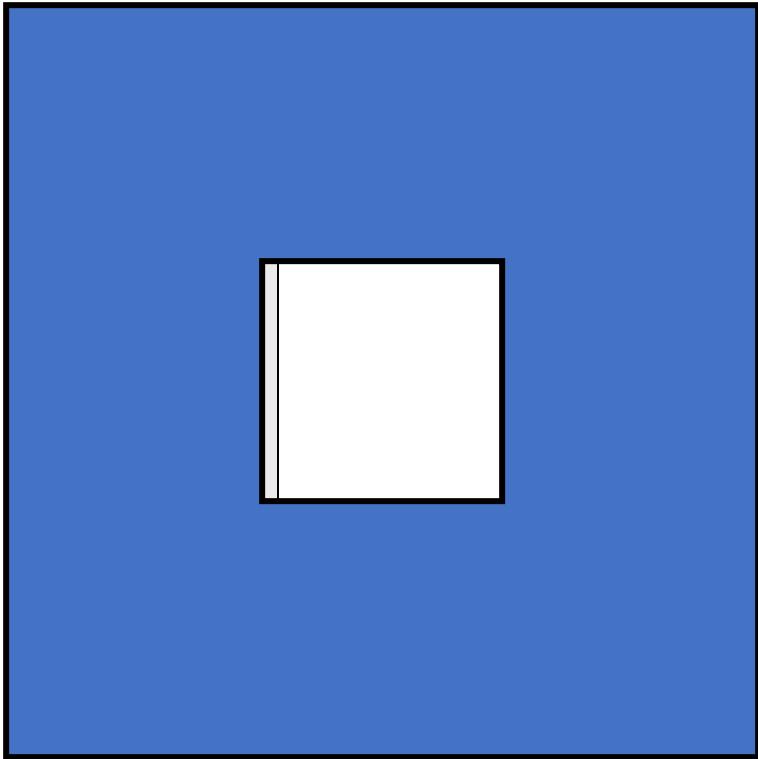
At each pixel, we have:

- One equation
- Two unknowns: u, v
- Multiple solutions
(because of less equation than unknowns)



Aperture Problem

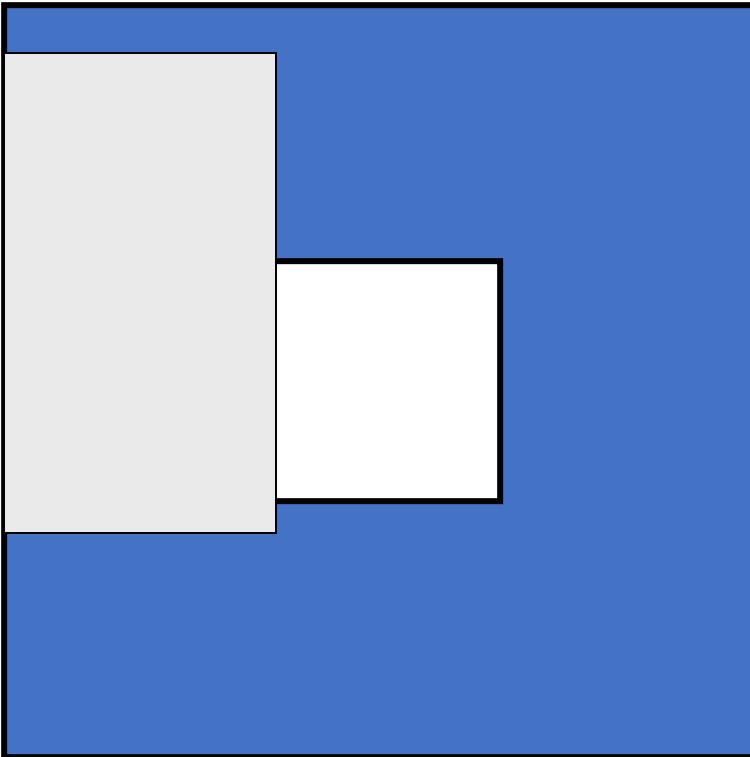
Example I





Aperture Problem

Example I





Aperture Problem

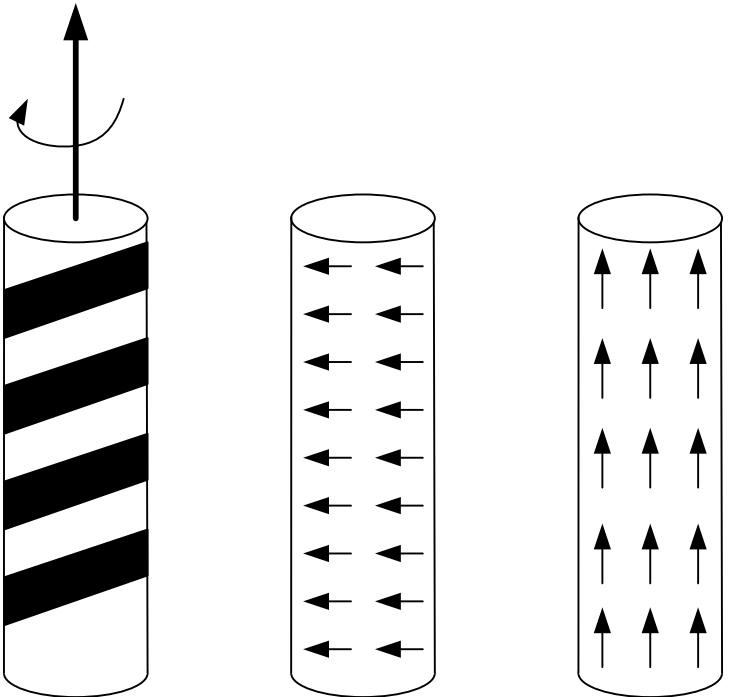
Example II



barber's pole

Aperture Problem

Example II



barber's pole motion field optical flow

Aperture Problem

$$I_x u + I_y v + I_t = 0$$

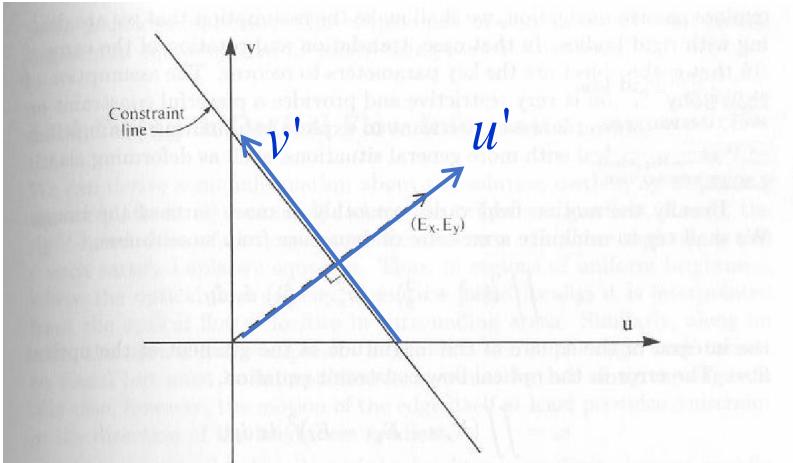


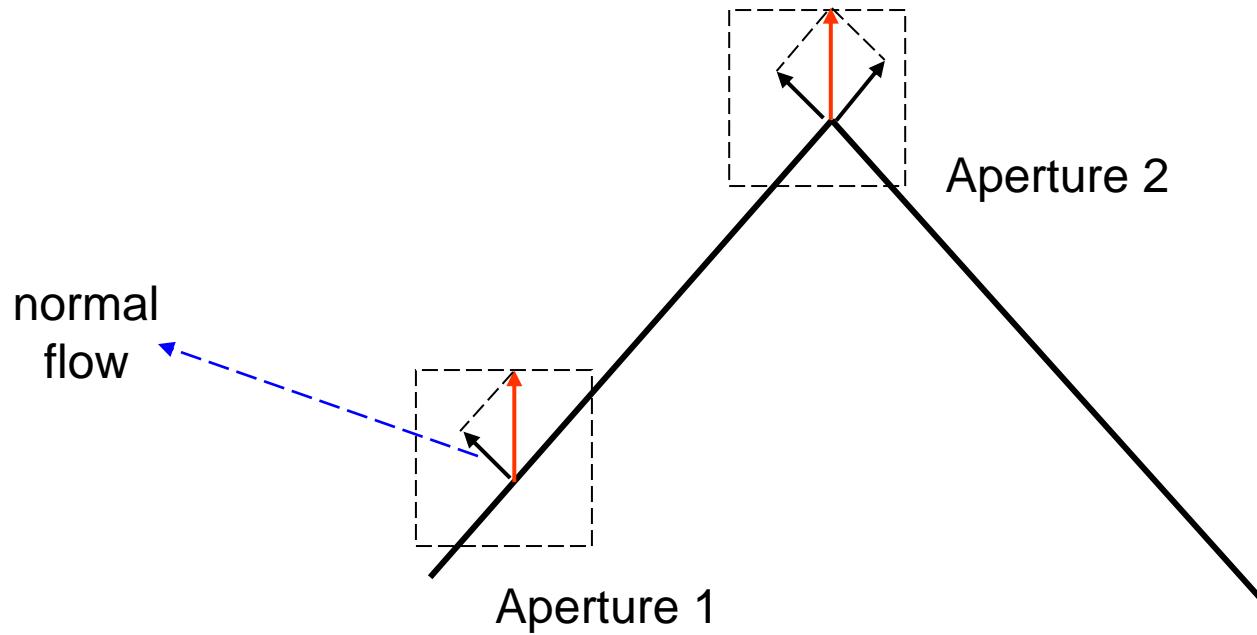
Figure 12-4. Local information on the brightness gradient and the rate of change of brightness with time provides only one constraint on the components of the optical flow vector. The flow velocity has to lie along a straight line perpendicular to the direction of the brightness gradient. We can only determine the component in the direction of the brightness gradient. Nothing is known about the flow component in the direction at right angles.

- The optical flow equation is a line equation in the u - v plane
- The velocity is confined on this line
- This line goes perpendicular to image gradient (I_x, I_y)
- Motion along image gradient can be determined (normal flow)
- Motion along the line direction cannot be determined

[Horn & Schunck 1981]

Aperture Problem

Ambiguity solved at corner



Questions?





Lucas & Kanade Method

Local continuous constraint –

- Each pixel gives a optical flow equation

$$I_x(i, j)u(i, j) + I_y(i, j)v(i, j) + I_t(i, j) = 0$$

- Two unknowns for each equation (aperture problem)
- Assume the motion is the constant locally to get more equations
- How many equations do we get from a 5x5 block?

$$I_x(p_1)u + I_y(p_1)v + I_t(p_1) = 0;$$

$$I_x(p_2)u + I_y(p_2)v + I_t(p_2) = 0;$$

⋮

$$I_x(p_{25})u + I_y(p_{25})v + I_t(p_{25}) = 0;$$



Lucas & Kanade Method

Matrix form

- Using a 5x5 patch, gives us 25 equations

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

$$A_{25 \times 2} \quad x_{2 \times 1} \quad b_{25 \times 1}$$



Lucas & Kanade Method

To obtain the solution, we need to solve:

$$A^T A \quad x \quad A^T b$$
$$\begin{bmatrix} \sum_p I_x I_x & \sum_p I_x I_y \\ \sum_p I_x I_y & \sum_p I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \sum_p I_x I_t \\ \sum_p I_y I_t \end{bmatrix}$$

which gives:

$$x = (A^T A)^{-1} A^T b$$



Lucas & Kanade Method

- When is $A^T A$ invertible?
 - λ_1, λ_2 should be large
- Where have you seen $A^T A$ before?
 - Harris corner detector
 - Corners are when λ_1, λ_2 are big; this is also when Lucas-Kanade optical flow works best
 - Corners are regions with two different directions of gradient (at least)
 - Corners are good places to compute flow!

$$A^T A = \begin{bmatrix} \sum_p I_x I_x & \sum_p I_x I_y \\ \sum_p I_x I_y & \sum_p I_y I_y \end{bmatrix}$$



Horn-Schunck vs Lucas-Kanade

- Lucas-Kanade assume locally constant flow
- Horn-Schunk assume globally smooth flow
 - The flow vectors at neighboring pixels should be similar
 - Just like the smoothness term on disparity at neighboring pixels

**Horn-Schunck
Optical Flow (1981)**

brightness constancy

small motion

'smooth' flow
(flow can vary from pixel to pixel)

global method
(dense)

**Lucas-Kanade
Optical Flow (1981)**

method of differences

'constant' flow
(flow is constant for all pixels)

local method
(sparse)



Horn & Schunck Method

- Two key constraints:
 - Optical flow equation (from brightness constancy)
 - Global smoothness
- Optical flow equation: $I_x u + I_y v + I_t = 0$
- Solving these equations at each pixel equivalents to minimize:

$$E_d = \sum_{i,j} [I_x(i,j)u(i,j) + I_y(i,j)v(i,j) + I_t(i,j)]^2$$

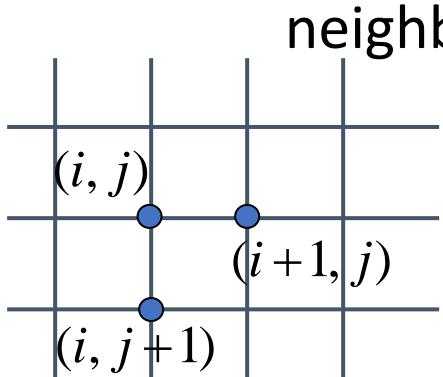
- Sometimes, written in integration form:

$$E_d = \iint [I_x(x,y)u(x,y) + I_y(x,y)v(x,y) + I_t(x,y)]^2 dx dy$$



Horn & Schunck Method

- Global smoothness constraint –



neighboring pixels have similar motion

$$\begin{aligned} u(i, j) &\approx u(i + 1, j); & v(i, j) &\approx v(i + 1, j); \\ u(i, j) &\approx u(i, j + 1); & v(i, j) &\approx v(i, j + 1); \end{aligned}$$

Or, equivalently, in continuous form:

$$u_x \approx u_y \approx 0; \quad v_x \approx v_y \approx 0$$

- Enforcing smoothness at each pixel equivalents to minimize:

$$E_s = \sum_{i,j} \left[(u_{ij} - u_{i+1,j})^2 + (u_{ij} - u_{i,j+1})^2 + (v_{ij} - v_{i+1,j})^2 + (v_{ij} - v_{i,j+1})^2 \right]$$

- Sometimes, written in integration form:

$$E_s = \iint [|\nabla u|^2 + |\nabla v|^2] dx dy = \iint [u_x^2 + u_y^2 + v_x^2 + v_y^2] dx dy$$



Horn & Schunck Method

Minimize the combined energy

$$\{u, v\} = \arg \min(\lambda E_d + E_s)$$

- In discrete form:

$$E_d = \sum_{i,j} [I_x(i,j)u(i,j) + I_y(i,j)v(i,j) + I_t(i,j)]^2$$

$$E_s = \sum_{i,j} [(u_{ij} - u_{i+1,j})^2 + (u_{ij} - u_{i,j+1})^2 + (v_{ij} - v_{i+1,j})^2 + (v_{ij} - v_{i,j+1})^2]$$

- In continuous form:

$$\lambda E_d + E_s = \iint \lambda(I_x u + I_y v + I_t)^2 + (|\nabla u|^2 + |\nabla v|^2) dx dy$$



Horn & Schunck Method

Minimize by set partial derivative to 0:

- In discrete form: $\frac{\partial(E_d + E_s)}{\partial u_{ij}} = 0$

$$\frac{\partial(E_d + E_s)}{\partial u_{ij}} = 2(u_{ij} - \bar{u}_{ij}) + 2\lambda(I_x u_{ij} + I_y v_{ij} + I_t)I_x = 0$$

$$\frac{\partial(E_d + E_s)}{\partial v_{ij}} = 2(v_{ij} - \bar{v}_{ij}) + 2\lambda(I_x u_{ij} + I_y v_{ij} + I_t)I_y = 0$$

$$\text{Local average: } \bar{u}_{ij} = \frac{1}{4}\{u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1}\}$$

After simple manipulations,

$$(1 + \lambda I_x^2)u_{ij} + \lambda I_x I_y v_{ij} = \bar{u}_{ij} - \lambda I_x I_t$$

$$\lambda I_x I_y u_{ij} + (1 + \lambda I_y^2)v_{ij} = \bar{v}_{ij} - \lambda I_y I_t$$

This becomes a large linear system

$$A\mathbf{x} = \mathbf{b}$$

A is a large sparse matrix!



Horn & Schunck Method

A simple iterative solver:

$$(1 + \lambda I_x^2) u_{ij}^{n+1} + \lambda I_x I_y v_{ij}^{n+1} = \bar{u}_{ij}^n - \lambda I_x I_t$$
$$\lambda I_x I_y u_{ij}^{n+1} + (1 + \lambda I_y^2) v_{ij}^{n+1} = \bar{v}_{ij}^n - \lambda I_y I_t$$

- Solve a 2×2 inverse matrix at each pixel

$$u_{ij}^{n+1} = \bar{u}_{ij}^n - \frac{I_x \bar{u}_{ij}^n + I_y \bar{v}_{ij}^n + I_t}{\lambda^{-1} + (I_x^2 + I_y^2)} I_x$$
$$v_{ij}^{n+1} = \bar{v}_{ij}^n - \frac{I_x \bar{u}_{ij}^n + I_y \bar{v}_{ij}^n + I_t}{\lambda^{-1} + (I_x^2 + I_y^2)} I_y$$



Summary

1. Precompute image gradient I_x, I_y
2. Precompute temporal gradient I_t
3. Initialize flow $u = v = 0$
4. While not converge

Compute flow field updates for each pixel:

$$u_{ij}^{n+1} = \bar{u}_{ij}^n - \frac{I_x \bar{u}_{ij}^n + I_y \bar{v}_{ij}^n + I_t}{1 + \lambda(I_x^2 + I_y^2)} I_x$$

$$v_{ij}^{n+1} = \bar{v}_{ij}^n - \frac{I_x \bar{u}_{ij}^n + I_y \bar{v}_{ij}^n + I_t}{1 + \lambda(I_x^2 + I_y^2)} I_y$$



Horn & Schunck Method

- The continuous version can be minimized similarly

- $$E = \lambda E_d + E_s = \iint \lambda(I_x u + I_y v + I_t)^2 + (|\nabla u|^2 + |\nabla v|^2) dx dy$$

- Want to evaluate $\frac{\partial E}{\partial u}$, but u itself is a function

- E is a function of a function, called functional

- Minimize a functional = solve its Euler equation $\frac{\partial E}{\partial u} = 0$

- In our case, the Euler equations are:

$$\Delta u - \lambda(I_x u + I_y v + I_t) = 0,$$

$$\Delta v - \lambda(I_x u + I_y v + I_t) = 0.$$

$$\Delta u = \bar{u} - u$$

Questions?



Modern Techniques for Optical Flow



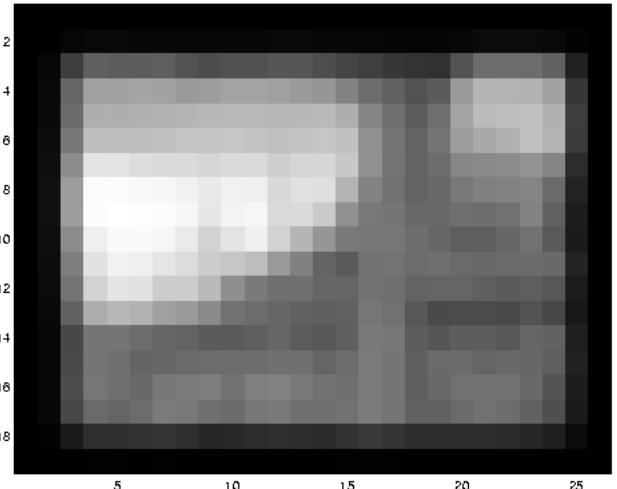
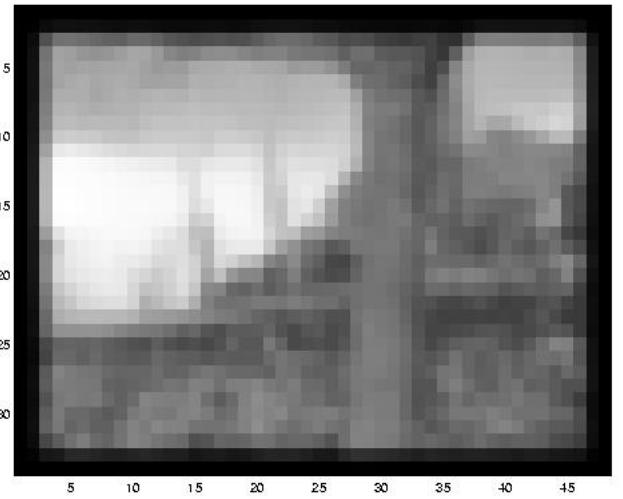
- Warping
- Gradient constancy
- Robust cost function
- Median / Bilateral filter

Warping

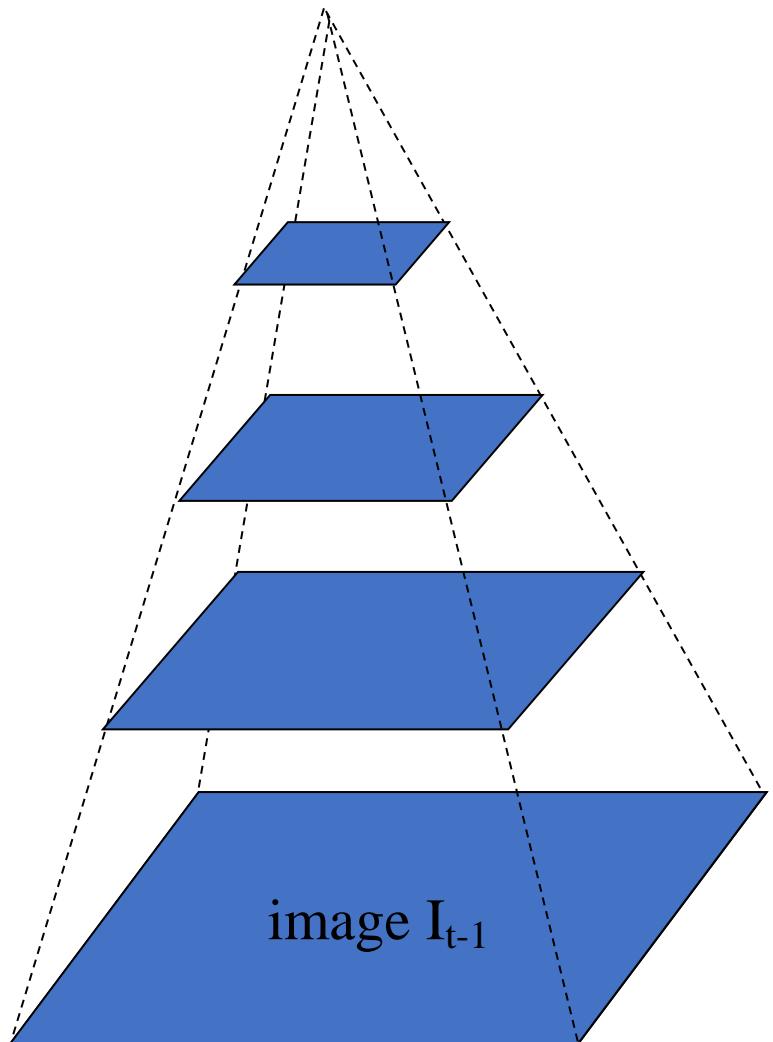


- Revisiting the small motion assumption
- Is this motion small enough?
 - Probably not—it's much larger than one pixel
 - How might we solve this problem?

reduce the resolution!



coarse-to-fine optical flow estimation



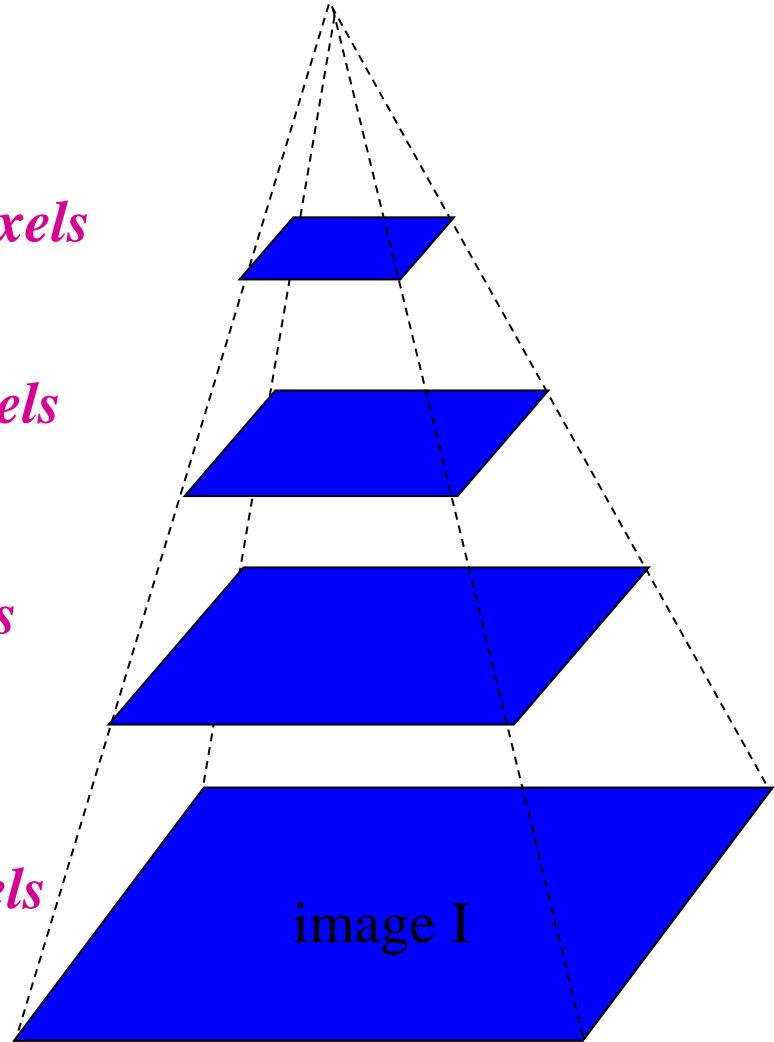
Gaussian pyramid of image I_{t-1}

$u=1.25 \text{ pixels}$

$u=2.5 \text{ pixels}$

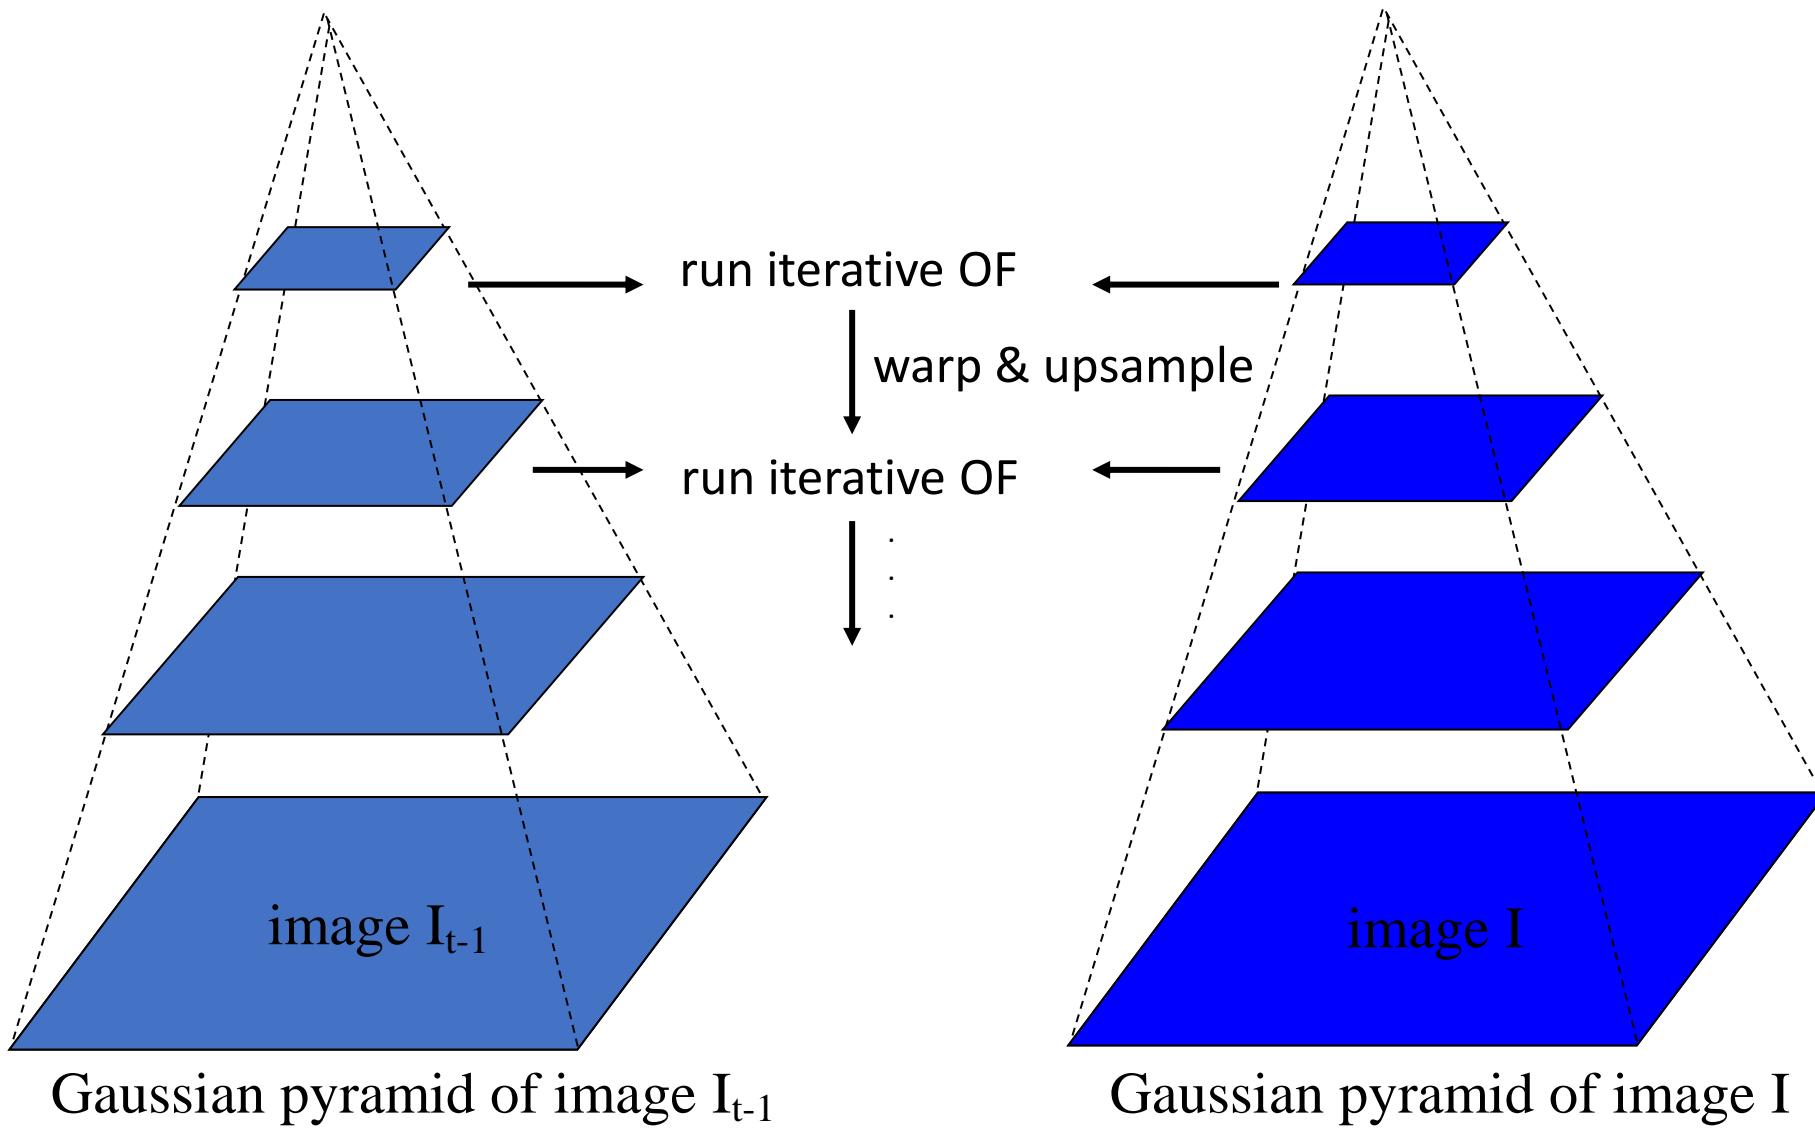
$u=5 \text{ pixels}$

$u=10 \text{ pixels}$



Gaussian pyramid of image I

coarse-to-fine optical flow estimation

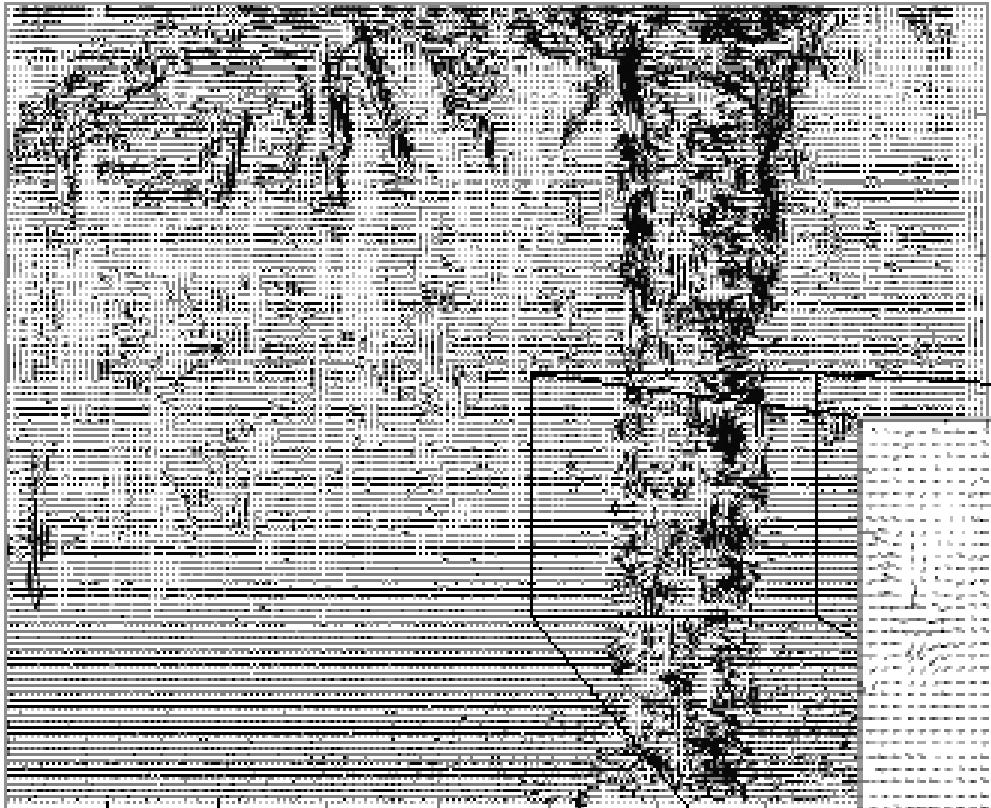




Warping based Optical Flow

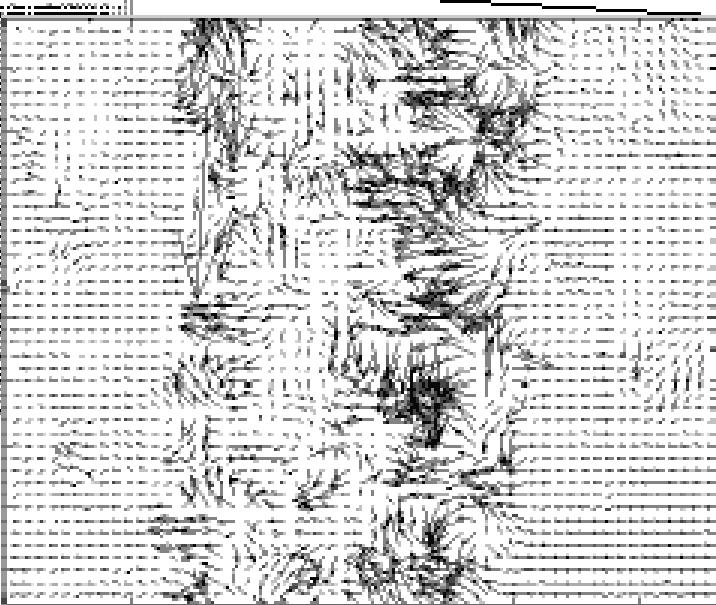
1. Compute optical flow at the highest pyramid level
2. At level i
 1. Take flow u_{i-1}, v_{i-1} from level $i - 1$
 2. Interpolate it to create u_i^*, v_i^* of twice resolution for level i
 3. Multiply u_i^*, v_i^* by 2
 4. Warp $I(t - 1)$ by u_i^*, v_i^* , denote the warped image as $I^*(t - 1)$
 5. Compute correction flow u'_i, v'_i between $I^*(t - 1)$ and $I(t)$
 6. Add corrections, i.e. $u_i = u_i^* + u'_i; v_i = v_i^* + v'_i$.

examples

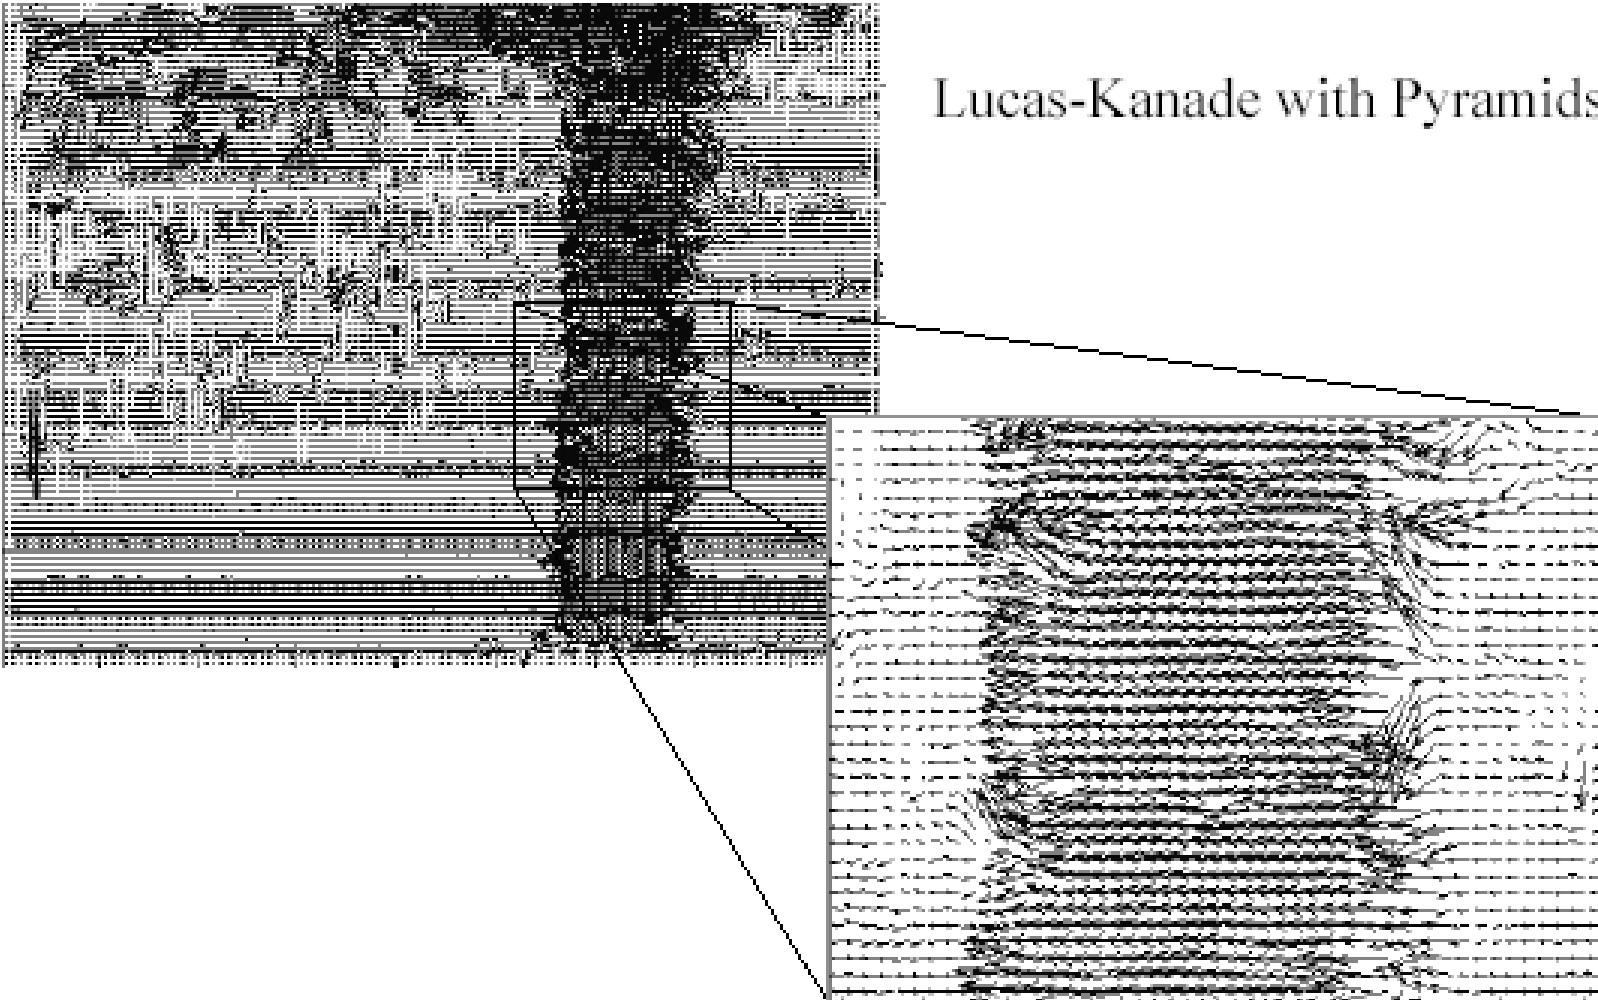


Lucas-Kanade
without pyramids

Fails in areas of large
motion



examples



Lucas-Kanade with Pyramids



Gradient Constancy

- The brightness constancy assumption might be wrong
 - Due to illumination/exposure/white balance changes
- Add gradient constancy

$$\nabla I_t(x, y) = \nabla I_{t+1}(x + d_x, y + d_y)$$

- Robust to illumination changes
- But violated by rotations
- It might be better to switch between color and gradient constancy instead of directly combine them

Motion Detail Preserving Optical Flow Estimation*

Li Xu Jiaya Jia Yasuyuki Matsushita
The Chinese University of Hong Kong Microsoft Research Asia
{xuli,leojia}@cse.cuhk.edu.hk yasumat@microsoft.com

[CVPR 2010]



Robust Cost Function

- Instead of enforcing L2 penalty on color constancy, M-estimator can be used here
 - Recall the Horn-Schunk formulation

$$\iint (I_x u + I_y v + I_t)^2 + (|\nabla u|^2 + |\nabla v|^2) dx dy$$

- A commonly used loss is L1 penalty, or the Charbonnier penalty

$$\rho(x) = \sqrt{x^2 + \epsilon^2}$$

- Can be used for both data and smoothness terms



Median/Bilateral Filter

- Top secrets of optical flow estimation: applying a median filter to the flow after the flow update (step 6 in page 82)
 - It removes outliers and preserves edges
 - A bilateral filter gives even better result
- Effectively, this changes the objective function to:

$$\frac{E_d(u, v) + \lambda_1 E_s(u, v) + \lambda_2 (|u - \hat{u}| + |v - \hat{v}|)}{\lambda_3 \sum_i \sum_{j \in N_i} \omega_{ij} (|\hat{u}_i - \hat{u}_j| + |\hat{v}_i - \hat{v}_j|)}$$

Energy formulation of the filtering
A better smoothness term

Coupling term



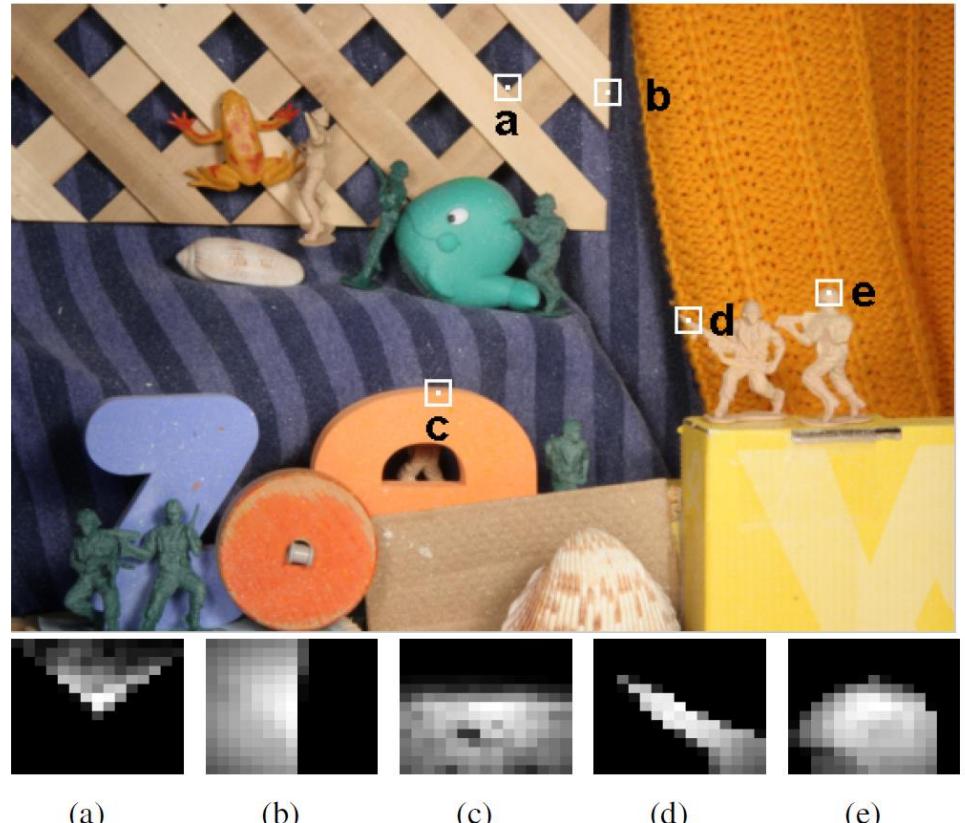
Median/Bilateral Filter

- The bilateral weights

$$\omega_{ij} = g_1(|i - j|)g_2(I_i - I_j)$$

Spatial Gauss

Range Gauss



Secrets of Optical Flow Estimation and Their Principles

Deqing Sun
Brown University

Stefan Roth
TU Darmstadt

Michael J. Black
Brown University
[CVPR 2010]

Questions?

