



Chapter 5: Advanced SQL

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan
See www.db-book.com for conditions on re-use



Chapter 5: Advanced SQL

- Accessing SQL From a Programming Language
- Functions and Procedures
- Triggers
- **Recursive Queries
- **Advanced Aggregation Features



Accessing SQL from a Programming Language

A database programmer must have access to a general-purpose programming language for at least two reasons

- Not all queries can be expressed in SQL, since SQL does not provide the full expressive power of a general-purpose language.
- Non-declarative actions -- such as printing a report, interacting with a user, or sending the results of a query to a graphical user interface -- cannot be done from within SQL.



Accessing SQL from a Programming Language (Cont.)

There are two approaches to accessing database from a general-purpose programming language

- **API** (Application Program Interface) - - A general-purpose program can connect to and communicate with a database server using a collection of functions.
- **Embedded SQL** -- provides a means by which a program can interact with a database server.
 - The SQL statements are translated at compile time into function calls.
 - At runtime, these function calls connect to the database using an API that provides dynamic SQL facilities.



JDBC and ODBC

- **API (application-program interface)** for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- **ODBC** (Open Database Connectivity) works with C, C++, C#
- **JDBC** (Java Database Connectivity) works with Java
- **Embedded SQL** in C
- **SQLJ** - embedded SQL in Java
- **JPA(Java Persistence API)** - OR mapping of Java



JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports **metadata retrieval**, such as querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
 - **Open a connection**
 - **Create a “statement” object**
 - **Execute queries** using the Statement object to send queries and fetch results
 - **Exception mechanism** to handle errors



JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try {
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement();
        ... Do Actual Work ....
        stmt.close();
        conn.close();
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```



JDBC Code (Cont.)

■ Update to database

```
try {  
    stmt.executeUpdate(  
        "insert into instructor values('77987', 'Kim', 'Physics', 98000)");  
} catch (SQLException sqle)  
{  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

■ Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg (salary)  
    from instructor  
    group by dept_name");  
while (rset.next()) {  
    System.out.println(rset.getString("dept_name") + " " +  
        rset.getFloat(2));  
}
```




JDBC Code Details

- Getting result fields:
 - `rset.getString("dept_name")` and `rset.getString(1)` equivalent if `dept_name` is the first argument of select result.
- Dealing with Null values
 - `int a = rset.getInt("a");`
`if (rset.wasNull()) Systems.out.println("Got null value");`



Prepared Statement

- **PreparedStatement** pStmt = conn.**prepareStatement**(
 "insert into instructor values(?,?,?,?)");
pStmt.**setString**(1, "88877"); pStmt.**setString**(2, "Perry");
pStmt.**setString**(3, "Finance"); pStmt.**setInt**(4, 125000);
pStmt.**executeUpdate**();
pStmt.**setString**(1, "88878");
pStmt.**executeUpdate**();

- **WARNING:** always use prepared statements when taking an input from the user and adding it to a query
 - **NEVER create a query by concatenating strings which you get as inputs**
 - "insert into instructor values(' " + ID + " ', ' " + name + " ', " +
 " ' " + dept name + " ', " + salary + ")"
 - What if name is "D'Souza"?



SQL Injection(SQL 注入)

- Suppose query is constructed using
 - "select * from instructor where name = '" + name + "'"
- Suppose the user, instead of entering a name, enters:
 - X' or 'Y' = 'Y'
- then the resulting string of the statement becomes:
 - "select * from instructor where name = '" + "X' or 'Y' = 'Y'" + ""
 - which is:
 - select * from instructor where name = 'X' or 'Y' = 'Y'
 - User could have even used
 - X'; update instructor set salary = salary + 10000;
- then →
 - select * from instructor where name = 'X';
update instructor set salary = salary + 10000;
- **Always use prepared statements, with user inputs as parameters**



Metadata Features

- **ResultSet metadata**

- E.g., after executing query to get a **ResultSet** rs:

- **ResultSetMetaData** rsmd = rs.**getMetaData()**;
 for(int i = 1; i <= rsmd.**getColumnCount()**; i++) {
 System.out.println(rsmd.**getColumnName(i)**);
 System.out.println(rsmd.**getColumnTypeName(i)**);
 }

- How is this useful?



Metadata (Cont)

- Database metadata

- `DatabaseMetaData dbmd = conn.getMetaData();`

```
ResultSet rs = dbmd.getColumns(null, "univdb", "department", "%");
```

```
// Arguments to getColumns: Catalog, Schema-pattern, Table-pattern,  
// and Column-Pattern
```

```
// Returns: One row for each column; row has a number of attributes
```

```
// such as COLUMN_NAME, TYPE_NAME
```

```
while( rs.next()) {
```

```
    System.out.println(rs.getString("COLUMN_NAME"),
```

```
                        rs.getString("TYPE_NAME");
```

```
}
```

- And where is this useful?



Transaction Control in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
 - bad idea for transactions with multiple updates
- Can turn off automatic commit on a connection
 - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
 - `conn.commit();` or
 - `conn.rollback();`
- `conn.setAutoCommit(true)` turns on automatic commit.



JDBC Resources

- JDBC Basics Tutorial
 - <https://docs.oracle.com/javase/tutorial/jdbc/index.html>



- ```
String deptName = iter.dept_name();
int avgSal = iter.avgSal();
System.out.println(deptName + " " + avgSal);
}
iter.close();
```





# SQLJ Resources

- SQLJ Developer's Guide
- <https://docs.oracle.com/en/database/oracle/oracle-database/20/jsqlj/toc.htm>



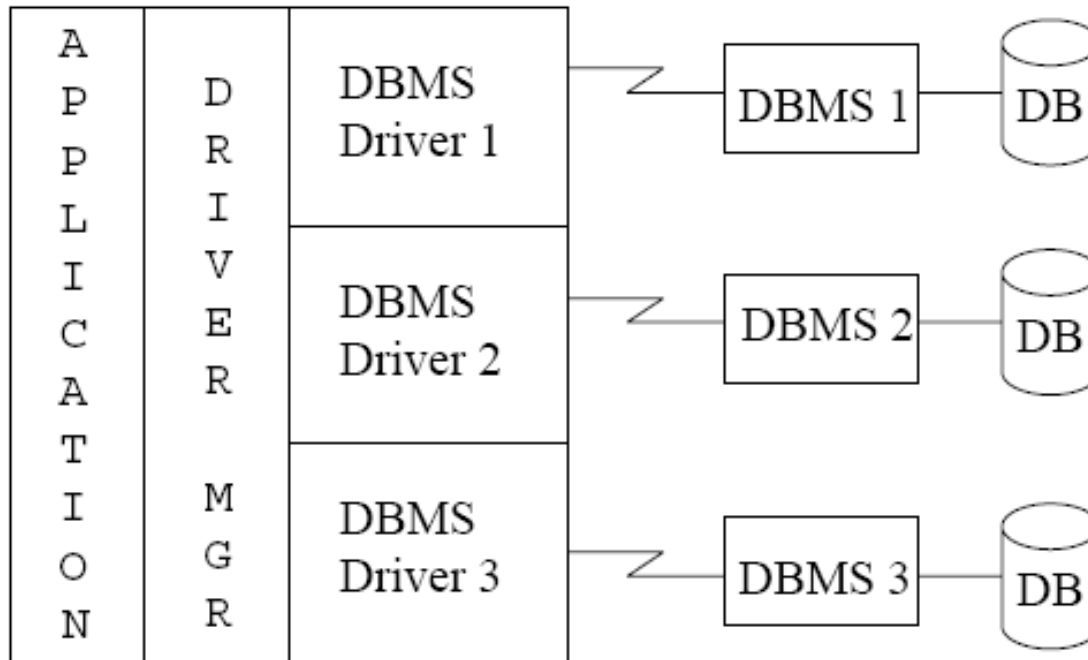
# ODBC

- **Open DataBase Connectivity(ODBC)** standard
  - standard for application program to communicate with a database server.
  - application program interface (API) to
    - ▶ open a connection with a database,
    - ▶ send queries and updates,
    - ▶ get back results.
- Applications such as **GUI, spreadsheets**, etc. can use ODBC
- Was defined originally for **Basic** and **C**, versions available for many languages.



# ODBC (Cont.)

- Each database system supporting ODBC provides a "driver" library that must be linked with the client program.





# ODBC (Cont.)

- Each database system supporting ODBC provides a "driver" library that must be linked with the client program.
- When client program makes an ODBC API call, the code in the library communicates with the server to carry out the requested action, and fetch results.
- ODBC program first allocates an SQL environment, then a database connection handle.
- Opens database connection using SQLConnect(). Parameters for SQLConnect:
  - connection handle,
  - the server to which to connect
  - the user identifier,
  - password
- Must also specify types of arguments:
  - SQL\_NTS denotes previous argument is a null-terminated string.



# ODBC Code

```
■ int ODBCexample()
{
 RETCODE error;
 HENV env; /* environment */
 HDBC conn; /* database connection */
 SQLAllocEnv(&env);
 SQLAllocConnect(env, &conn);
 SQLConnect(conn, "db.yale.edu", SQL_NTS, "avi", SQL_NTS,
 "avipasswd", SQL_NTS);
 { Do actual work ... }

 SQLDisconnect(conn);
 SQLFreeConnect(conn);
 SQLFreeEnv(env);
}
```



# ODBC Code (Cont.)

- Program sends SQL commands to database by using `SQLExecDirect`
- Result tuples are fetched using `SQLFetch()`
- `SQLBindCol()` binds C language variables to attributes of the query result
  - When a tuple is fetched, its attribute values are automatically stored in corresponding C variables.
  - Arguments to `SQLBindCol()`
    - ▶ ODBC stmt variable, attribute position in query result
    - ▶ The type conversion from SQL to C.
    - ▶ The address of the variable.
    - ▶ For variable-length types like character arrays,
      - The maximum length of the variable
      - Location to store actual length when a tuple is fetched.
      - Note: A negative value returned for the length field indicates null value
- Good programming requires checking results of every function call for errors; we have omitted most checks for brevity.



# ODBC Code (Cont.)

## ■ Main body of program

```
char deptname[80];
float salary;
int lenOut1, lenOut2;
HSTMT stmt;
char * sqlquery = "select dept_name, sum (salary)
 from instructor
 group by dept_name";
SQLAllocStmt(conn, &stmt);
error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
if (error == SQL_SUCCESS) {
 SQLBindCol(stmt, 1, SQL_C_CHAR, deptname, 80, &lenOut1);
 SQLBindCol(stmt, 2, SQL_C_FLOAT, &salary, 0, &lenOut2);
 while (SQLFetch(stmt) == SQL_SUCCESS) {
 printf (" %s %g\n", deptname, salary);
 }
}
SQLFreeStmt(stmt, SQL_DROP);
```



# ODBC Prepared Statements

## ■ Prepared Statement

- SQL statement prepared: compiled at the database
- Can have placeholders: E.g. insert into account values(?,?,?)
- Repeatedly executed with actual values for the placeholders

## ■ To prepare a statement

`SQLPrepare(stmt, <SQL String>);`

## ■ To bind parameters

`SQLBindParameter(stmt, <parameter#>, ... type information and value omitted for simplicity..)`

## ■ To execute the statement

`retcode = SQLExecute( stmt);`





# More ODBC Features

## ■ Metadata features

- finding all the relations in the database and
- finding the names and types of columns of a query result or a relation in the database.

## ■ By default, each SQL statement is treated as a separate transaction that is committed automatically.

- Can turn off automatic commit on a connection
  - ▶ `SQLSetConnectOption(conn, SQL_AUTOCOMMIT, 0)`
- Transactions must then be committed or rolled back explicitly by
  - ▶ `SQLTransact(conn, SQL_COMMIT)` or
  - ▶ `SQLTransact(conn, SQL_ROLLBACK)`



# ODBC Resources

- Oracle Database Programmer's Guide - Embedded SQL
- [https://docs.oracle.com/cd/E17952\\_01/connector-odbc-en/index.html](https://docs.oracle.com/cd/E17952_01/connector-odbc-en/index.html)



# Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as **C**, **C++**, **Java**, **Fortran**, and **PL/1**,
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- The basic form of these languages follows that of the System R embedding of SQL into PL/1.
- **EXEC SQL** statement is used in the host language to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement >;

Note: this varies by language:

- In some languages, like **COBOL**, the semicolon is replaced with END-EXEC
- In **Java** embedding uses **# SQL { .... }**;



# Embedded SQL (Cont.)

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL **connect to** *server* **user** *user-name* **using** *password*;

Here, *server* identifies the server to which a connection is to be established.

- Variables of the **host language** can be used within embedded SQL statements. They are preceded by a colon (:) to distinguish from SQL variables (e.g., *:credit\_amount* )
- **Host Variables** used as above must be declared within **DECLARE** section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax.

**EXEC-SQL BEGIN DECLARE SECTION**

*int credit\_amount* ;

**EXEC-SQL END DECLARE SECTION;**



# Embedded SQL (Cont.)

- To write an embedded SQL query, we use the  
**declare c cursor for <SQL query>**  
statement. The variable *c* is used to identify the query
- Example:
  - From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable *credit\_amount* in the host language
  - Specify the query in SQL as follows:

EXEC SQL

```
declare c cursor for
select ID, name
from student
where tot_cred > :credit_amount;
```



# Embedded SQL (Cont.)

- The **open** statement for our example is as follows:

```
EXEC SQL open c ;
```

This statement causes the database system to execute the query and to save the results within a temporary relation. The query uses the value of the host-language variable *credit-amount* at the time the **open** statement is executed.

- The fetch statement causes the values of one tuple in the query result to be placed on host language variables.

```
EXEC SQL fetch c into :si, :sn ;
```

Repeated calls to fetch get successive tuples in the query result



## Embedded SQL (Cont.)

- A variable called **SQLSTATE** in the SQL communication area (**SQLCA**) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

```
EXEC SQL close c ;
```



# Updates Through Embedded SQL

- Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- Can update tuples fetched by cursor by declaring that the cursor is for update

## EXEC SQL

```
declare c cursor for
select *
from instructor
where dept_name = 'Music'
for update
```

- We then iterate through the tuples by performing **fetch** operations on the cursor , and after fetching each tuple we execute the following code:

```
update instructor
set salary = salary + 1000
where current of c
```





# Embedded SQL Resources

- Oracle Database Programmer's Guide - Embedded SQL
- <https://docs.oracle.com/en/database/oracle/oracle-database/20/Inpcc/embedded-SQL.html#GUID-C671CABF-202A-4503-A16B-DC78D3F1AB13>



# Procedural Constructs in SQL



# Procedural Extensions and Stored Procedures

- SQL provides a **module** language
  - Permits definition of procedures in SQL, with **if-then-else** statements, **for** and **while** loops, etc.
- **Stored Procedures**
  - Can store procedures in the database
  - then execute them using the **call** statement
  - permit external applications to operate on the database without knowing about internal details



# Functions and Procedures

- Functions and procedures allow “business logic” to be stored in the database and executed from SQL statements.
- These can be defined either by the procedural component of SQL or by an external programming language such as Java, C, or C++.
- The syntax we present here is defined by the SQL standard.
  - Most databases implement nonstandard versions of this syntax.



# SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))
returns integer
begin
 declare d_count integer;
 select count (*) into d_count
 from instructor
 where instructor.dept_name = dept_name
 return d_count;
end
```

- Find the department name and budget of all departments with more than 12 instructors.

```
select dept_name, budget
from department
where dept_count (dept_name) > 1
```



# Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all accounts owned by a given customer

**create function** *instructors\_of* (*dept\_name* **char**(20) )

**returns table** ( *ID* **varchar**(5),  
                  *name* **varchar**(20),  
                  *dept\_name* **varchar**(20),  
                  *salary* **numeric**(8,2))

**return table**

(**select** *ID, name, dept\_name, salary*  
**from** *instructor*  
**where** *instructor.dept\_name = instructors\_of.dept\_name*)

- Usage

**select** \*  
**from table** (*instructors\_of* ('Music'))



# SQL Procedures

- The *dept\_count* function could instead be written as procedure:  
**create procedure** *dept\_count\_proc* (**in** *dept\_name* **varchar**(20),  
**out** *d\_count* **integer**)

**begin**

```
select count(*) into d_count
from instructor
where instructor.dept_name = dept_count_proc.dept_name
```

**end**

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare d_count integer;
call dept_count_proc('Physics', d_count);
```

Procedures and functions can be invoked also from dynamic SQL



# Procedural Constructs

- Warning: most database systems implement their own variant of the standard syntax below
  - read your system manual to see what works on your system
- Compound statement: **begin ... end**,
  - May contain multiple SQL statements between **begin** and **end**.
  - Local variables can be declared within a compound statements

- **While** and **repeat** statements :

```
declare n integer default 0;
```

```
while $n < 10$ do
```

```
 set $n = n + 1$
```

```
end while
```

```
 repeat
```

```
 set $n = n - 1$
```

```
 until $n = 0$
```

```
end repeat
```





# Procedural Constructs (Cont.)

- **For** loop
  - Permits iteration over all results of a query
  - Example:

```
declare n integer default 0;
for r as
 select budget from department
 where dept_name = 'Music'
do
 set n = n - r.budget
end for
```



# Procedural Constructs (cont.)

- Conditional statements (**if-then-else**)

```
 if boolean expression
 then statement or compound statement
elseif boolean expression
 then statement or compound statement
 else statement or compound statement
end if
```

- SQL:1999 also supports a **case** statement similar to C case statement



# Example procedure

```
create function registerStudent(
 in s_id varchar(5),
 in s_courseid varchar (8),
 in s_secid varchar (8),
 in s_semester varchar (6),
 in s_year numeric (4,0),
 out errorMsg varchar(100)

returns integer
begin
 declare currEnrol int;
 select count(*) into currEnrol
 from takes
 where course_id = s_courseid and sec_id = s_secid
 and semester = s_semester and year = s_year;
 declare limit int;
 select capacity into limit
 from classroom natural join section
 where course_id = s_courseid and sec_id = s_secid
 and semester = s_semester and year = s_year;
 if (currEnrol < limit)
 begin
 insert into takes values
 (s_id, s_courseid, s_secid, s_semester, s_year, null);
 return(0);
 end
 -- Otherwise, section capacity limit already reached
 set errorMsg = 'Enrollment limit reached for course ' || s_courseid
 || ' section ' || s_secid;
 return(-1);
end;
```

- Registers a student after ensuring
- classroom capacity is not exceeded.
- Returns 0 on success, and -1 if
- capacity is exceeded.



# External Language Functions/Procedures

- SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- Declaring external language procedures and functions

```
create procedure dept_count_proc(in dept_name varchar(20),
 out count integer)
```

```
language C
external name ' /usr/avi/bin/dept_count_proc'
```

```
create function dept_count(dept_name varchar(20))
returns integer
language C
external name '/usr/avi/bin/dept_count'
```



# External Language Routines (Cont.)

- Benefits of external language functions/procedures:
  - more efficient for many operations, and more expressive power.
- Drawbacks
  - Code to implement function may need to be loaded into database system and executed in the database system's address space.
    - ▶ risk of accidental corruption of database structures
    - ▶ security risk, allowing users access to unauthorized data
  - There are alternatives, which give good security at the cost of potentially worse performance.
  - Direct execution in the database system's space is used when efficiency is more important than security.



# Security with External Language Routines

- To deal with security problems
  - Use **sandbox** techniques
    - ▶ that is use a safe language like Java, which cannot be used to access/damage other parts of the database code.
  - Or, run external language functions/procedures in a separate process, with no access to the database process' memory.
    - ▶ Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space.



# Triggers



# Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- Trigger - ECA rule
  - E: Event ( **insert, delete , update** )
  - C: Condition
  - A: Action
- To design a trigger mechanism, we must:
  - Specify the conditions under which the trigger is to be executed.
  - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.





# Trigger Example

- `account_log(account, amount, datetime)`
- **create trigger *account\_trigger* after update on *account(balance)* referencing new row as *nrow***

**referencing old row as *orow***

**for each row**

**when *nrow.balance* - *orow.balance* > =200000 or**

***orow.balance* -*nrow.balance* >=50000**

**begin**

**insert into *account\_log* values (*nrow.account-number*,  
*nrow.balance-orow.balance* , *current\_time()* )**

**end**



# Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
  - **E.g., after update of *takes* on *grade***
- Values of attributes before and after an update can be referenced
  - **referencing old row as** : for deletes and updates
  - **referencing new row as** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints. E.g. convert blank grades to null.

```
create trigger setnull_trigger before update of takes
referencing new row as nrow
for each row
when (nrow.grade = ' ')
 begin atomic
 set nrow.grade = null;
 end;
```



# Trigger to Maintain `credits_earned` value

- **create trigger *credits\_earned* after update of *takes* on (*grade*)**  
**referencing new row as *nrow***  
**referencing old row as *orow***  
**for each row**  
**when *nrow.grade* <> 'F' and *nrow.grade* is not null**  
**and (*orow.grade* = 'F' or *orow.grade* is null)**  
**begin atomic**  
**update *student***  
**set *tot\_cred*= *tot\_cred* +**  
**(select *credits***  
**from *course***  
**where *course.course\_id*= *nrow.course\_id*)**  
**where *student.id* = *nrow.id*;**  
**end;**



# Statement Level Triggers

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
  - Use **for each statement** instead of **for each row**
  - Use **referencing old table** or **referencing new table** to refer to temporary tables (called *transition tables*) containing the affected rows
  - Can be more efficient when dealing with SQL statements that update a large number of rows



# Statement Level Triggers

- `register(sid, __ cno, score)`
- **create trigger *score\_trigger* after update on *register(score)* referencing new table as *ntable***

```
for each statement
 when some(select average(score)
 from ntable
 group by cno) < 60
 begin
 rollback
 end
```



# When Not To Use Triggers

- Triggers were used earlier for tasks such as
  - maintaining summary data (e.g., total salary of each department)
  - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
  - Databases today provide built in **materialized view** facilities to maintain summary data
  - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
  - Define methods to update fields
  - Carry out actions as part of the update methods instead of through a trigger



# When Not To Use Triggers

- Risk of unintended execution of triggers, for example, when
  - loading data from a backup copy
  - replicating updates at a remote site
  - Trigger execution can be disabled before such actions.
- Other risks with triggers:
  - Error leading to failure of critical transactions that set off the trigger
  - Cascading execution



# Recursive Queries





# Recursion in SQL

- SQL:1999 permits recursive view definition
- Example: find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
with recursive rec_prereq(course_id, prereq_id) as (
 select course_id, prereq_id
 from prereq
 union
 select rec_prereq.course_id, prereq.prereq_id,
 from rec_prereq, prereq
 where rec_prereq.prereq_id = prereq.course_id
)
select *
from rec_prereq;
```

This example view, *rec\_prereq*, is called the *transitive closure* of the *prereq* relation



# The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
  - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *prereq* with itself
    - ▶ This can give only a fixed number of levels of prerequisites.
    - ▶ Given a fixed non-recursive query, we can construct a database with a greater number of levels of prerequisites on which the query will not work
    - ▶ Alternative: write a procedure to iterate as many times as required
      - See procedure *findAllPrereqs* in book on Page.189



# The Power of Recursion

```
create function findAllPrereqs(cid varchar(8))
-- Finds all courses that are prerequisite (directly or indirectly) for cid
returns table (course_id varchar(8))
-- The relation prereq(course_id, prereq_id) specifies which course is
-- directly a prerequisite for another course.
begin
 create temporary table c_prereq (course_id varchar(8));
 -- table c_prereq stores the set of courses to be returned
 create temporary table new_c_prereq (course_id varchar(8));
 -- table new_c_prereq contains courses found in the previous iteration
 create temporary table temp (course_id varchar(8));
 -- table temp is used to store intermediate results
 insert into new_c_prereq
 select prereq_id
 from prereq
 where course_id = cid;
 repeat
 insert into c_prereq
 select course_id
 from new_c_prereq;

 insert into temp
 (select prereq.course_id
 from new_c_prereq, prereq
 where new_c_prereq.course_id = prereq.prereq_id
)
 except (
 select course_id
 from c_prereq
);
 delete from new_c_prereq;
 insert into new_c_prereq
 select *
 from temp;
 delete from temp;

 until not exists (select * from new_c_prereq)
 end repeat;
 return table c_prereq;
end
```



# OLAP\*\*



# Data Analysis and OLAP

## ■ Online Analytical Processing (OLAP)

- Interactive analysis of data, allowing data to be summarized and viewed in different ways in an online fashion (with negligible delay)

## ■ Data that can be modeled as dimension attributes and measure attributes are called **multidimensional data**.

### ● **Measure attributes**

- ▶ measure some value
- ▶ can be aggregated upon
- ▶ e.g., the attribute *number* of the *sales* relation

### ● **Dimension attributes**

- ▶ define the dimensions on which measure attributes (or aggregates thereof) are viewed
- ▶ e.g., attributes *item\_name*, *color*, and *size* of the *sales* relation



## Example sales relation

| <i>item_name</i> | <i>color</i> | <i>clothes_size</i> | <i>quantity</i> |
|------------------|--------------|---------------------|-----------------|
| skirt            | dark         | small               | 2               |
| skirt            | dark         | medium              | 5               |
| skirt            | dark         | large               | 1               |
| skirt            | pastel       | small               | 11              |
| skirt            | pastel       | medium              | 9               |
| skirt            | pastel       | large               | 15              |
| skirt            | white        | small               | 2               |
| skirt            | white        | medium              | 5               |
| skirt            | white        | large               | 3               |
| dress            | dark         | small               | 2               |
| dress            | dark         | medium              | 6               |
| dress            | dark         | large               | 12              |
| dress            | pastel       | small               | 4               |
| dress            | pastel       | medium              | 3               |
| dress            | pastel       | large               | 3               |
| dress            | white        | small               | 2               |
| dress            | white        | medium              | 3               |
| dress            | white        | large               | 0               |
| shirt            | dark         | small               | 2               |
| shirt            | dark         | medium              | 4               |
| ...              | ...          | ...                 | ...             |
| ...              | ...          | 5.76                | ...             |



# Cross Tabulation of sales by *item\_name* and color

*clothes\_size* **all**

|                  |       | <i>color</i> |        |       |       |
|------------------|-------|--------------|--------|-------|-------|
|                  |       | dark         | pastel | white | total |
| <i>item_name</i> | skirt | 8            | 35     | 10    | 53    |
|                  | dress | 20           | 10     | 5     | 35    |
|                  | shirt | 14           | 7      | 28    | 49    |
|                  | pants | 20           | 2      | 5     | 27    |
|                  | total | 62           | 54     | 48    | 164   |

- The table above is an example of a **cross-tabulation** (**cross-tab**), also referred to as a **pivot-table**.
  - Values for one of the dimension attributes form the row headers
  - Values for another dimension attribute form the column headers
  - Other dimension attributes are listed on top
  - Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.



# Data Cube

- A **data cube** is a multidimensional generalization of a cross-tab
- Can have  $n$  dimensions; we show 3 below
- Cross-tabs can be used as views on a data cube

|       |        |           |       |       |       |     |   |  |
|-------|--------|-----------|-------|-------|-------|-----|---|--|
|       |        | 2         |       |       |       |     |   |  |
|       |        | 4         |       |       |       |     |   |  |
|       |        | 2         |       |       |       |     |   |  |
| color | dark   | 8         | 20    | 14    | 20    | 62  | 3 |  |
|       | pastel | 35        | 10    | 7     | 2     | 54  | 2 |  |
|       | white  | 10        | 8     | 28    | 5     | 48  | 7 |  |
|       | all    | 53        | 38    | 49    | 27    | 164 | a |  |
|       |        | skirt     | dress | shirt | pants | all |   |  |
|       |        | item_name |       |       |       |     |   |  |





# Cross Tabulation With Hierarchy

- Cross-tabs can be easily extended to deal with hierarchies
  - Can drill down or roll up on a hierarchy

*clothes\_size:* **all**

|                 |                  | <i>color</i> |        |       |       |     |
|-----------------|------------------|--------------|--------|-------|-------|-----|
| <i>category</i> | <i>item_name</i> | dark         | pastel | white | total |     |
| womenswear      | skirt            | 8            | 8      | 10    | 53    | 88  |
|                 | dress            | 20           | 20     | 5     | 35    |     |
|                 | subtotal         | 28           | 28     | 15    |       |     |
| menswear        | pants            | 14           | 14     | 28    | 49    | 76  |
|                 | shirt            | 20           | 20     | 5     | 27    |     |
|                 | subtotal         | 34           | 34     | 33    |       |     |
| total           |                  | 62           | 62     | 48    |       | 164 |



# Relational Representation of Cross-tabs

- Cross-tabs can be represented as relations
  - We use the value **all** is used to represent aggregates.
  - The SQL standard actually uses null values in place of **all** despite confusion with regular null values.

| <i>item_name</i> | <i>color</i> | <i>clothes_size</i> | <i>quantity</i> |
|------------------|--------------|---------------------|-----------------|
| skirt            | dark         | <b>all</b>          | 8               |
| skirt            | pastel       | <b>all</b>          | 35              |
| skirt            | white        | <b>all</b>          | 10              |
| skirt            | <b>all</b>   | <b>all</b>          | 53              |
| dress            | dark         | <b>all</b>          | 20              |
| dress            | pastel       | <b>all</b>          | 10              |
| dress            | white        | <b>all</b>          | 5               |
| dress            | <b>all</b>   | <b>all</b>          | 35              |
| shirt            | dark         | <b>all</b>          | 14              |
| shirt            | pastel       | <b>all</b>          | 7               |
| shirt            | White        | <b>all</b>          | 28              |
| shirt            | <b>all</b>   | <b>all</b>          | 49              |
| pant             | dark         | <b>all</b>          | 20              |
| pant             | pastel       | <b>all</b>          | 2               |
| pant             | white        | <b>all</b>          | 5               |
| pant             | <b>all</b>   | <b>all</b>          | 27              |
| <b>all</b>       | dark         | <b>all</b>          | 62              |
| <b>all</b>       | pastel       | <b>all</b>          | 54              |
| <b>all</b>       | white        | <b>all</b>          | 48              |
| <b>all</b>       | <b>all</b>   | <b>all</b>          | 164             |



# Extended Aggregation to Support OLAP

- The **cube** operation computes union of **group by**'s on every subset of the specified attributes
- Example relation for this section  
*sales(item\_name, color, clothes\_size, quantity)*
- E.g. consider the query

```
select item_name, color, size, sum(number)
from sales
group by cube(item_name, color, size)
```

This computes the union of eight different groupings of the *sales* relation:

```
{ (item_name, color, size), (item_name, color),
 (item_name, size), (color, size),
 (item_name), (color),
 (size), () }
```

where ( ) denotes an empty **group by** list.

- For each grouping, the result contains the null value for attributes not present in the grouping.



## Extended Aggregation (Cont.)

- The **rollup** construct generates union on every prefix of specified list of attributes

- E.g.,

```
select item_name, color, size, sum(number)
from sales
group by rollup(item_name, color, size)
```

Generates union of four groupings:

$\{ (item\_name, color, size), (item\_name, color), (item\_name), ( ) \}$

- Rollup can be used to generate aggregates at multiple levels of a hierarchy.
- E.g., suppose table *itemcategory*(*item\_name*, *category*) gives the category of each item. Then

```
select category, item_name, sum(number)
from sales, itemcategory
where sales.item_name = itemcategory.item_name
group by rollup(category, item_name)
```

would give a hierarchical summary by *item\_name* and by *category*.



# Extended Aggregation (Cont.)

- Multiple rollups and cubes can be used in a single group by clause
  - Each generates set of group by lists, cross product of sets gives overall set of group by lists
- E.g.,

```
select item_name, color, size, sum(number)
from sales
group by rollup(item_name), rollup(color, size)
```

generates the groupings

$$\{item\_name, ()\} \times \{(color, size), (color), ()\}$$
$$= \{ (item\_name, color, size), (item\_name, color), (item\_name), (color, size), (color), () \}$$



# Online Analytical Processing Operations

- **Pivoting:** changing the dimensions used in a cross-tab is called
- **Slicing:** creating a cross-tab for fixed values only
  - Sometimes called **dicing**, particularly when values for multiple dimensions are fixed.
- **Rollup:** moving from finer-granularity data to a coarser granularity
- **Drill down:** The opposite operation - that of moving from coarser-granularity data to finer-granularity data



# End of Chapter

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



## Figure 5.22

| <i>item_name</i> | <i>clothes_size</i> | <i>dark</i> | <i>pastel</i> | <i>white</i> |
|------------------|---------------------|-------------|---------------|--------------|
| skirt            | small               | 2           | 11            | 2            |
| skirt            | medium              | 5           | 9             | 5            |
| skirt            | large               | 1           | 15            | 3            |
| dress            | small               | 2           | 4             | 2            |
| dress            | medium              | 6           | 3             | 3            |
| dress            | large               | 12          | 3             | 0            |
| shirt            | small               | 2           | 4             | 17           |
| shirt            | medium              | 6           | 1             | 1            |
| shirt            | large               | 6           | 2             | 10           |
| pant             | small               | 14          | 1             | 3            |
| pant             | medium              | 6           | 0             | 0            |
| pant             | large               | 0           | 1             | 2            |





## Figure 5.23

| <i>item_name</i> | <i>quantity</i> |
|------------------|-----------------|
| skirt            | 53              |
| dress            | 35              |
| shirt            | 49              |
| pant             | 27              |



# Figure 5.24

| <i>item_name</i> | <i>color</i> | <i>quantity</i> |
|------------------|--------------|-----------------|
| skirt            | dark         | 8               |
| skirt            | pastel       | 35              |
| skirt            | white        | 10              |
| dress            | dark         | 20              |
| dress            | pastel       | 10              |
| dress            | white        | 5               |
| shirt            | dark         | 14              |
| shirt            | pastel       | 7               |
| shirt            | white        | 28              |
| pant             | dark         | 20              |
| pant             | pastel       | 2               |
| pant             | white        | 5               |



# Another Recursion Example

- Given relation  
*manager(employee\_name, manager\_name)*
- Find all employee-manager pairs, where the employee reports to the manager directly or indirectly (that is manager's manager, manager's manager's manager, etc.)  

```
with recursive empl (employee_name, manager_name) as (
 select employee_name, manager_name
 from manager
union
 select manager.employee_name, empl.manager_name
 from manager, empl
 where manager.manager_name = empl.employee_name)
select *
from empl
```

This example view, *empl*, is the *transitive closure* of the *manager* relation



# Merge statement (now in Chapter 24)

- Merge construct allows batch processing of updates.
- Example: relation *funds\_received* (*account\_number*, *amount*) has batch of deposits to be added to the proper account in the *account* relation

```
merge into account as A
 using (select *
 from funds_received as F)
 on (A.account_number = F.account_number)
 when matched then
 update set balance = balance + F.amount
```